

Table of Contents

- [1 Final Project Submission](#)
 - [1.1 Summary:](#)
- [2 Import Libraries](#)
- [3 Define Functions](#)
 - [3.1 Melt](#)
 - [3.2 Stationarity Check](#)
- [4 Choose Zip Codes by Percent Increase in Median Price](#)
 - [4.1 Create Melted Dataframe of ZipCodes](#)
 - [4.2 Print Decomposition Chart for each Zip Code](#)
 - [4.3 Seasonality Trends Fascinating](#)
- [5 Adding Breadth of Zip Codes](#)
 - [5.1 Mountain Towns](#)
 - [5.1.1 Create Mountain Towns df](#)
 - [5.1.2 Summary:](#)
 - [5.1.2.1 Provo, Utah 84601 and 84606](#)
 - [5.1.2.2 Santa Fe, New Mexico 87506](#)
 - [5.1.2.3 Boise, Idaho 83703](#)
 - [5.2 Mid-Size Cities](#)
 - [5.2.1 Predicted Values](#)
 - [5.2.2 Bar Chart: Rate of Return](#)
 - [5.2.3 Summary:](#)
- [6 SARIMA Model](#)
 - [6.1 GridSearch optimal parameters, Model, Predict, calculate RMSE](#)
 - [6.2 RMSE and RMSE Ratio to Median Value](#)
 - [6.3 Create top_prospects df](#)
- [7 Price Forecast - Top Zip Codes](#)
 - [7.1 Calculate predicted rate of return after 5 years](#)
 - [7.2 Bar Chart: Rate of Return](#)
- [8 Recommendations](#)
 - [8.1 Oakland, CA 94606](#)
 - [8.2 Lake Worth, FL 33460](#)
 - [8.3 Boise, ID 83703](#)
- [9 Further Study](#)
 - [9.1 Adding to the model](#)
 - [9.2 Ongoing testing](#)
 - [9.3 Additional zip codes](#)
 - [9.4 Top returning zip codes](#)

1 Final Project Submission

- Student name: Jim Petoskey

- Student pace: Self-paced
- Scheduled project review date/time: Thursday May 5, 2022 at 9:00 am EST
- Instructor name: Abhineet Kulkarni
- [Blog post URL: Time Series Decomposition - Spotting Seasonality \(<https://dev.to/jpetoskey/time-series-decomposition-spotting-seasonality-4e71>\)](https://dev.to/jpetoskey/time-series-decomposition-spotting-seasonality-4e71)
- [Time Series Zillow Starter Repository \(<https://github.com/learn-co-curriculum/dsc-phase-4-choosing-a-dataset/tree/main/time-series>\)](https://github.com/learn-co-curriculum/dsc-phase-4-choosing-a-dataset/tree/main/time-series)

1.1 Summary:

- Top Zip Codes to Invest in:
 - Oakland, CA 94606
 - High Return (68.1%), Low Risk
 - Lake Worth, FL 33460
 - Highest Return (93%), Modest Risk
 - Boise, ID 83703
 - Moderately high return (55.1%), Low Risk
- Process:
 1. Measure rate of return for every zip code in US from April 2012 to April 2018
 2. Choose top 15 zip codes by rate of return
 3. Added top Mountain Towns and Mid-size Cities from Adding Breadth of Zip Codes Section
 4. Model price predictions, measure accuracy with RMSE
 5. Choose zip codes which are modeled most accurately
 6. Re-model those zip codes, forecast rate of return, plot predictions and confidence interval
 7. Propose investment in zip codes with lower risk, as measured by the confidence interval and RMSE, and high predicted returns over 5 year investment timeline

2 Import Libraries

In [4]:

```

1 # Import necessary libraries
2 import warnings
3 warnings.filterwarnings('ignore')
4 import itertools
5 import pandas as pd
6 import numpy as np
7 import statsmodels.api as sm
8 import matplotlib.pyplot as plt
9 from matplotlib.pylab import rcParams
10 plt.style.use('default')
11 from statsmodels.tsa.seasonal import seasonal_decompose

```

executed in 7ms, finished 10:45:47 2022-05-03

3 Define Functions

3.1 Melt

In [5]:

```

1 def melt_data(df):
2     """
3         Takes the zillow_data dataset in wide form or a subset of the zillow_data
4         Returns a long-form datetime dataframe
5         with the datetime column names as the index and the values as the 'values'
6
7         If more than one row is passes in the wide-form dataset, the values column
8         will be the mean of the values from the datetime columns in all of the row
9     """
10
11     melted = pd.melt(df, id_vars=['ZipCode', 'SizeRank', 'City', 'State', 'Metr
12     melted['time'] = pd.to_datetime(melted['time'], infer_datetime_format=True
13     melted = melted.dropna(subset=['value'])
14     return melted.groupby('time').aggregate({'value':'mean'}))

```

executed in 7ms, finished 10:45:49 2022-05-03

3.2 Stationarity Check

In [6]:

```

1 def stationarity_check(TS):
2
3     # Import adfuller
4     from statsmodels.tsa.stattools import adfuller
5
6     # Calculate rolling statistics
7     roll_mean = TS.rolling(window=8, center=False).mean()
8     roll_std = TS.rolling(window=8, center=False).std()
9
10    # Perform the Dickey Fuller test
11    dfoutput = adfuller(TS)
12
13    # Plot rolling statistics:
14    fig = plt.figure(figsize=(12,6))
15    orig = plt.plot(TS, color='blue',label='Original')
16    mean = plt.plot(roll_mean, color='red', label='Rolling Mean')
17    std = plt.plot(roll_std, color='black', label = 'Rolling Std')
18    plt.legend(loc='best')
19    plt.title('Rolling Mean & Standard Deviation')
20    plt.show(block=False)
21
22    # Print Dickey-Fuller test results
23    print('Results of Dickey-Fuller Test: \n')
24
25    dfoutput = pd.Series(dfoutput[0:4], index=['Test Statistic', 'p-value',
26                           '#Lags Used', 'Number of Observat
27    for key, value in dfoutput[4].items():
28        dfoutput['Critical Value (%s)'%key] = value
29    print(dfoutput)
30
31    return None

```

executed in 8ms, finished 10:45:50 2022-05-03

4 Choose Zip Codes by Percent Increase in Median Price

- Time Period: April 2012 to April 2018

4.1 Create Melted Dataframe of ZipCodes

In [7]:

```
1 # Create dataframe from zillow median home price data
2 df = pd.read_csv('zillow_data.csv')
3
4 # make copy of df
5 df1 = df.copy()
6
7 # Rename RegionName as ZipCode
8 df1 = df1.rename(columns={"RegionName": "ZipCode"})
9
10 # add column with percent change between 2017 and 2012 median value
11 df1['Percent Change'] = (df['2018-04'] / df['2012-04']) * 100
12
13 # create new dataframe with top 10 zip codes for Percent Change
14 df_percent = df1.sort_values('Percent Change', ascending = False)
15
16 # new df_percent with only top ten results
17 df_percent = df_percent.head(15)
18
19 # add mountain towns and mid-size cities
20 mount_mid = df1.loc[[9058,4502,4293,4442,6810,5764]]
21 df_percent = pd.concat([df_percent,mount_mid])
22
23 # Remove Raw Difference and Percent Change for melting
24 prospects = df_percent.drop(columns=['Percent Change', 'RegionID'])
25
26 # Remove Column with missing data
27 prospects = prospects.drop([1946])
```

executed in 967ms, finished 10:45:52 2022-05-03

In [8]:

1	prospects
executed in 70ms, finished 10:45:53 2022-05-03	

Out[8]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05
1477	94601	Oakland	CA	San Francisco	Alameda	1478	114600.0	114500.0
1239	94590	Vallejo	CA	Vallejo	Solano	1240	108200.0	107800.0
2627	33460	Lake Worth	FL	Miami-Fort Lauderdale	Palm Beach	2628	59800.0	59900.0
1681	94606	Oakland	CA	San Francisco	Alameda	1682	120400.0	120300.0
1853	94804	Richmond	CA	San Francisco	Contra Costa	1854	171300.0	170300.0
1960	89104	Las Vegas	NV	Las Vegas	Clark	1961	94800.0	94700.0
474	85008	Phoenix	AZ	Phoenix	Maricopa	475	61600.0	61900.0
4298	95824	Sacramento	CA	Sacramento	Sacramento	4299	73800.0	73400.0
2409	33404	Riviera Beach	FL	Miami-Fort Lauderdale	Palm Beach	2410	56800.0	56700.0
2661	33705	Saint Petersburg	FL	Tampa	Pinellas	2662	54600.0	54800.0
2934	89107	Las Vegas	NV	Las Vegas	Clark	2935	92300.0	92400.0
815	89115	Las Vegas	NV	Las Vegas	Clark	816	92500.0	92500.0
2775	95820	Sacramento	CA	Sacramento	Sacramento	2776	71800.0	71800.0
5830	48240	Redford	MI	Detroit	Wayne	5831	67800.0	68200.0
9058	84101	Salt Lake City	UT	Salt Lake City	Salt Lake	9059	85100.0	85400.0
4502	48009	Birmingham	MI	Detroit	Oakland	4503	242700.0	242600.0
4293	66104	Kansas City	KS	Kansas City	Wyandotte	4294	41300.0	41200.0
4442	84606	Provo	UT	Provo	Utah	4443	142700.0	141400.0
6810	87506	Santa Fe	NM	Santa Fe	Santa Fe	6811	227000.0	227400.0
5764	83703	Boise	ID	Boise City	Ada	5765	107500.0	107500.0

20 rows × 271 columns

In [9]:

```
1 # Check for missing data
2 print(prospects.isnull().sum().sum())
```

executed in 12ms, finished 10:45:53 2022-05-03

0

4.2 Print Decomposition Chart for each Zip Code

In [10]:

```
1 # Melt and decompose each zip from prospects df
2 for zipcode in prospects['ZipCode']:
3     zips = prospects.loc[prospects['ZipCode'] == zipcode]
4     m1 = melt_data(zips)
5     m1[zips.iloc[0,1],zips.iloc[0,2],str(zipcode)] = m1['value']
6     m1 = m1.drop('value', axis=1)
7
8     #Apply seasonal_decompose()
9     decomposition = seasonal_decompose(np.log(m1))
10
11     # Gather the trend, seasonality, and residuals
12     trend = decomposition.trend
13     seasonal = decomposition.seasonal
14     residual = decomposition.resid
15
16     # Plot gathered statistics
17     plt.figure(figsize=(12,8))
18     plt.subplot(411)
19     plt.title(str(m1.columns.values))
20     plt.plot(np.log(m1), label='Original', color='blue')
21     plt.legend(loc='best')
22     plt.subplot(412)
23     plt.plot(trend, label='Trend', color='blue')
24     plt.legend(loc='best')
25     plt.subplot(413)
26     plt.plot(seasonal,label='Seasonality', color='blue')
27     plt.legend(loc='best')
28     plt.subplot(414)
29     plt.plot(residual, label='Residuals', color='blue')
30     plt.legend(loc='best')
31     plt.tight_layout()
```

executed in 17.0s, finished 10:46:10 2022-05-03

4.3 Seasonality Trends Fascinating

- Seasonality present in every zip code, but how the trend occurs is different for each zip code.
- Certainly justifies using SARIMA instead of ARIMA due to seasonal pattern.

5 Adding Breadth of Zip Codes

- Mountain Towns and Mid-Size Cities

5.1 Mountain Towns

- Boise, ID
- Santa Fe, NM
- Boulder, CO
- Provo, UT
- Bend, OR

5.1.1 Create Mountain Towns df

In [26]:

```

1 mountain_list = []
2 mountain_list.append(['Boise', 'Boulder', 'Provo', 'Bend'])
3
4 mountain_towns = pd.DataFrame()
5
6 mt = pd.DataFrame(mountain_list)
7 mt = mt.T
8
9 for i in mt[0]:
10     # Create df from new city
11     ok = df1[df1['City'] == str(i)]
12
13     # add column with percent change between 2017 and 2012 median value
14     ok['Percent Change'] = (ok['2018-04'] / ok['2012-04']) * 100
15
16     # create new dataframe with top 10 zip codes for Percent Change
17     ok = ok.sort_values('Percent Change', ascending = False)
18
19     # new df_percent with only top ten results
20     ok = ok.head(2)
21
22     # concatenate to mountain towns df
23     mountain_towns = pd.concat([ok, mountain_towns])
24
25 # Create df from new city
26 santa_fe = df1[df1['City'] == 'Santa Fe']
27
28 # Choose top zip code and my zip code from Santa Fe, NM
29 santa_fe = santa_fe[(santa_fe['ZipCode'] == 87505) | (santa_fe['ZipCode'] == 8
30
31 # concatenate to mountain towns df
32 mountain_towns = pd.concat([santa_fe, mountain_towns])
33
34 mountain_towns = mountain_towns.drop(['RegionID', 'Percent Change'], axis=1)
35
36 mountain_towns

```

executed in 110ms, finished 11:50:25 2022-05-03

Out[26]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	1996-
1776	87505	Santa Fe	NM	Santa Fe	Santa Fe	1777	188200.0	188400.0	188500
6810	87506	Santa Fe	NM	Santa Fe	Santa Fe	6811	227000.0	227400.0	227900
1446	97701	Bend	OR	Bend	Deschutes	1447	136300.0	136100.0	135900
706	97702	Bend	OR	Bend	Deschutes	707	160300.0	160300.0	160300

ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	1996-
4081	84601	Provo	UT	Provo	Utah	4082	142700.0	140500.0
4442	84606	Provo	UT	Provo	Utah	4443	142700.0	141400.0
6115	80305	Boulder	CO	Boulder	Boulder	6116	187900.0	187700.0
3363	80303	Boulder	CO	Boulder	Boulder	3364	231100.0	230300.0
5764	83703	Boise	ID	Boise City	Ada	5765	107500.0	107500.0
2962	83705	Boise	ID	Boise City	Ada	2963	83000.0	83000.0

10 rows × 271 columns

In [27]:

```

1 # Create list for final predicted value
2 pred_list_2 = []
3
4 # Construct model from entirety of data - no train/test
5 # iterate over prospects to extract optimal SARIMA values,
6 # instantiate SARIMAX model, make predictions
7 # and produce forecast with confidence intervals
8 for zipcode in mountain_towns['ZipCode']:
9     zips = mountain_towns.loc[mountain_towns['ZipCode'] == zipcode]
10    m1 = melt_data(zips)
11    m1[str(zipcode)] = m1['value']
12    m1 = m1.drop('value', axis=1)
13
14    # do NOT create train and test sets
15    #train = m1.iloc[:238]
16    #test = m1.iloc[238:]
17
18    # Run a grid with pdq and seasonal pdq parameters calculated above and ge
19    ans = []
20    outputs = []
21    comb_list = []
22    combs_list = []
23
24    for comb in pdq:
25        for combs in pdqs:
26            try:
27                mod = sm.tsa.statespace.SARIMAX(m1,
28                                                order=comb,
29                                                seasonal_order=combs,
30                                                enforce_stationarity=False,
31                                                enforce_invertibility=False)
32
33                output = mod.fit()
34                ans.append([comb, combs, output.aic])
35                outputs.append([output.aic])
36                comb_list.append([comb])
37                combs_list.append([combs])
38                AIC_df = pd.DataFrame(list(outputs, comb_list, combs_list))
39                #print('ARIMA {} x {}12 : AIC Calculated ={}'.format(comb, co
40            except:
41                continue
42
43    # Print the lowest AIC and it's parameters for SARIMA
44    AIC_df[0] = AIC_df[0].str.get(0)
45    best_AIC = AIC_df.copy()
46    lowest_AIC = min(best_AIC[0])
47    lowest_AIC_row = AIC_df.loc[AIC_df[0] == lowest_AIC]
48    #print(lowest_AIC_row)
49
50    # call value for comb and combs for model input
51    lowest_AIC_row[1] = lowest_AIC_row[1].str.get(0)
52    lowest_AIC_row[2] = lowest_AIC_row[2].str.get(0)

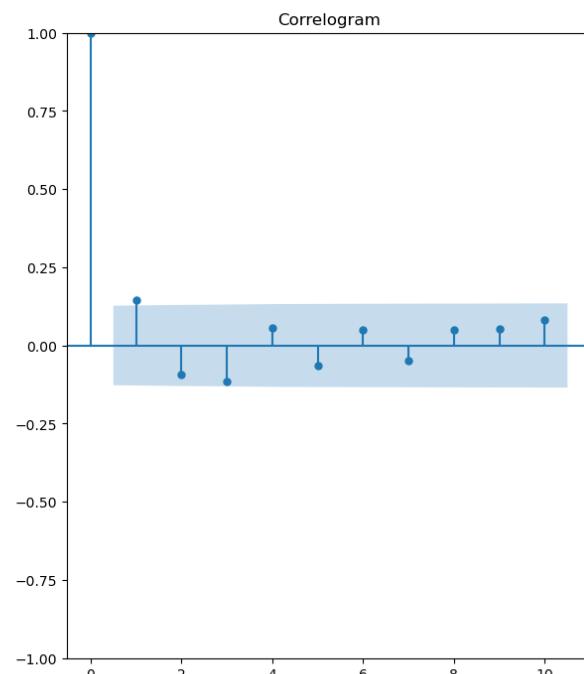
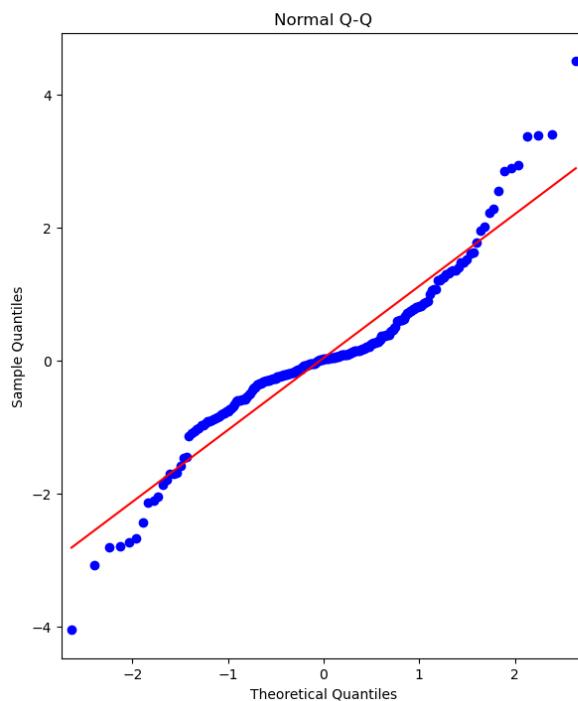
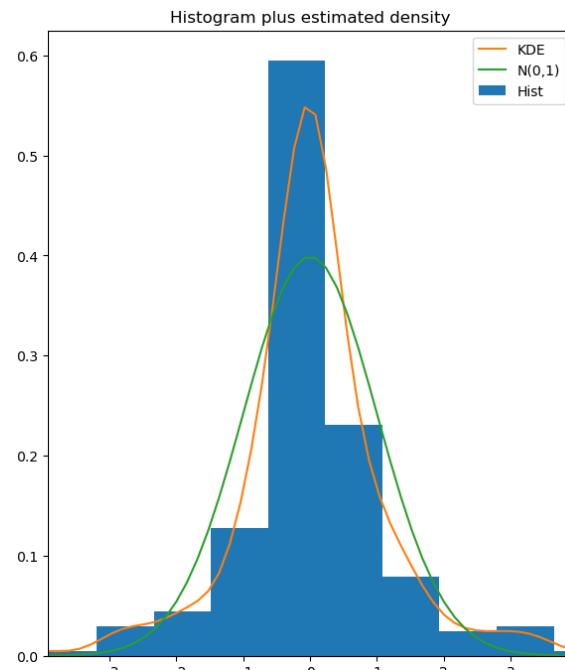
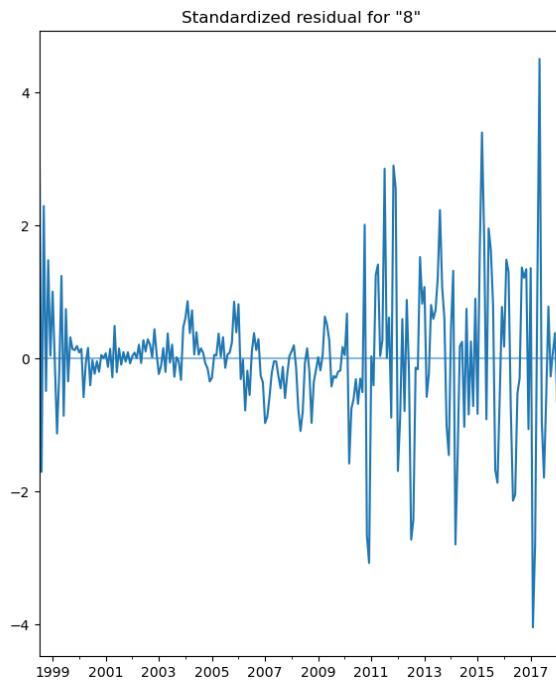
```

```
53     arima = lowest_AIC_row.iat[0,1]
54     s = lowest_AIC_row.iat[0,2]
55     print(str(m1.columns.values))
56     print(arima, s)
57
58     # Plug the optimal parameter values into a new SARIMAX model
59     ARIMA_MODEL = sm.tsa.statespace.SARIMAX(m1,
60                                              order=arima,
61                                              seasonal_order=s,
62                                              enforce_stationarity=False,
63                                              enforce_invertibility=False)
64
65     # Fit the model and print results
66     output = ARIMA_MODEL.fit()
67
68     print(output.summary().tables[1])
69
70     # Call plot_diagnostics() on the results calculated above
71     output.plot_diagnostics(figsize=(15, 18))
72     plt.show()
73
74     # Get actual results in dataframe
75     price_truth = m1['2018-04-01':]
76
77     # Get forecast 500 steps ahead in future
78     prediction = output.get_forecast(steps=60)
79
80     # Create predictions df
81     pred = prediction.predicted_mean
82
83     # Add prediction for 5 year forecasted value to list
84     pred_list_2.append(pred['2023-04-01'])
85
86     # Get confidence intervals of forecasts
87     pred_conf = prediction.conf_int()
88
89     # Only use data from April 2010 to improve visual of forecast
90     plot_vals = m1.loc['2010-04-01':]
91
92     # Plot future predictions with confidence intervals
93     ax = plot_vals.plot(label='observed', figsize=(20, 15), color='black')
94     prediction.predicted_mean.plot(ax=ax, label='Forecast', color='green')
95     ax.fill_between(pred_conf.index,
96                      pred_conf.iloc[:, 0],
97                      pred_conf.iloc[:, 1], color='k', alpha=0.25)
98     ax.set_xlabel('Date').set_fontsize(30)
99     ax.set_ylabel('Median House Price ($)').set_fontsize(30)
100    ax.set_title('Five Year Median House Price Forecast').set_fontsize(40)
101
102   for item in (ax.get_xticklabels() + ax.get_yticklabels()):
103       item.set_fontsize(20)
104
105   right_side = ax.spines["right"]
106   right_side.set_visible(False)
```

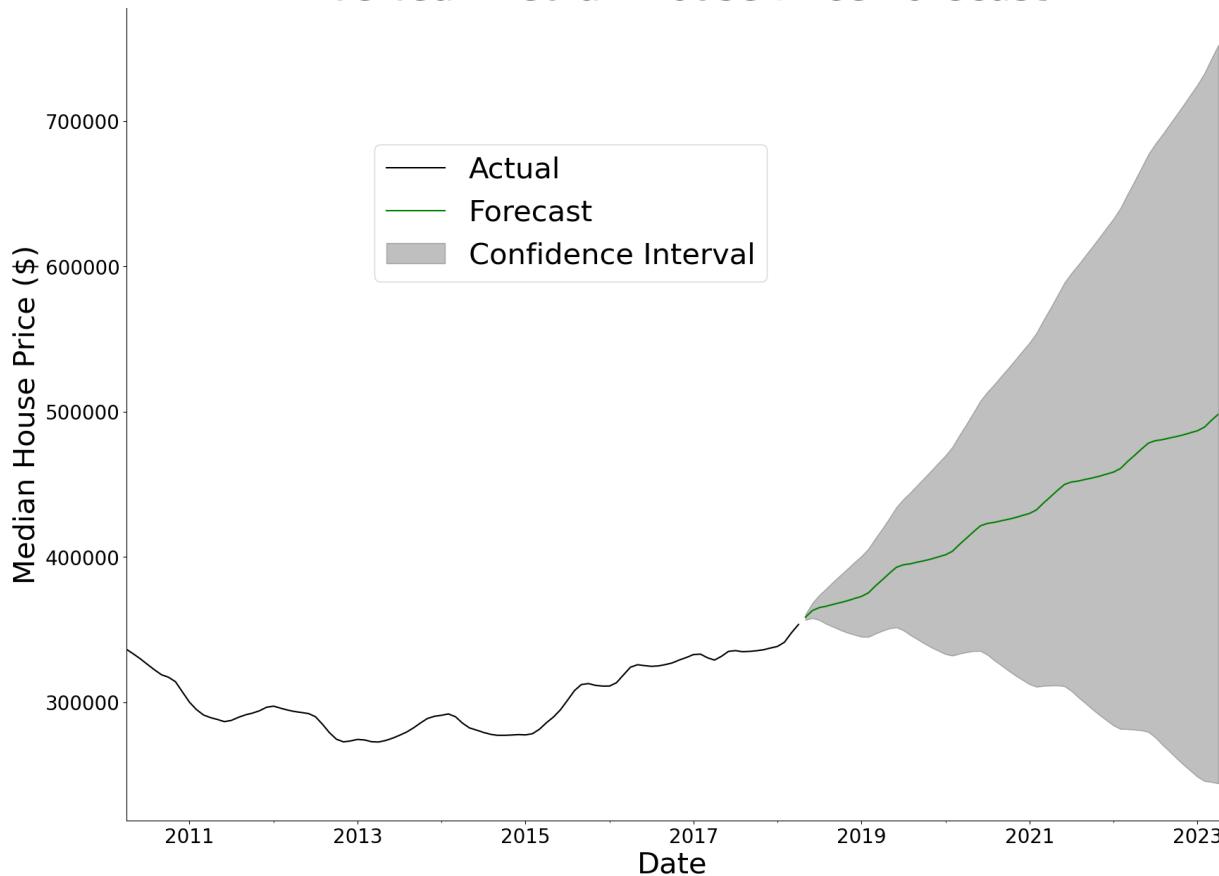
```
107
108     top_side = ax.spines["top"]
109     top_side.set_visible(False)
110
111     mylabels = 'Actual', 'Forecast', 'Confidence Interval'
112     plt.legend(fontsize=30, labels=mylabels, bbox_to_anchor=(0.6,0.85))
113     plt.show()
114     print()
115     print()
116
117
```

executed in 4m 43s, finished 11:55:23 2022-05-03

```
['87505']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
          coef      std err           z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.7393      0.033      22.200      0.000      0.674
0.805
ma.L1      0.7373      0.028      26.530      0.000      0.683
0.792
ma.S.L12   -0.2625      0.027     -9.647      0.000     -0.316
-0.209
sigma2    1.005e+06  6.32e+04      15.899      0.000    8.81e+05
1.13e+06
=====
```

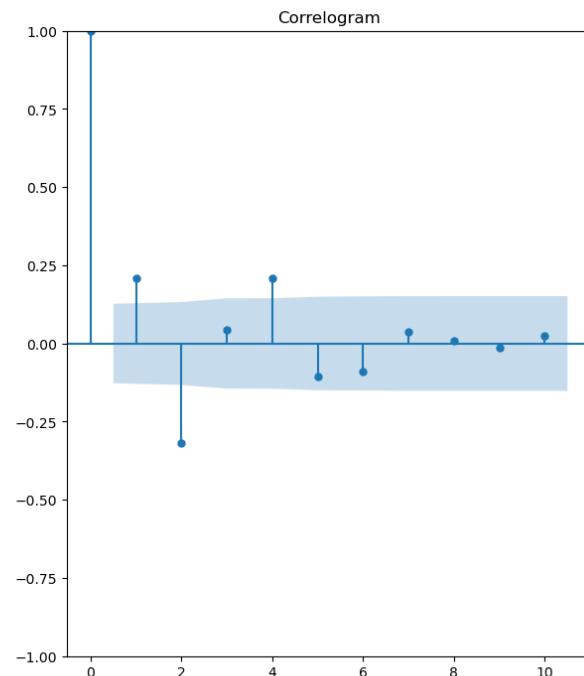
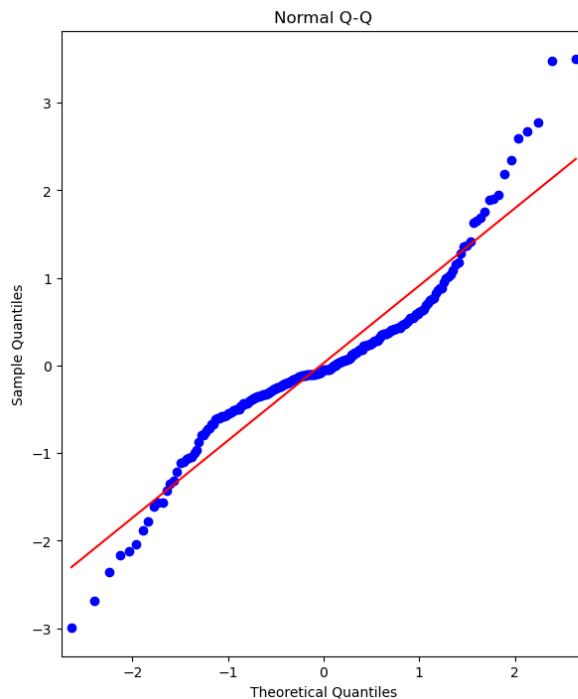
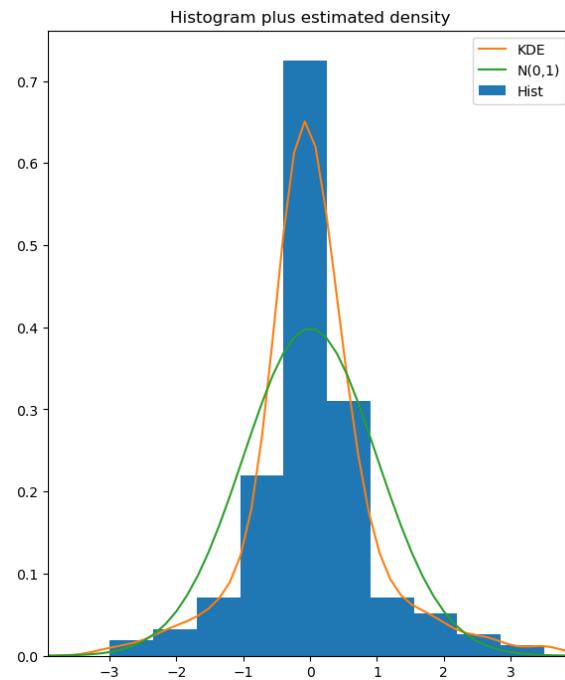
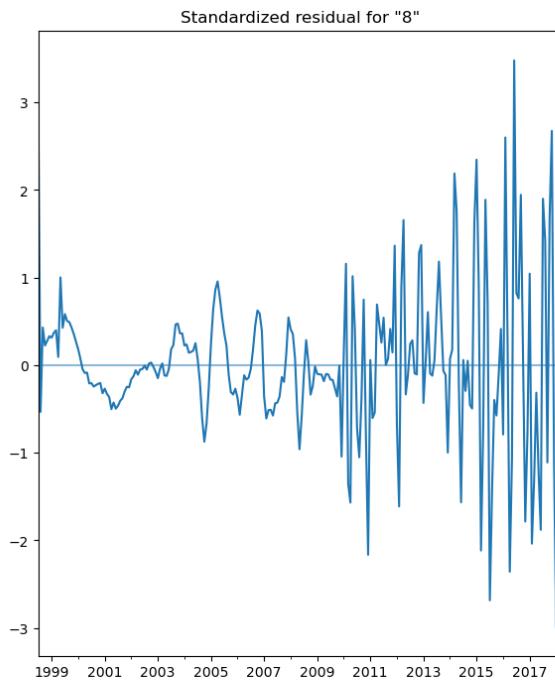


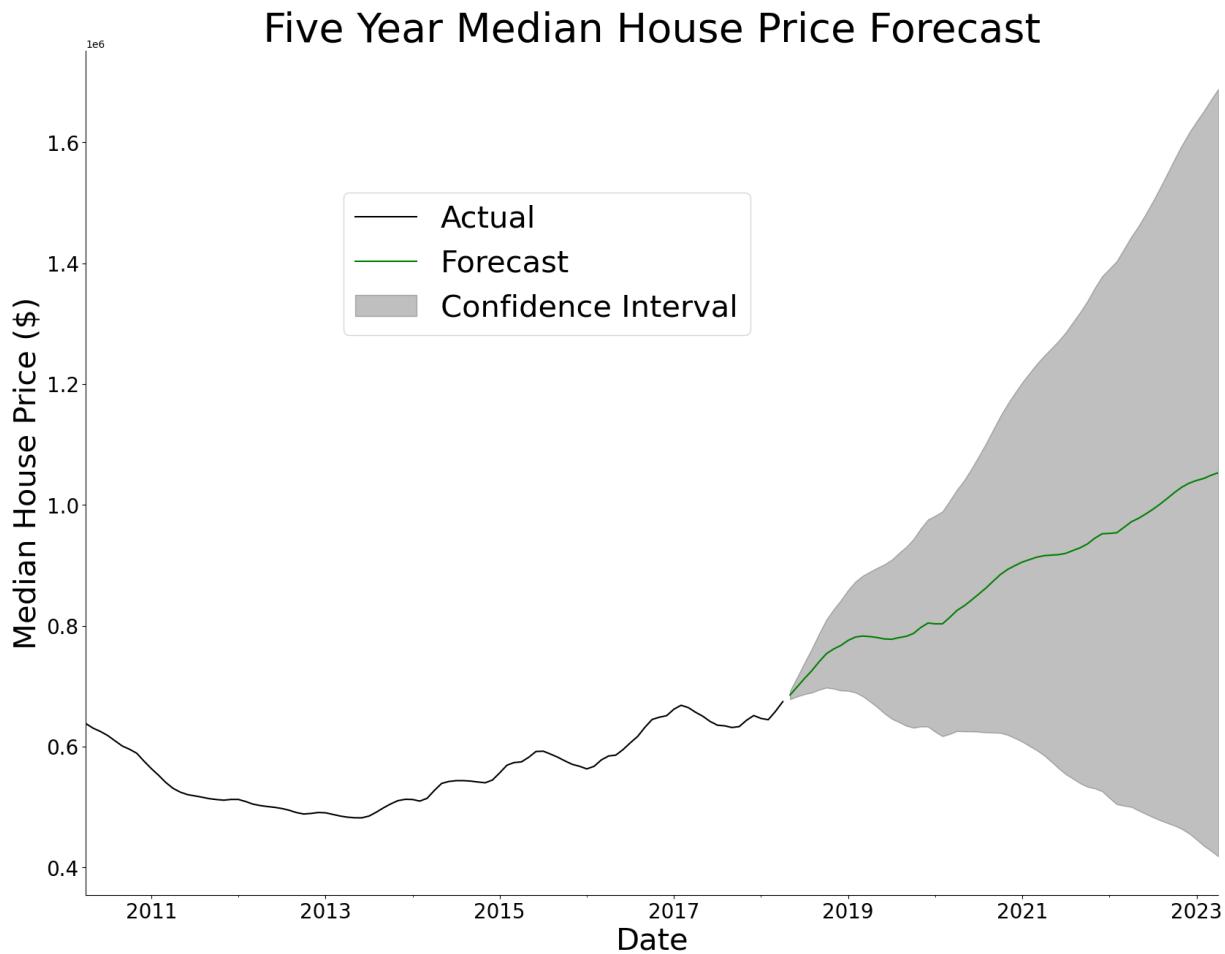
Five Year Median House Price Forecast



```

['87506']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
-----
ar.L1      0.7830      0.044      17.772      0.000      0.697
0.869
ma.L1      0.2943      0.021      14.125      0.000      0.253
0.335
ar.S.L12   -0.7339      0.050     -14.795      0.000     -0.831
-0.637
ma.S.L12   -0.0188      0.032      -0.590      0.555     -0.081
0.044
sigma2     1.292e+07    9.44e-10    1.37e+16    0.000     1.29e+07
1.29e+07
=====
=====
```



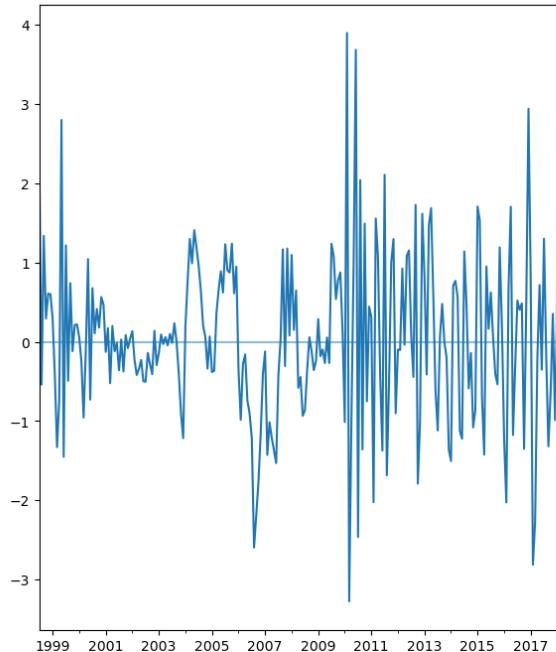


```
['97701']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
      coef      std err          z      P>|z|      [ 0.025
0.975]
```

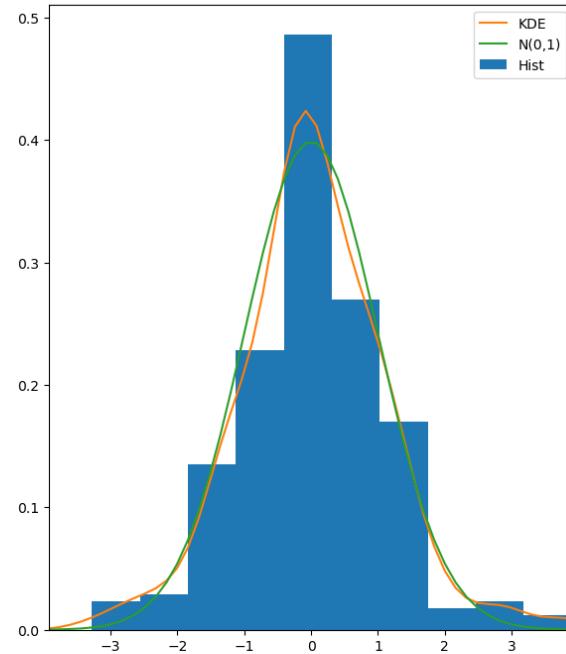
ar.L1	0.9299	0.022	43.130	0.000	0.888
0.972					
ma.L1	0.5517	0.034	16.282	0.000	0.485
0.618					
ma.S.L12	-0.4288	0.030	-14.182	0.000	-0.488
-0.370					
sigma2	4.249e+05	2.95e+04	14.381	0.000	3.67e+05
4.83e+05					

=====

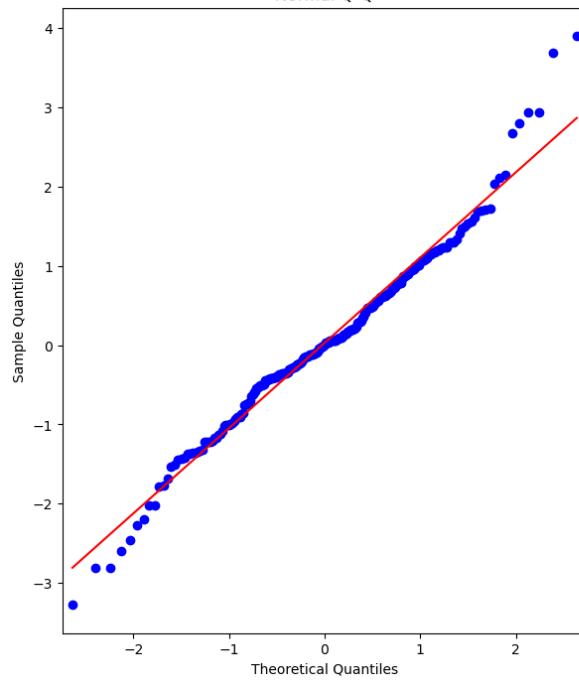
Standardized residual for "9"



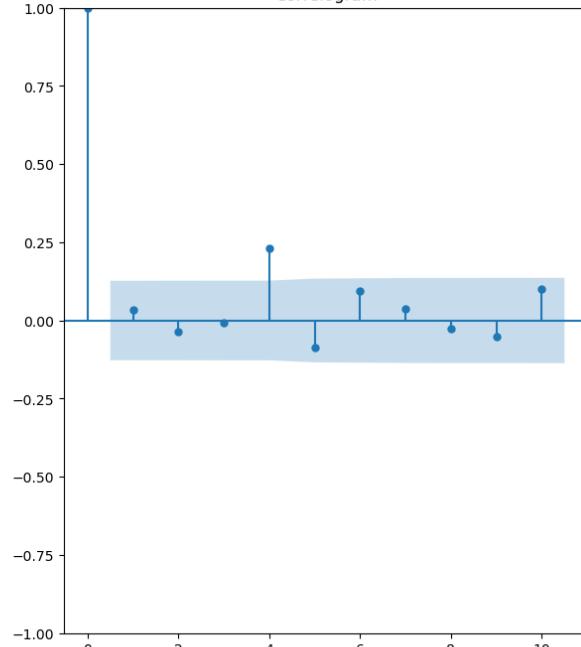
Histogram plus estimated density



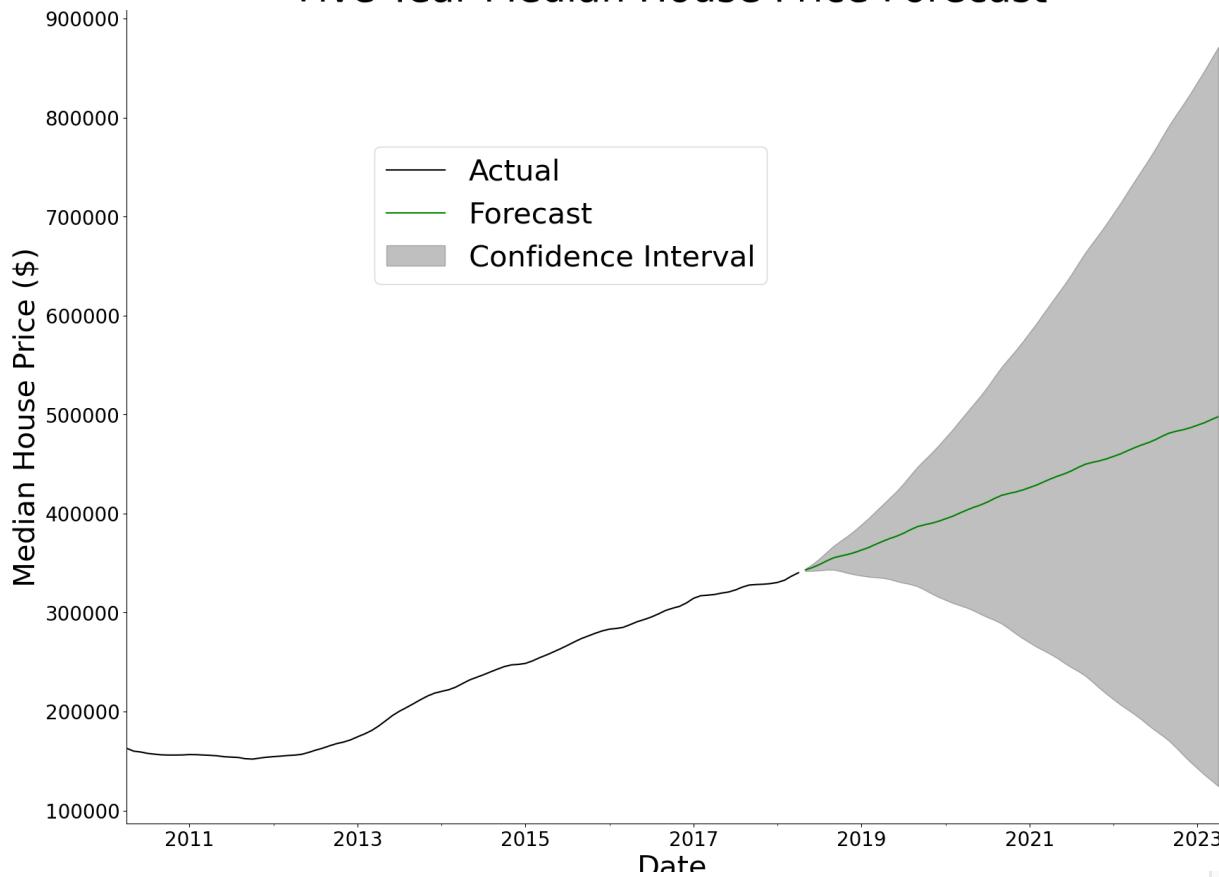
Normal Q-Q



Correlogram



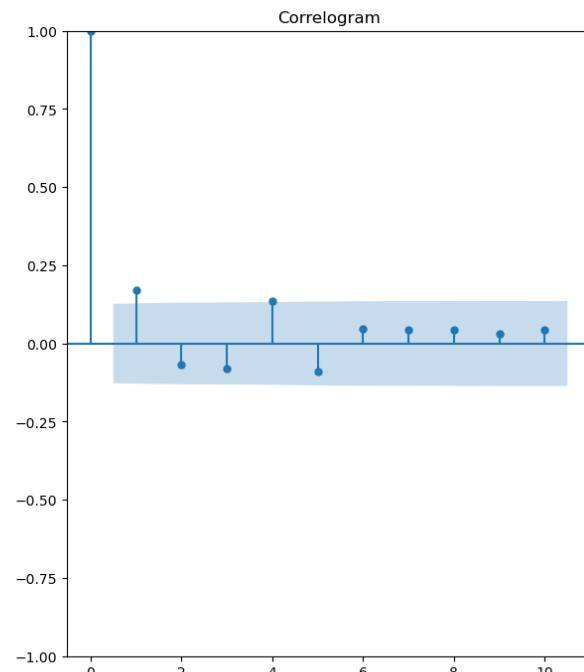
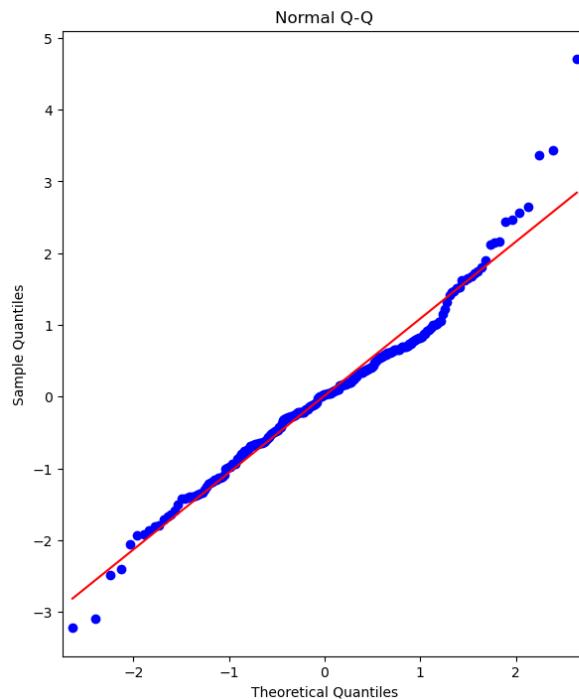
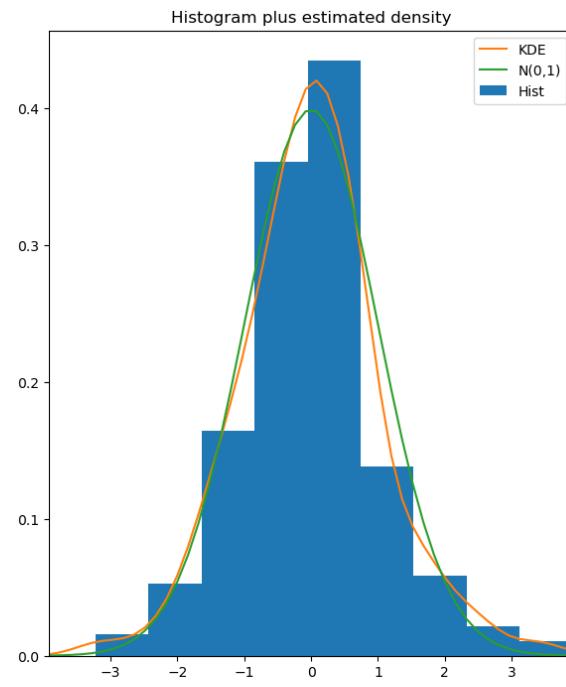
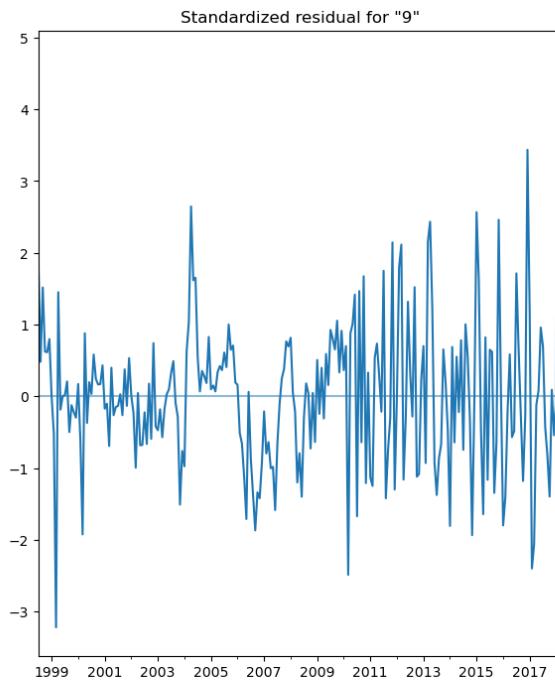
Five Year Median House Price Forecast

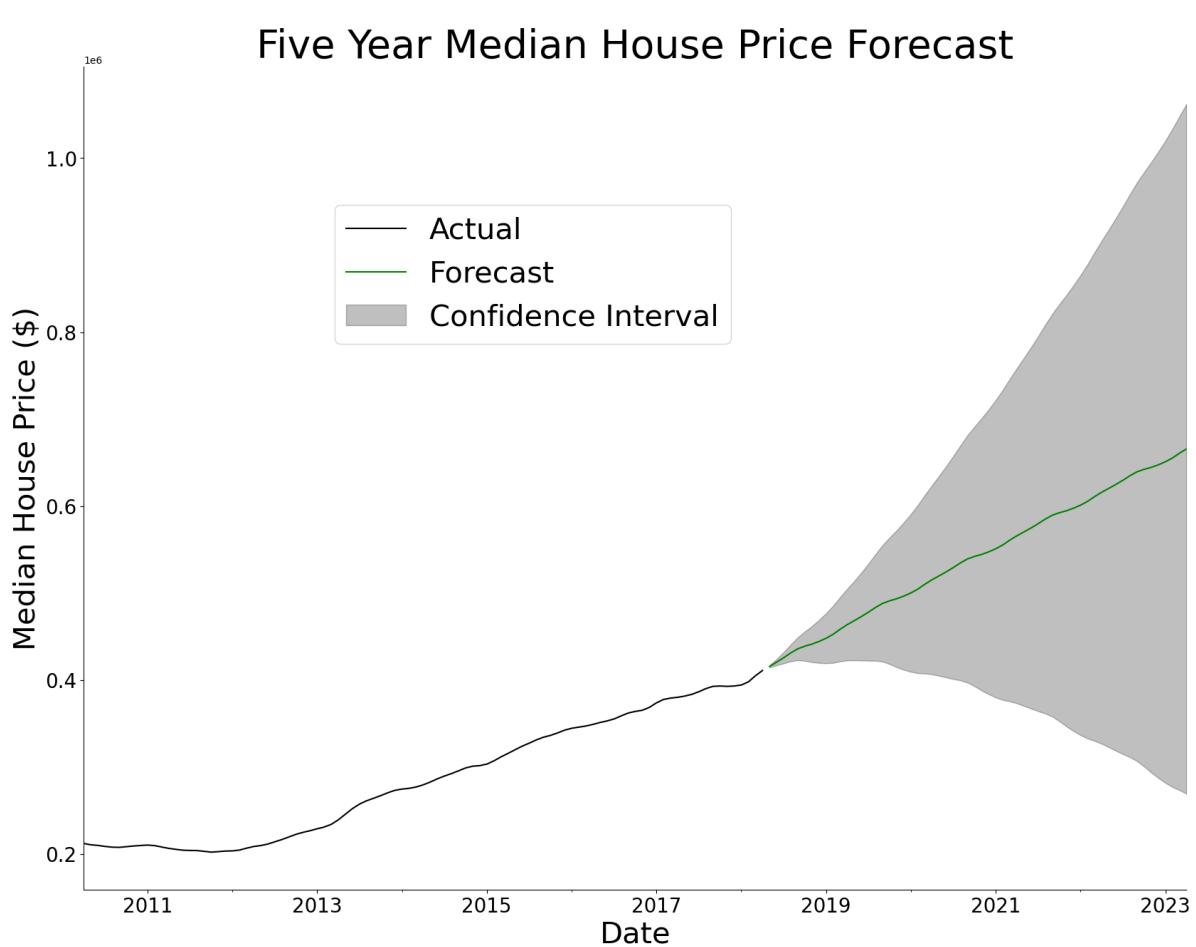


```

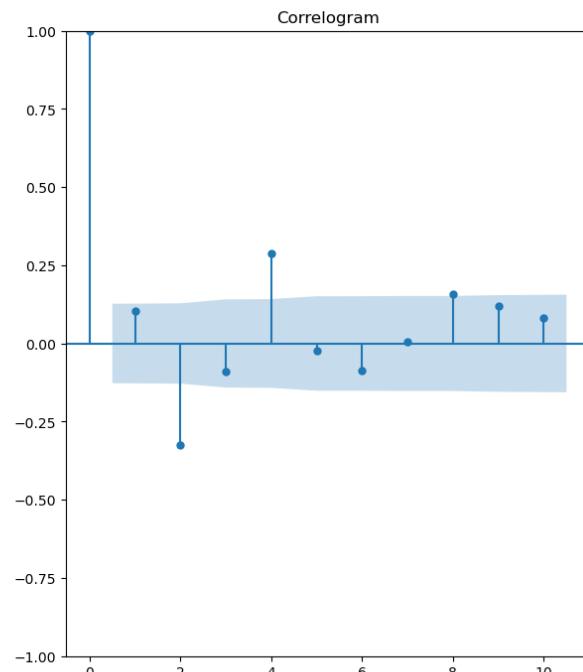
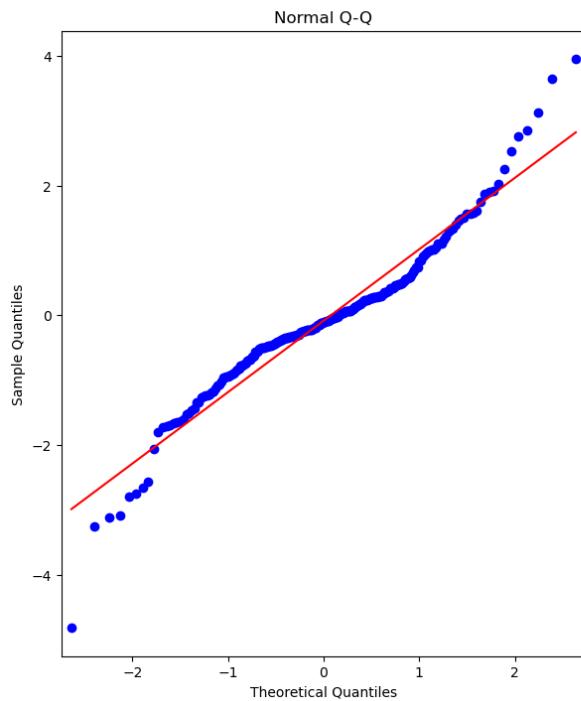
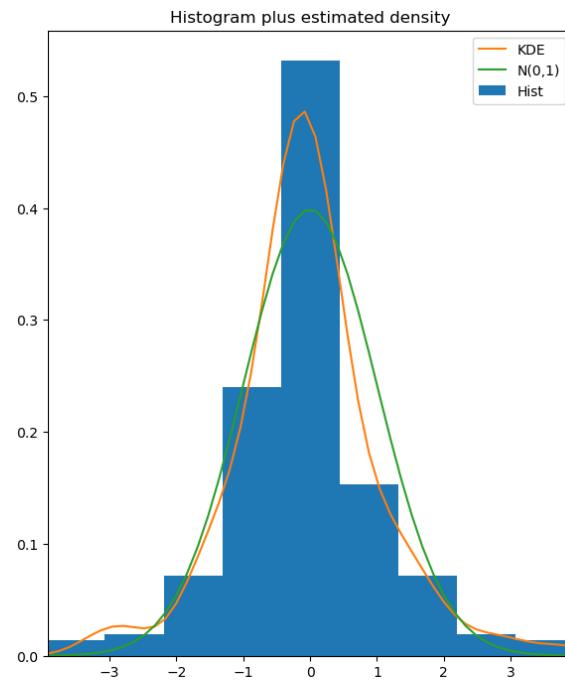
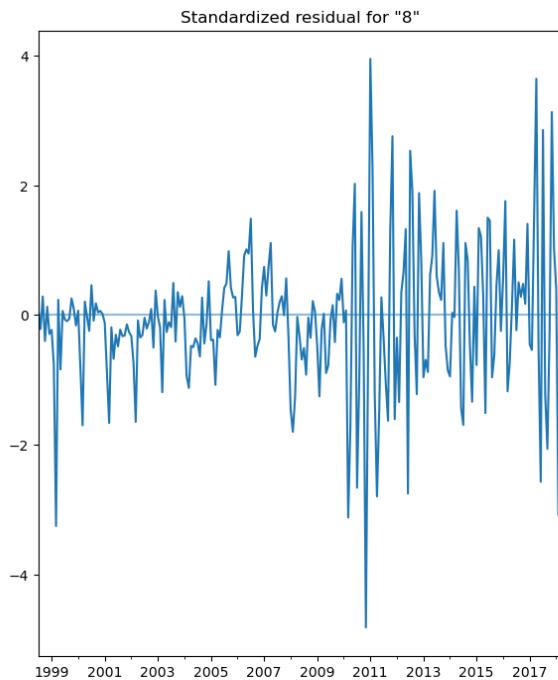
['97702']
(1, 1, 1) (1, 1, 1, 12)
=====
=====

      coef    std err        z   P>|z|      [ 0.025
0.975]
-----
-----
ar.L1      0.9307    0.024    39.489    0.000    0.885
0.977
ma.L1      0.7650    0.039    19.513    0.000    0.688
0.842
ar.S.L12    0.0607    0.011     5.459    0.000    0.039
0.082
ma.S.L12   -0.5339    0.029   -18.252    0.000   -0.591
-0.477
sigma2     4.122e+05  2.93e+04   14.086    0.000   3.55e+05
4.7e+05
=====
=====
```

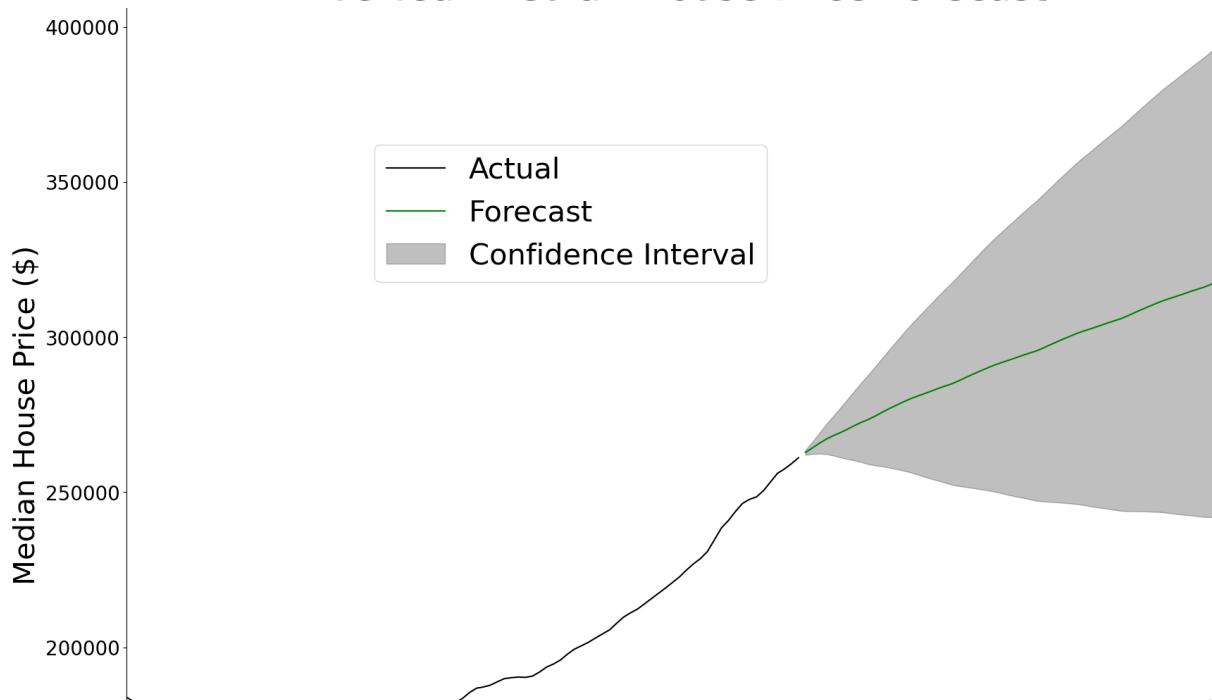




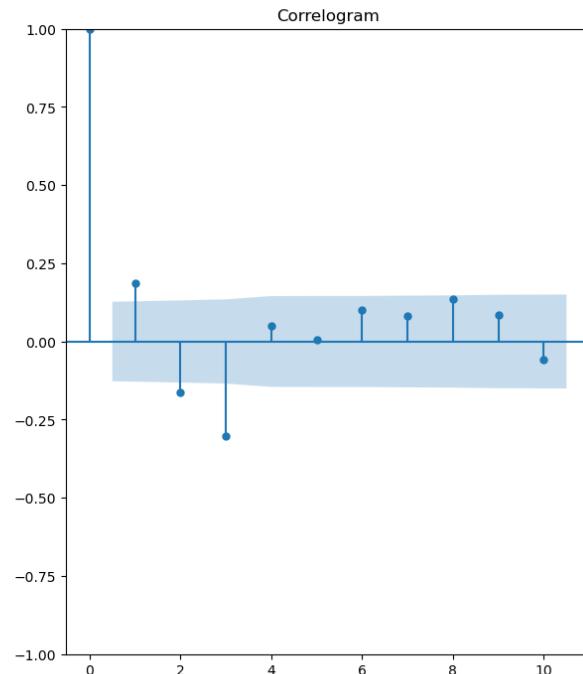
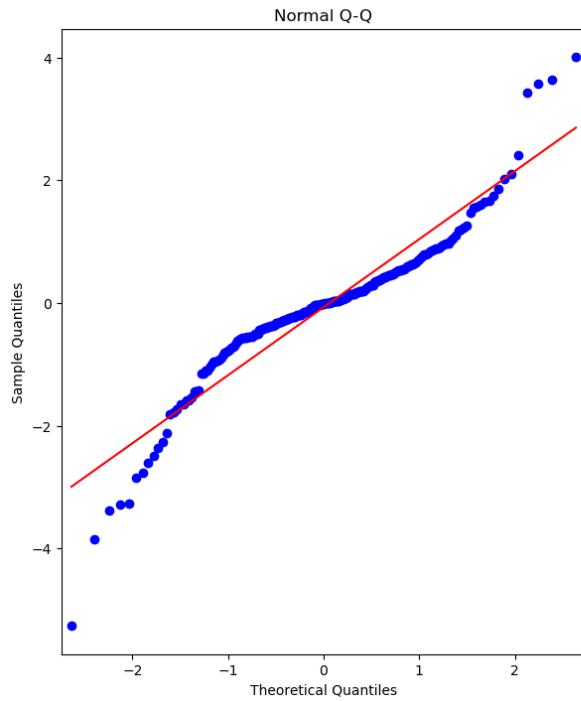
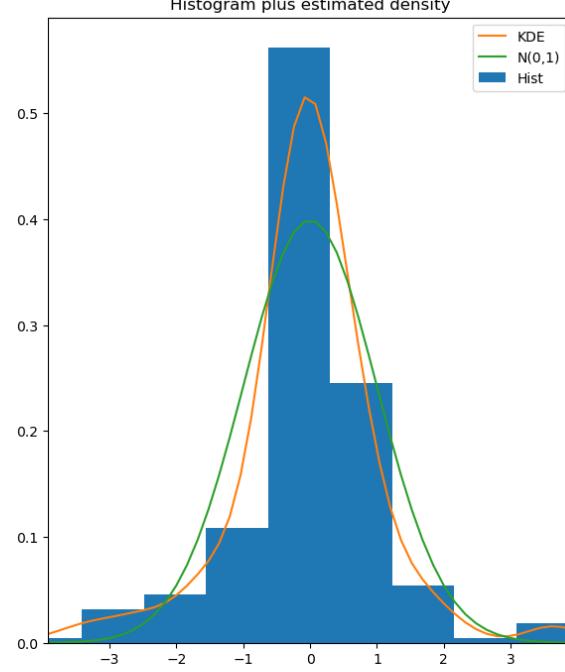
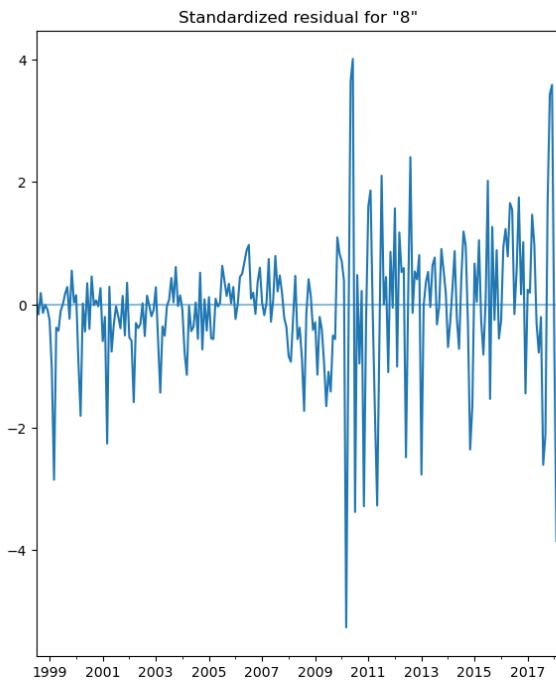
```
[ '84601' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
          coef      std err       z     P>|z|      [ 0.025
0.975]
-----
ar.L1      0.8761      0.024    36.982      0.000      0.830
0.923
ma.L1      0.4131      0.036    11.379      0.000      0.342
0.484
ar.S.L12   0.0880      0.014     6.510      0.000      0.061
0.114
ma.S.L12   -0.9493     0.085   -11.121      0.000     -1.117
-0.782
sigma2     1.636e+05   1.29e+04    12.725      0.000     1.38e+05
1.89e+05
=====
```



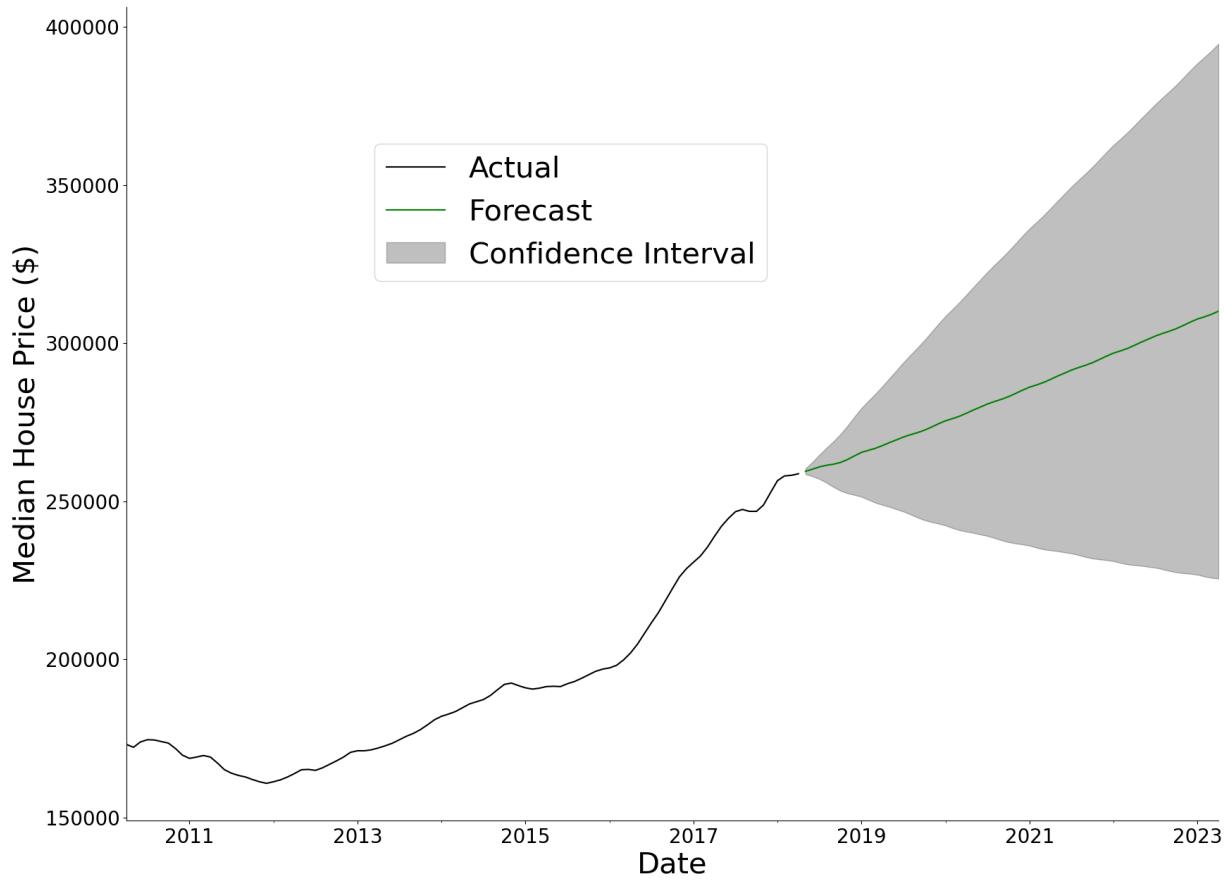
Five Year Median House Price Forecast



```
['84606']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
0.975]
-----
ar.L1      0.8480    0.028    30.792    0.000    0.794
0.902
ma.L1      0.4739    0.041    11.471    0.000    0.393
0.555
ar.S.L12   0.0933    0.018    5.135     0.000    0.058
0.129
ma.S.L12   -0.8693   0.048   -18.100    0.000   -0.963
-0.775
sigma2     2.099e+05  1.09e+04  19.290    0.000  1.89e+05
2.31e+05
=====
```



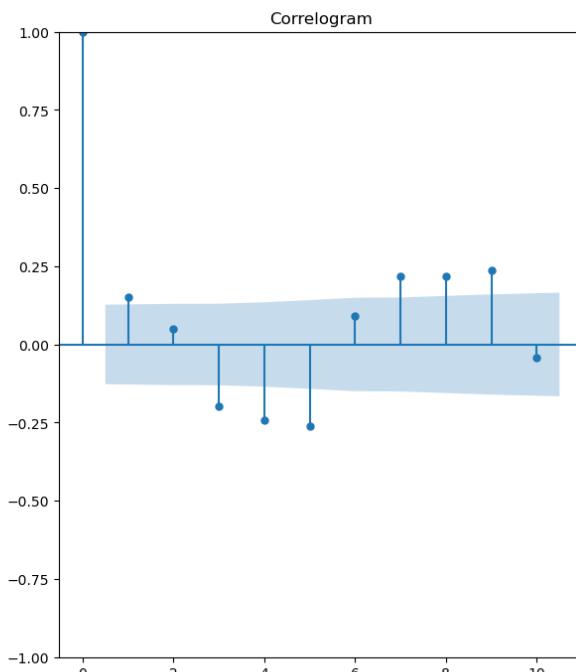
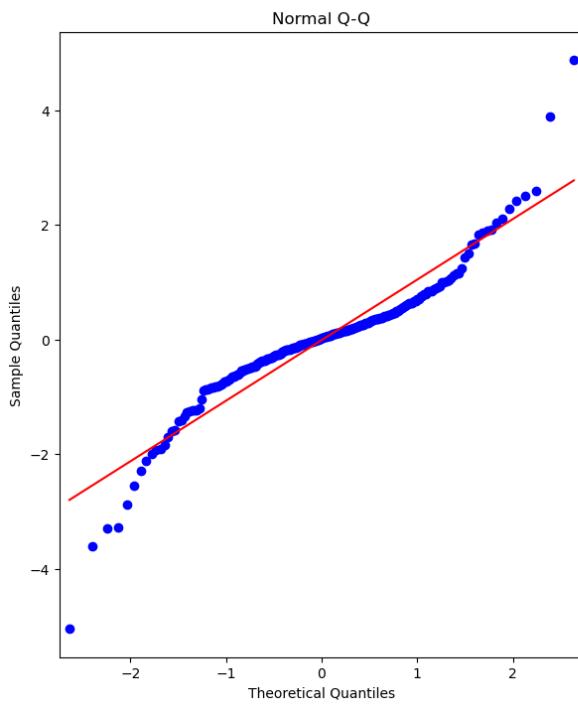
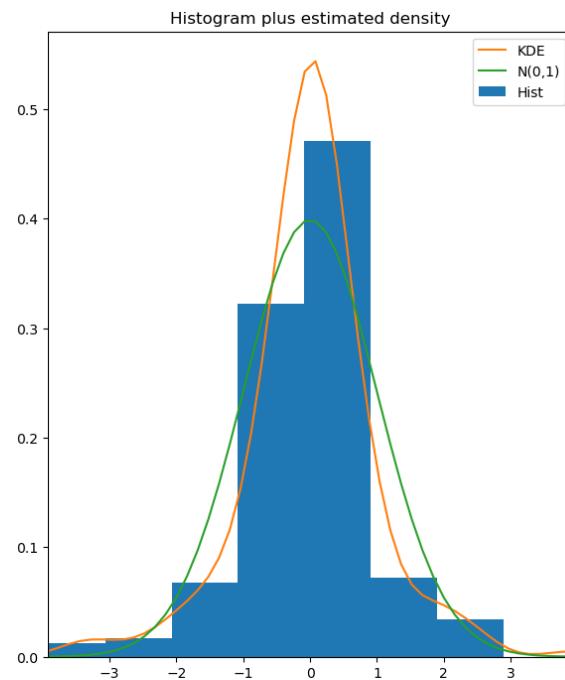
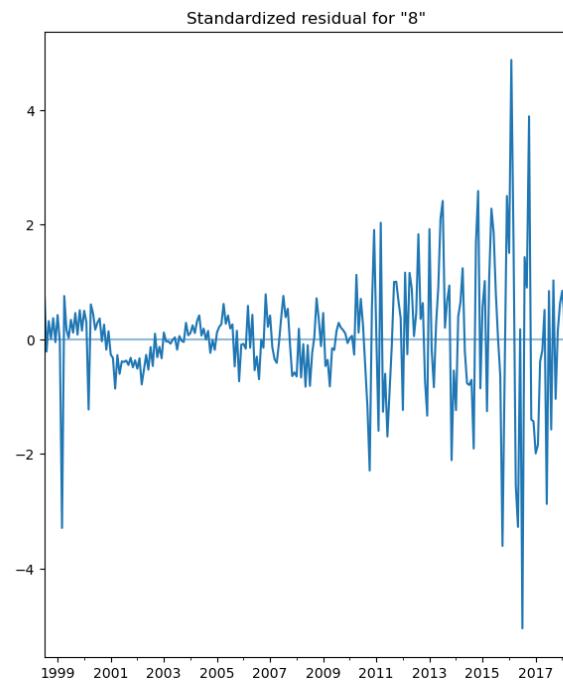
Five Year Median House Price Forecast



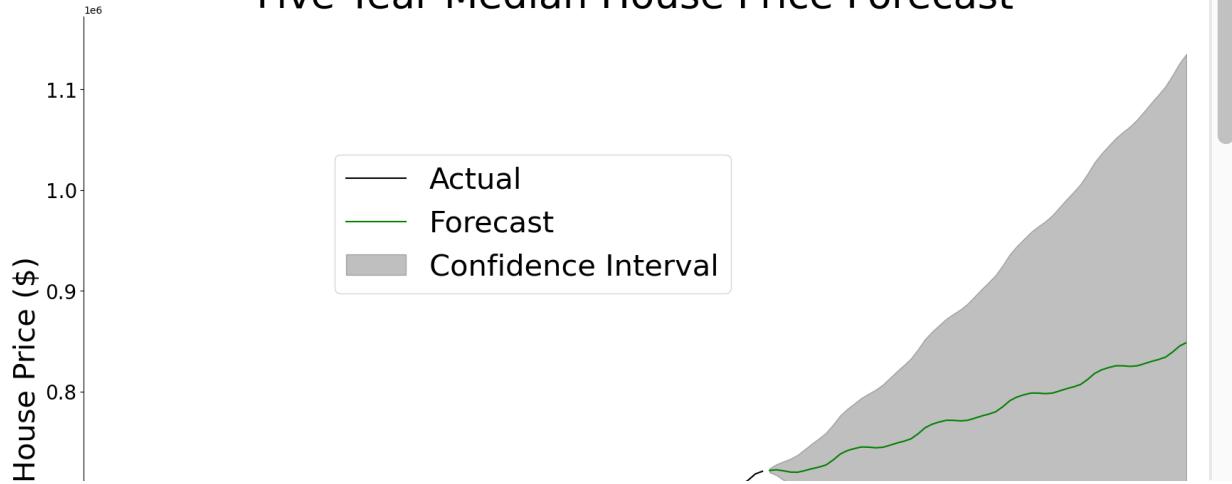
```
['80305']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
          coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
ar.L1      0.7820      0.029      26.897      0.000      0.725
0.839
ma.L1      0.7818      0.028      27.511      0.000      0.726
0.838
ar.S.L12   0.0834      0.023      3.676      0.000      0.039
```

```

0.128
ma.S.L12      -0.4268     0.021    -20.510     0.000    -0.468
-0.386
sigma2        1.052e+06   5.02e+04   20.971     0.000    9.54e+05
1.15e+06
=====
=====
```



Five Year Median House Price Forecast



```
[ '80303' ]
```

```
(1, 1, 1) (1, 1, 1, 12)
```

```
=====
```

```
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

```
=====
```

ar.L1	0.8025	0.035	22.671	0.000	0.733
-------	--------	-------	--------	-------	-------

```
0.872
```

ma.L1	0.7827	0.046	16.927	0.000	0.692
-------	--------	-------	--------	-------	-------

```
0.873
```

ar.S.L12	0.0784	0.036	2.178	0.029	0.008
----------	--------	-------	-------	-------	-------

```
0.149
```

ma.S.L12	-0.2799	0.042	-6.647	0.000	-0.362
----------	---------	-------	--------	-------	--------

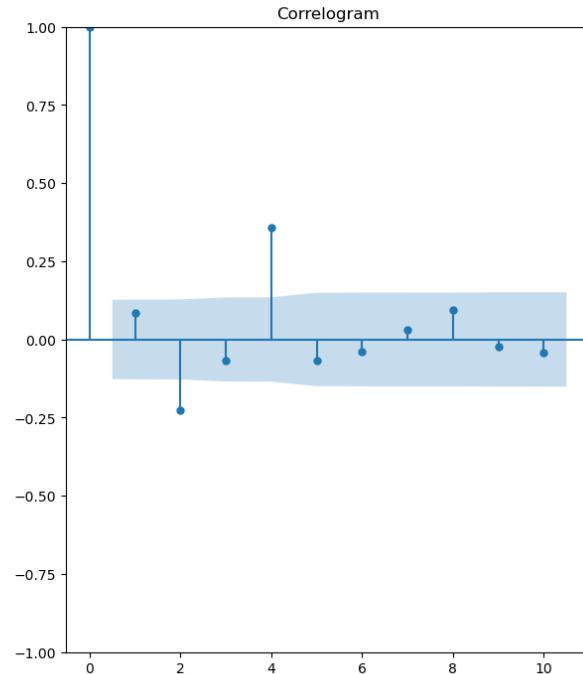
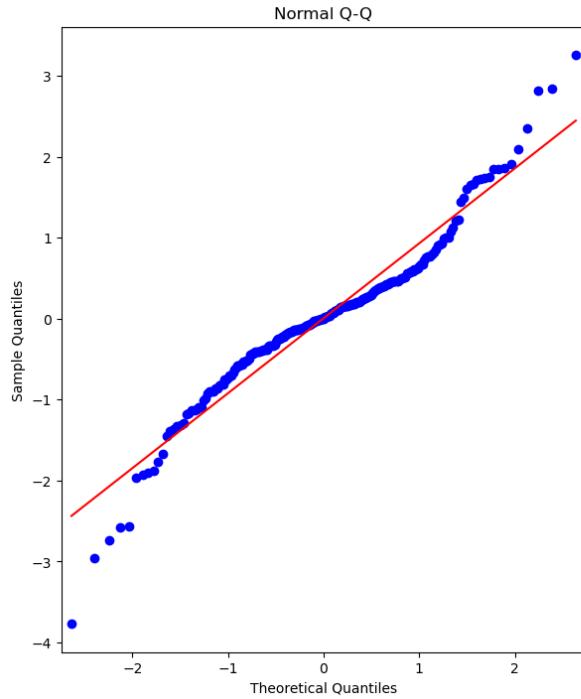
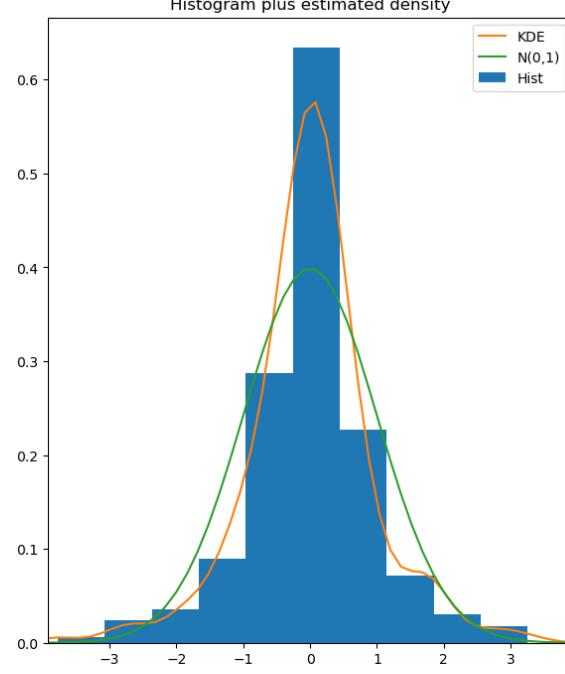
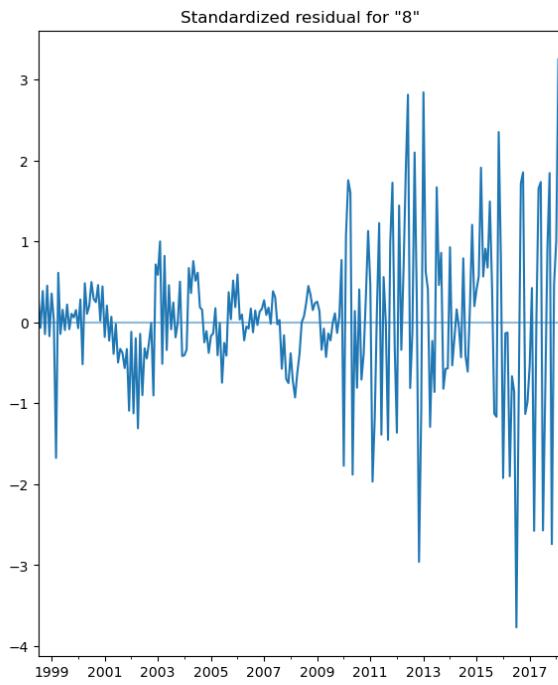
```
-0.197
```

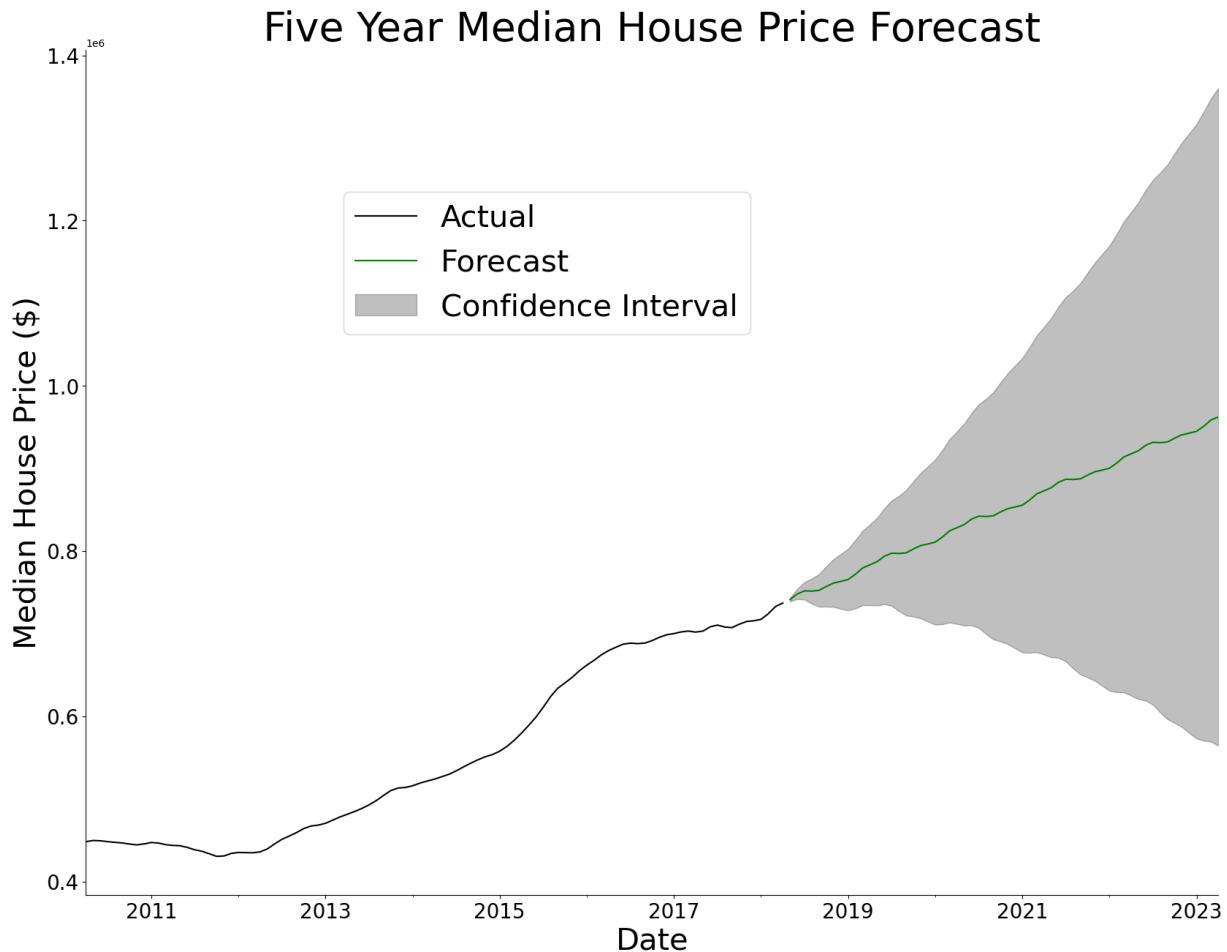
sigma2	1.304e+06	9.64e+04	13.530	0.000	1.11e+06
--------	-----------	----------	--------	-------	----------

```
1.49e+06
```

```
=====
```

```
=====
```

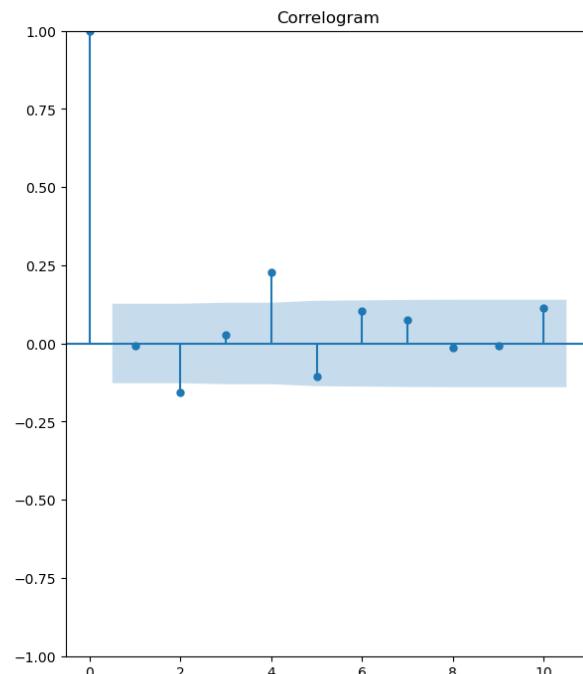
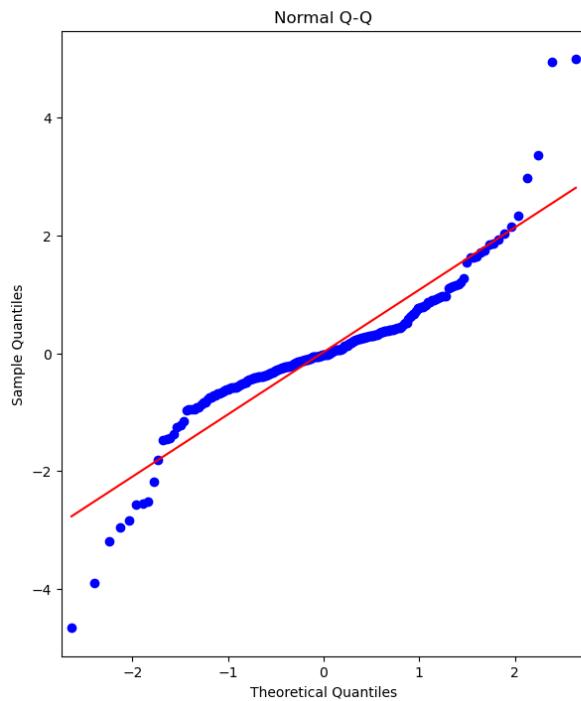
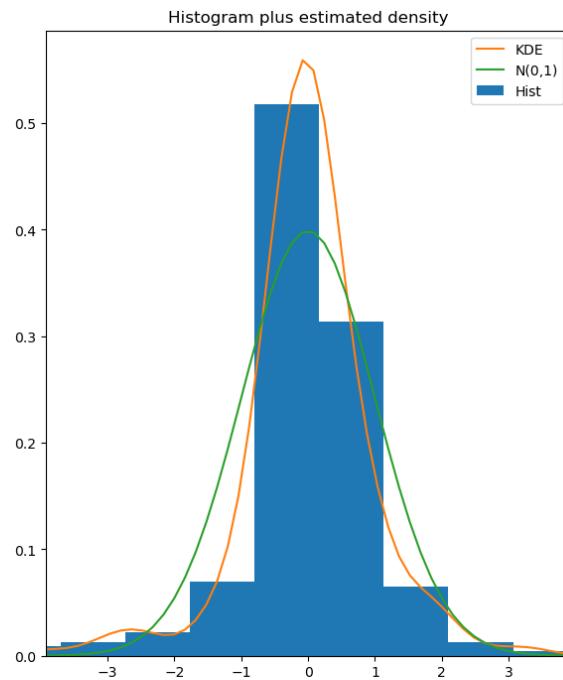
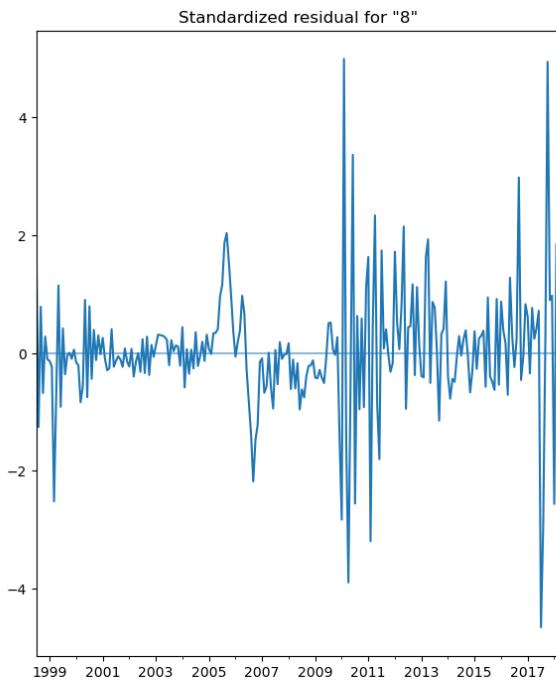


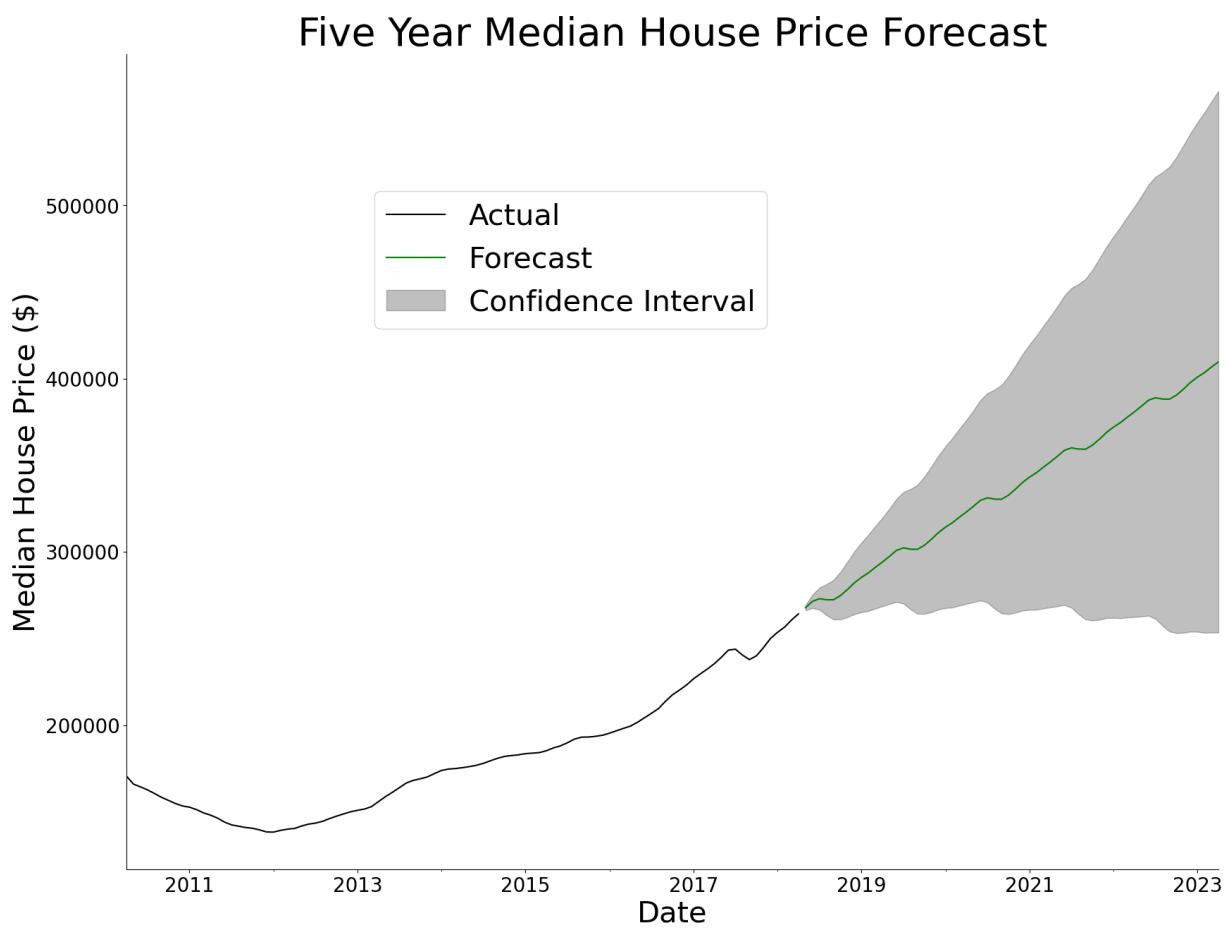


```
['83703']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

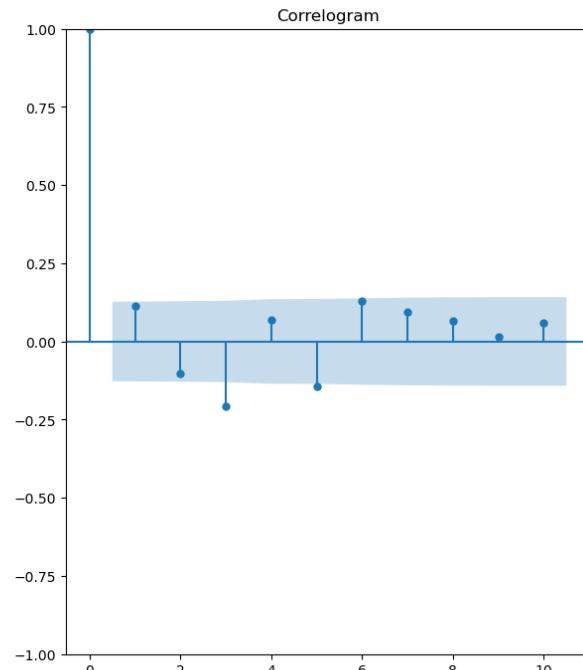
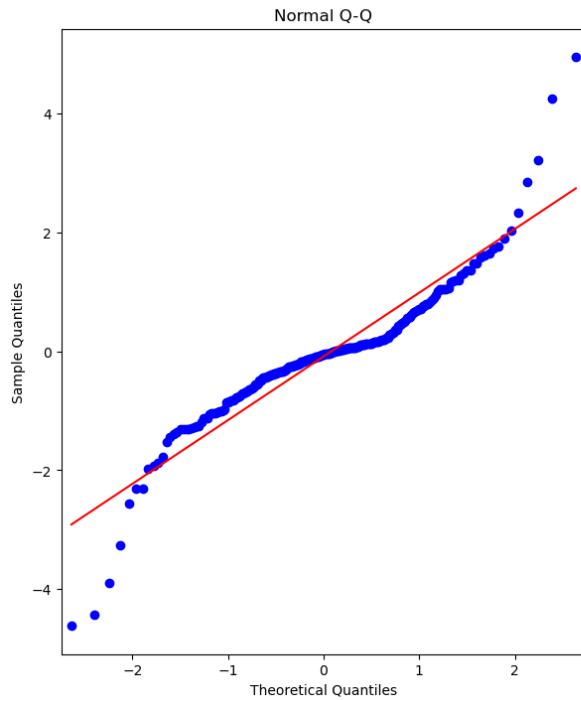
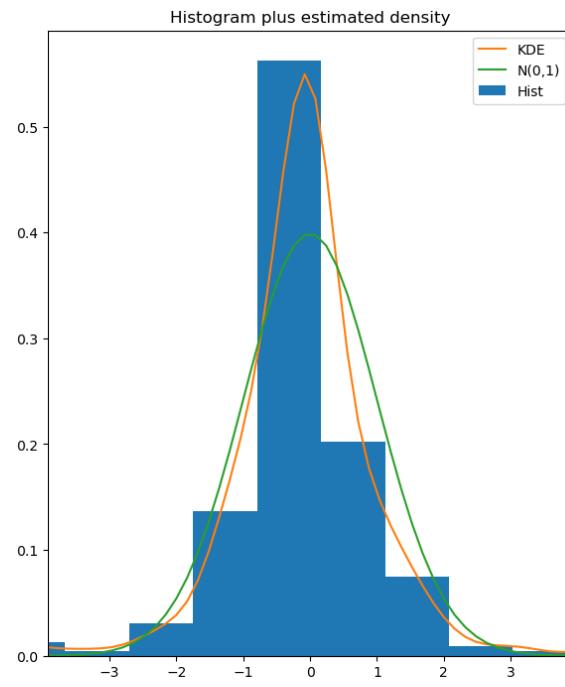
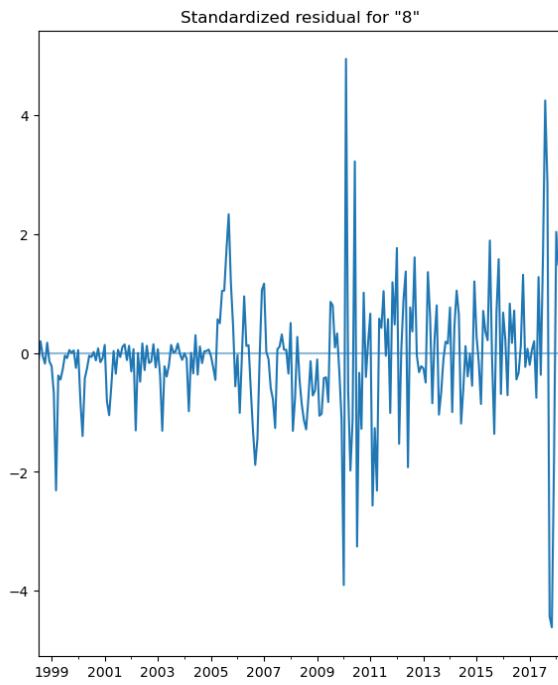
ar.L1	0.7428	0.038	19.304	0.000	0.667
0.818					
ma.L1	0.6268	0.033	18.860	0.000	0.562
0.692					
ma.S.L12	-0.4391	0.032	-13.887	0.000	-0.501
-0.377					
sigma2	5.895e+05	3.31e+04	17.792	0.000	5.25e+05
6.54e+05					
=====					
=====					

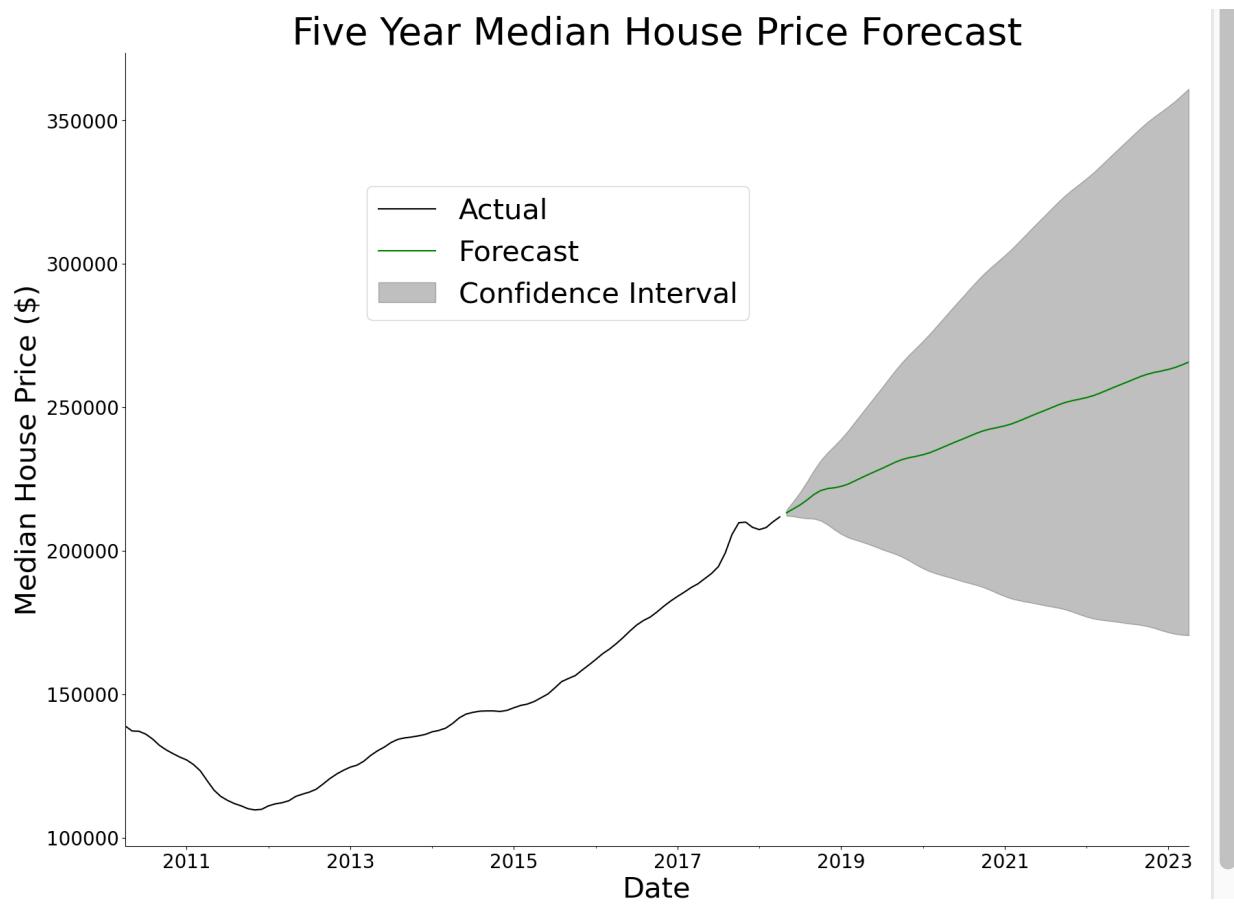




```
['83705']
(1, 1, 1) (1, 1, 1, 12)
```

	coef	std err	z	P> z	[0.025
0.975]					
ar.L1	0.8608	0.021	41.553	0.000	0.820
0.901					
ma.L1	0.4534	0.029	15.615	0.000	0.396
0.510					
ar.S.L12	0.1321	0.032	4.097	0.000	0.069
0.195					
ma.S.L12	-0.9555	0.113	-8.433	0.000	-1.178
-0.733					
sigma2	2.787e+05	2.93e+04	9.506	0.000	2.21e+05
3.36e+05					





In [28]:

```

1 # Create column for predicted price after 5 years
2 mountain_towns['5_Year_Price'] = pred_list_2
3
4 # Calculate Percent Return from time of investment
5 mountain_towns['Rate_Return'] = ((
6     mountain_towns['5_Year_Price']-
7     mountain_towns['2018-04'])/
8     mountain_towns['2018-04'])*100
9
10 mountain_towns

```

executed in 62ms, finished 11:55:23 2022-05-03

Out[28]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	1996-06
1776	87505	Santa Fe	NM	Santa Fe	Santa Fe	1777	188200.0	188400.0	188500.0
6810	87506	Santa Fe	NM	Santa Fe	Santa Fe	6811	227000.0	227400.0	227900.0
1446	97701	Bend	OR	Bend	Deschutes	1447	136300.0	136100.0	135900.0
706	97702	Bend	OR	Bend	Deschutes	707	160300.0	160300.0	160300.0
4081	84601	Provo	UT	Provo	Utah	4082	142700.0	140500.0	138400.0
4442	84606	Provo	UT	Provo	Utah	4443	142700.0	141400.0	140100.0
6115	80305	Boulder	CO	Boulder	Boulder	6116	187900.0	187700.0	187500.0
3363	80303	Boulder	CO	Boulder	Boulder	3364	231100.0	230300.0	229500.0
5764	83703	Boise	ID	Boise City	Ada	5765	107500.0	107500.0	107500.0
2962	83705	Boise	ID	Boise City	Ada	2963	83000.0	83000.0	83000.0

10 rows × 273 columns

In [29]:

```
1 rate = mountain_towns[ [ 'Rate_Return', 'City' ] ]  
2 rate
```

executed in 23ms, finished 11:55:23 2022-05-03

Out[29]:

	Rate_Return	City
1776	40.979317	Santa Fe
6810	56.185596	Santa Fe
1446	46.366086	Bend
706	61.959046	Bend
4081	21.852131	Provo
4442	19.863232	Provo
6115	17.734878	Boulder
3363	30.563084	Boulder
5764	55.076165	Boise
2962	25.482573	Boise

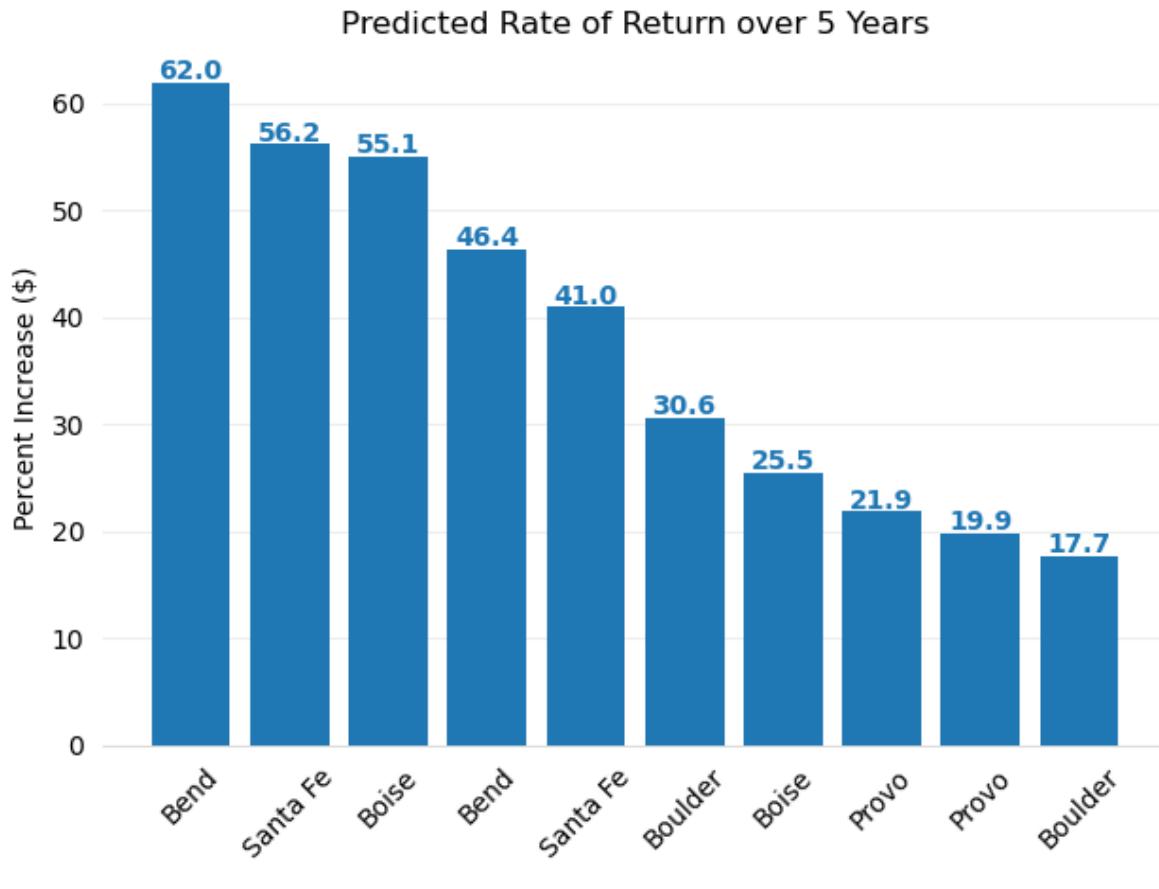
In [69]:

```

1 # create df and sort by Rate_Return in Descending Order
2 rate_2 = mountain_towns[['Rate_Return','City']]
3 #df.sort_values('importances',inplace=True)
4 rate_2.sort_values('Rate_Return',inplace=True,ascending=False)
5
6 fig, ax = plt.subplots()
7
8 # Save the chart so we can loop through the bars below.
9 bars = ax.bar(
10     x=np.arange(rate_2['Rate_Return'].size),
11     height=rate_2['Rate_Return'],
12     tick_label=rate_2['City']
13 )
14
15 # Axis formatting.
16 ax.spines['top'].set_visible(False)
17 ax.spines['right'].set_visible(False)
18 ax.spines['left'].set_visible(False)
19 ax.spines['bottom'].set_color('#DDDDDD')
20 ax.tick_params(bottom=False, left=False)
21 ax.set_axisbelow(True)
22 ax.yaxis.grid(True, color='#EEEEEE')
23 ax.xaxis.grid(False)
24
25 # Grab the color of the bars so we can make the
26 # text the same color.
27 bar_color = bars[0].get_facecolor()
28
29 # Add text annotations to the top of the bars.
30 # Note, you'll have to adjust this slightly (the 0.3)
31 # with different data.
32 for bar in bars:
33     ax.text(
34         bar.get_x() + bar.get_width() / 2,
35         bar.get_height() + 0.3,
36         round(bar.get_height(), 1),
37         horizontalalignment='center',
38         color=bar_color,
39         weight='bold'
40     )
41
42 plt.xticks(rotation = 45)
43 plt.title('Predicted Rate of Return over 5 Years')
44 plt.ylabel('Percent Increase ($)')
45
46 fig.tight_layout()

```

executed in 459ms, finished 13:56:16 2022-05-03



5.1.2 Summary

- Due to the work-from-home movement following the covid pandemic, these zip codes are currently ~50% above the predicted price, which represents a rate of return well over 100% in most zip codes.

5.1.2.1 Provo, Utah 84601 and 84606

- Both of the Provo zip codes had good rates of return (21.9% and 19.9%) and had the least potential downside.

5.1.2.2 Santa Fe, New Mexico 87506

- This zip code had the highest rate of return (56.2%), but did have a greater potential downside than Provo, UT.

5.1.2.3 Boise, Idaho 83703

- Predicted increase of 55.1% with very little risk.
- Probably the best investment potential for mountain towns.

5.2 Mid-Size Cities

- Jacksonville, FL
- Birmingham, AL
- Kansas City, MO
- Oklahoma City, OK
- Salt Lake City, UT

In [30]:

```

1 midsize_list = []
2 midsize_list.append(['Salt Lake City', 'Oklahoma City', 'Jacksonville', 'Birmi
3
4 midsize_cities = pd.DataFrame()
5
6 mid = pd.DataFrame(midsize_list)
7 mid = mid.T
8
9 for i in mid[0]:
10     # Create df from new city
11     ok = df1[df1['City'] == str(i)]
12
13     # add column with percent change between 2017 and 2012 median value
14     ok['Percent Change'] = (ok['2018-04'] / ok['2012-04']) * 100
15
16     # create new dataframe with top 10 zip codes for Percent Change
17     ok = ok.sort_values('Percent Change', ascending = False)
18
19     # new df_percent with only top ten results
20     ok = ok.head(2)
21
22     # concatenate to mountain towns df
23     midsize_cities = pd.concat([ok,midsize_cities])
24
25 midsize_cities = midsize_cities.drop(['RegionID','Percent Change'], axis=1)
26
27 midsize_cities

```

executed in 204ms, finished 11:55:24 2022-05-03

Out[30]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05
4121	66102	Kansas City	KS	Kansas City	Wyandotte	4122	38800.0	38800.0
4293	66104	Kansas City	KS	Kansas City	Wyandotte	4294	41300.0	41200.0
4502	48009	Birmingham	MI	Detroit	Oakland	4503	242700.0	242600.0
8459	35222	Birmingham	AL	Birmingham	Jefferson	8460	103400.0	104200.0
8259	32204	Jacksonville	FL	Jacksonville	Duval	8260	41400.0	41400.0
2596	32211	Jacksonville	FL	Jacksonville	Duval	2597	62000.0	61700.0
9540	73103	Oklahoma City	OK	Oklahoma City	Oklahoma	9541	46800.0	47000.0
5702	73114	Oklahoma City	OK	Oklahoma City	Oklahoma	5703	NaN	NaN
9058	84101	Salt Lake City	UT	Salt Lake City	Salt Lake	9059	85100.0	85400.0

ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	
6557	84111	Salt Lake City	UT	Salt Lake City	Salt Lake	6558	116700.0	116900.0

10 rows × 271 columns

5.2.1 Predicted Values

In [31]:

```

1 # Create list for final predicted value
2 pred_list_3 = []
3
4 # Construct model from entirety of data - no train/test
5 # iterate over prospects to extract optimal SARIMA values,
6 # instantiate SARIMAX model, make predictions
7 # and produce forecast with confidence intervals
8 for zipcode in midsize_cities['ZipCode']:
9     zips = midsize_cities.loc[midsize_cities['ZipCode'] == zipcode]
10    m1 = melt_data(zips)
11    m1[str(zipcode)] = m1['value']
12    m1 = m1.drop('value', axis=1)
13
14    # do NOT create train and test sets
15    #train = m1.iloc[:238]
16    #test = m1.iloc[238:]
17
18    # Run a grid with pdq and seasonal pdq parameters calculated above and ge
19    ans = []
20    outputs = []
21    comb_list = []
22    combs_list = []
23
24    for comb in pdq:
25        for combs in pdqs:
26            try:
27                mod = sm.tsa.statespace.SARIMAX(m1,
28                                                order=comb,
29                                                seasonal_order=combs,
30                                                enforce_stationarity=False,
31                                                enforce_invertibility=False)
32
33                output = mod.fit()
34                ans.append([comb, combs, output.aic])
35                outputs.append([output.aic])
36                comb_list.append([comb])
37                combs_list.append([combs])
38                AIC_df = pd.DataFrame(list(outputs, comb_list, combs_list))
39                #print('ARIMA {} x {}12 : AIC Calculated ={}'.format(comb, co
40            except:
41                continue
42
43    # Print the lowest AIC and it's parameters for SARIMA
44    AIC_df[0] = AIC_df[0].str.get(0)
45    best_AIC = AIC_df.copy()
46    lowest_AIC = min(best_AIC[0])
47    lowest_AIC_row = AIC_df.loc[AIC_df[0] == lowest_AIC]
48    #print(lowest_AIC_row)
49
50    # call value for comb and combs for model input
51    lowest_AIC_row[1] = lowest_AIC_row[1].str.get(0)
52    lowest_AIC_row[2] = lowest_AIC_row[2].str.get(0)

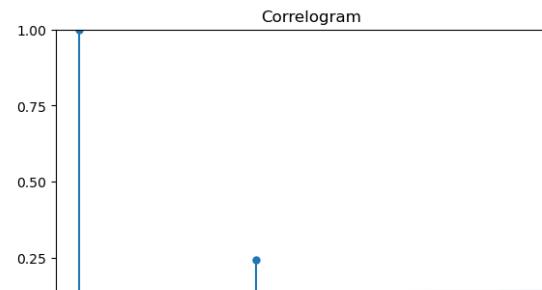
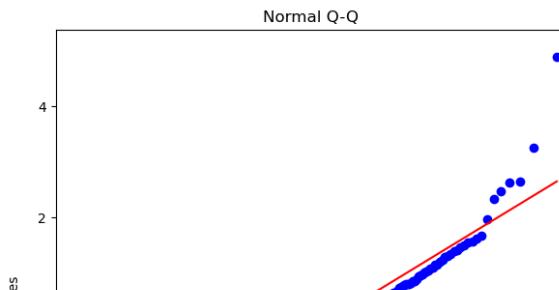
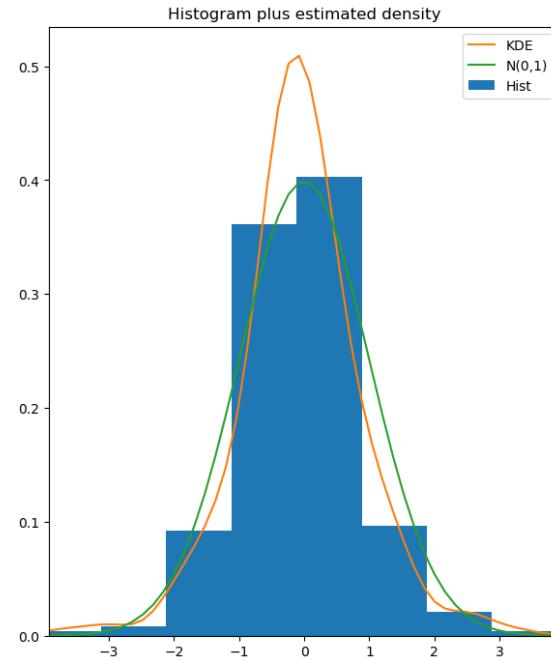
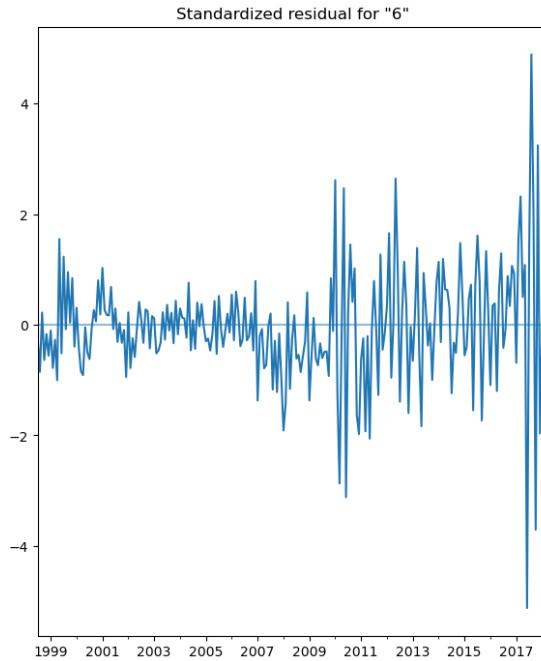
```

```
53     arima = lowest_AIC_row.iat[0,1]
54     s = lowest_AIC_row.iat[0,2]
55     print(str(m1.columns.values))
56     print(arima, s)
57
58     # Plug the optimal parameter values into a new SARIMAX model
59     ARIMA_MODEL = sm.tsa.statespace.SARIMAX(m1,
60                                              order=arima,
61                                              seasonal_order=s,
62                                              enforce_stationarity=False,
63                                              enforce_invertibility=False)
64
65     # Fit the model and print results
66     output = ARIMA_MODEL.fit()
67
68     print(output.summary().tables[1])
69
70     # Call plot_diagnostics() on the results calculated above
71     output.plot_diagnostics(figsize=(15, 18))
72     plt.show()
73
74     # Get actual results in dataframe
75     price_truth = m1['2018-04-01':]
76
77     # Get forecast 500 steps ahead in future
78     prediction = output.get_forecast(steps=60)
79
80     # Create predictions df
81     pred = prediction.predicted_mean
82
83     # Add prediction for 5 year forecasted value to list
84     pred_list_3.append(pred['2023-04-01'])
85
86     # Get confidence intervals of forecasts
87     pred_conf = prediction.conf_int()
88
89     # Only use data from April 2010 to improve visual of forecast
90     plot_vals = m1.loc['2010-04-01':]
91
92     # Plot future predictions with confidence intervals
93     ax = plot_vals.plot(label='observed', figsize=(20, 15), color='black')
94     prediction.predicted_mean.plot(ax=ax, label='Forecast', color='green')
95     ax.fill_between(pred_conf.index,
96                      pred_conf.iloc[:, 0],
97                      pred_conf.iloc[:, 1], color='k', alpha=0.25)
98     ax.set_xlabel('Date').set_fontsize(30)
99     ax.set_ylabel('Median House Price ($)').set_fontsize(30)
100    ax.set_title('Five Year Median House Price Forecast').set_fontsize(40)
101
102   for item in (ax.get_xticklabels() + ax.get_yticklabels()):
103       item.set_fontsize(20)
104
105   right_side = ax.spines["right"]
106   right_side.set_visible(False)
```

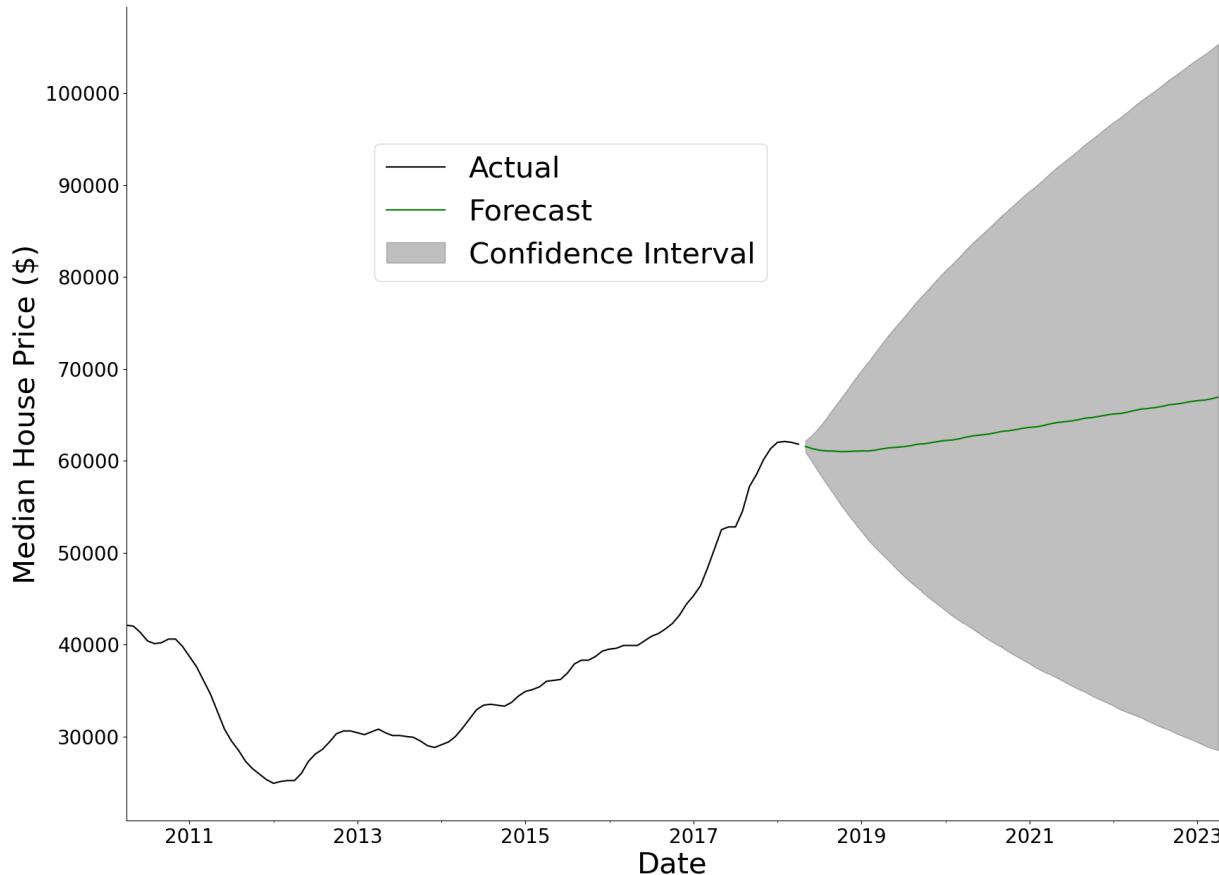
```
107      top_side = ax.spines["top"]
108      top_side.set_visible(False)
109
110      mylabels = 'Actual', 'Forecast', 'Confidence Interval'
111      plt.legend(fontsize=30, labels=mylabels, bbox_to_anchor=(0.6,0.85))
112      plt.show()
113      print()
114      print()
```

executed in 4m 5s, finished 11:59:28 2022-05-03

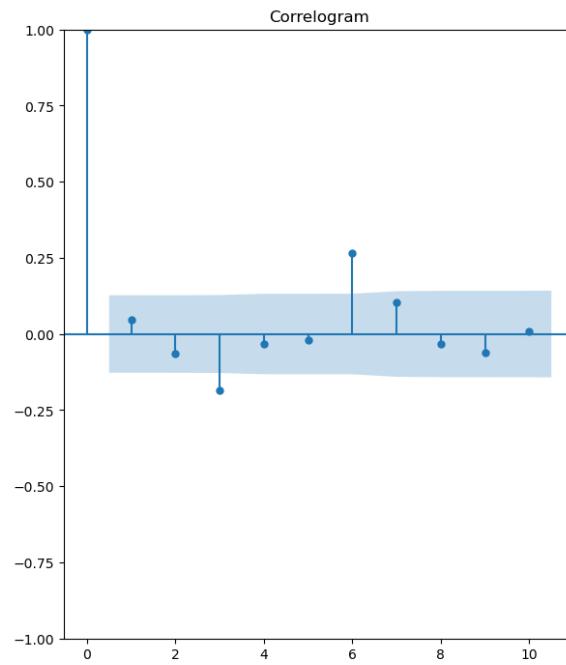
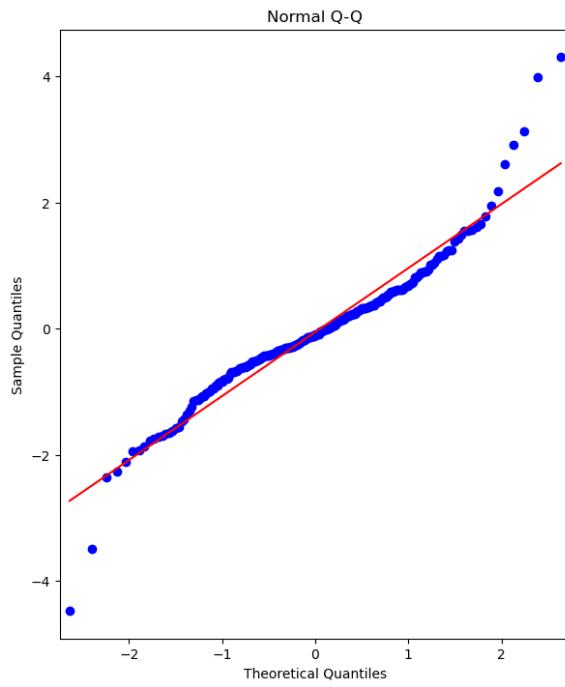
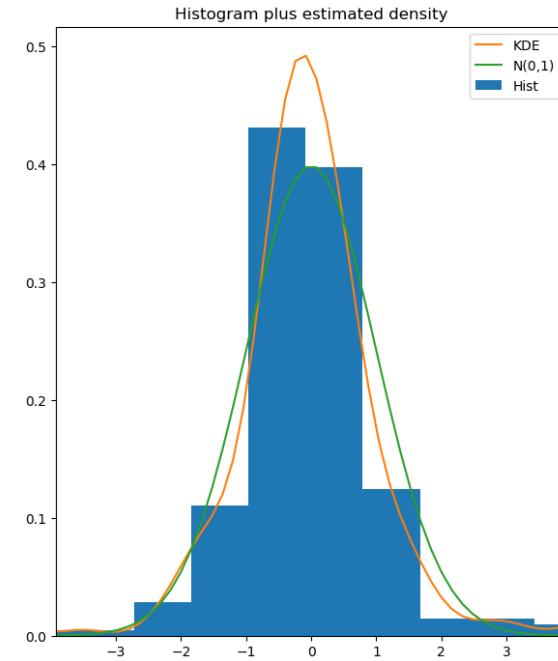
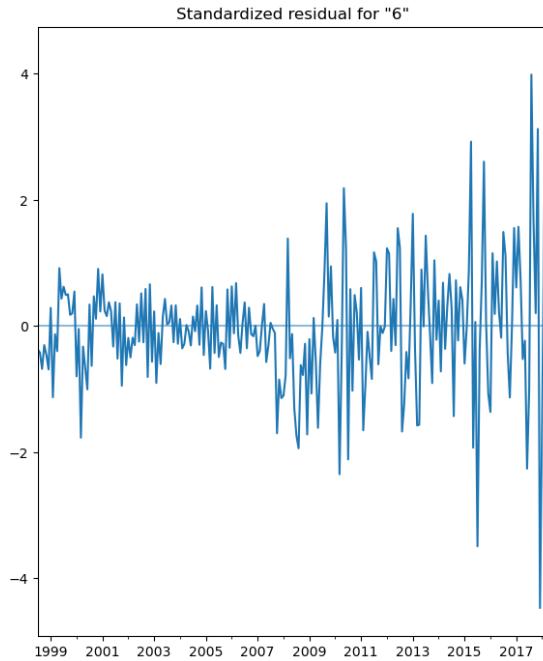
```
['66102']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.8185      0.029      28.179      0.000      0.762
0.875
ma.L1      0.4528      0.035      12.856      0.000      0.384
0.522
ma.S.L12   -0.9918      0.559     -1.775      0.076     -2.087
0.103
sigma2     9.407e+04    4.85e+04      1.939      0.052    -1012.390
1.89e+05
=====
=====
```



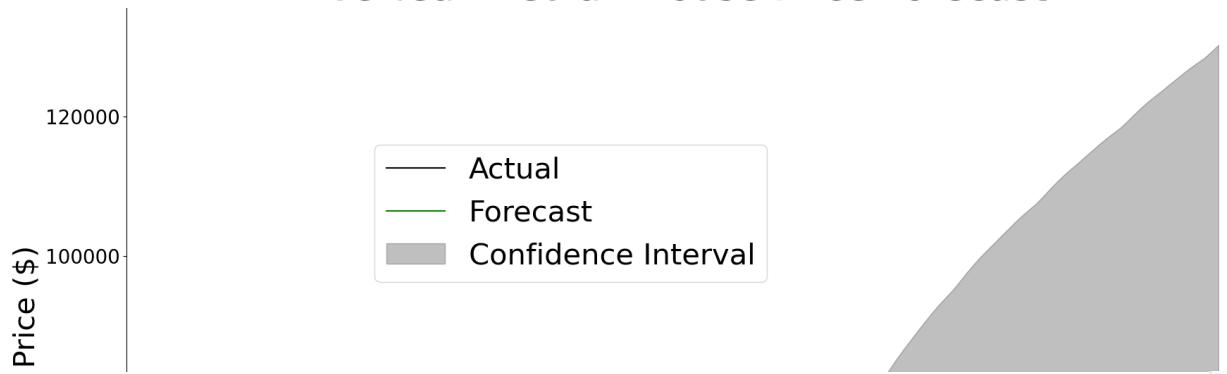
Five Year Median House Price Forecast



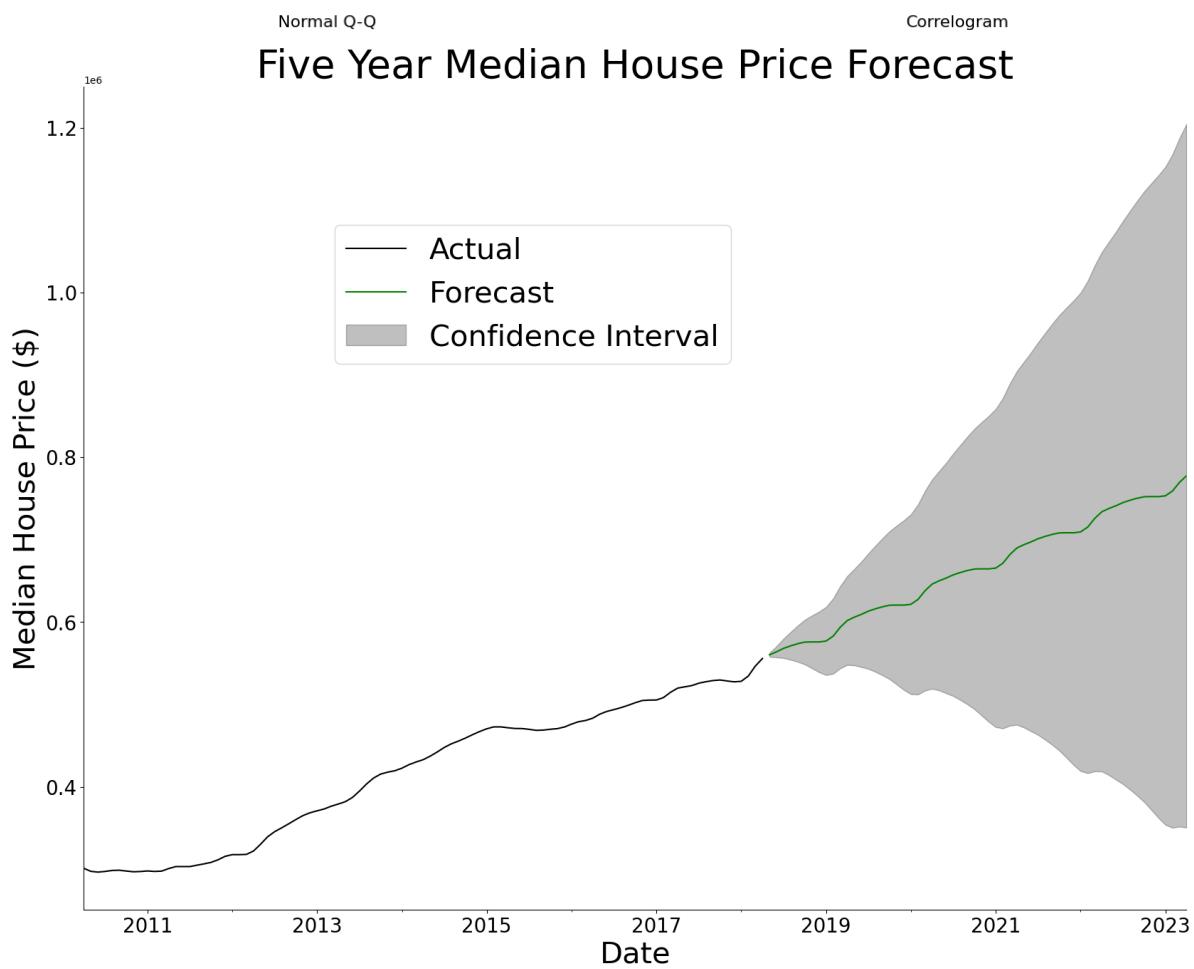
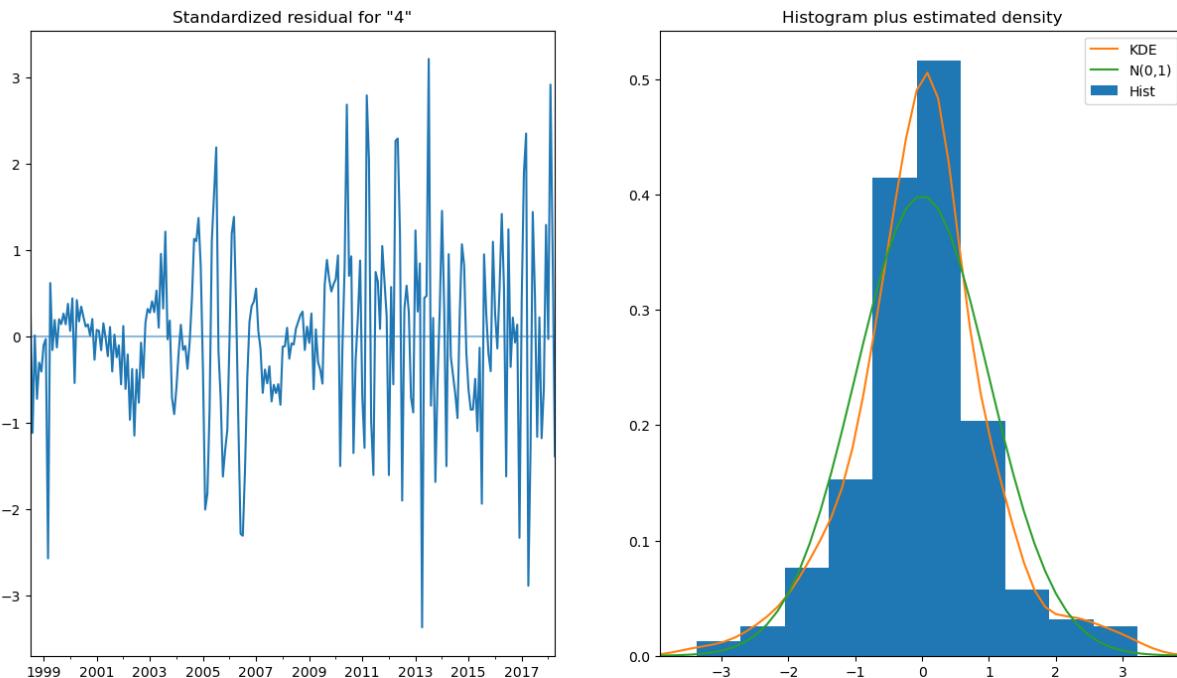
```
['66104']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
          coef      std err       z     P>|z|      [ 0.025
0.975]
-----
ar.L1      0.8532      0.028    29.979      0.000      0.797
0.909
ma.L1      0.6579      0.048    13.745      0.000      0.564
0.752
ma.S.L12   -0.9592      0.115    -8.349      0.000     -1.184
-0.734
sigma2     8.856e+04  8088.236    10.949      0.000    7.27e+04
1.04e+05
=====
```



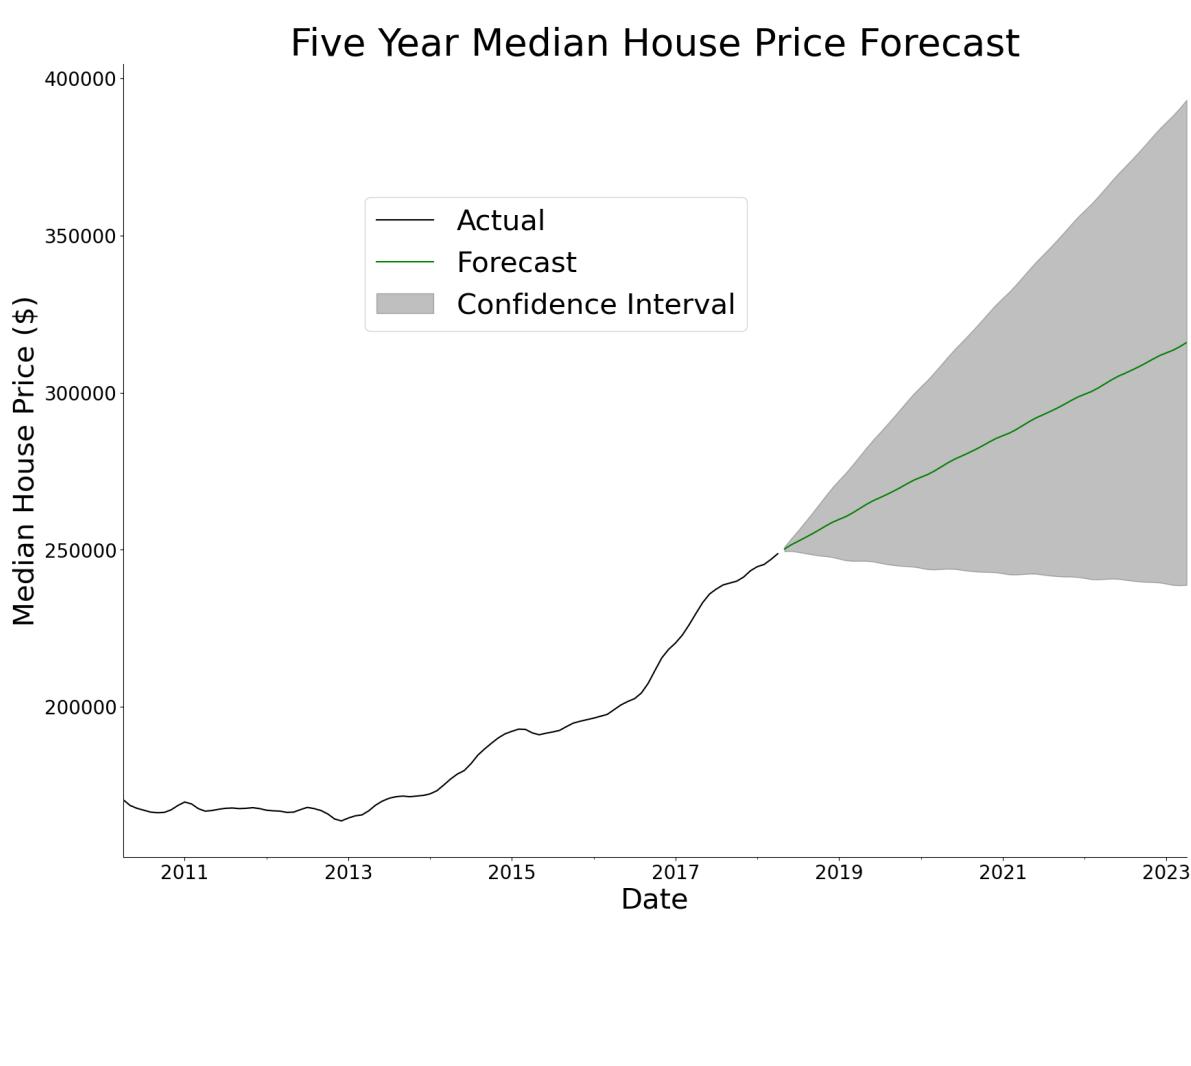
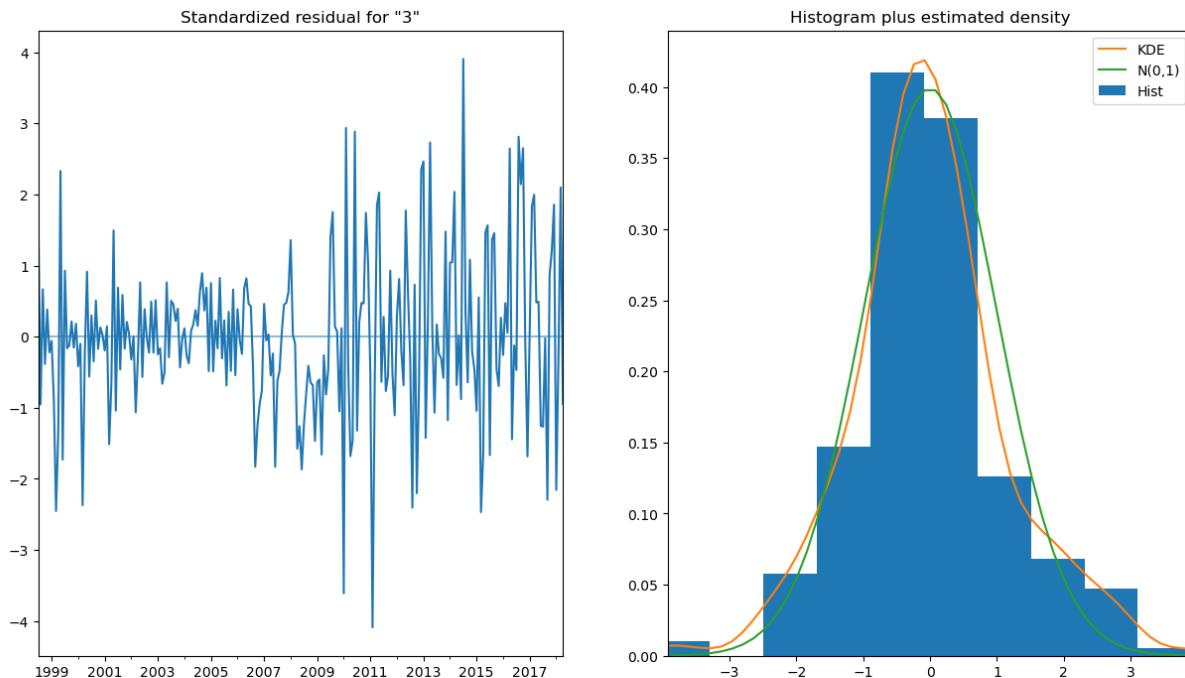
Five Year Median House Price Forecast



```
[ '48009' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.7974      0.041      19.250      0.000      0.716
0.879
ma.L1      0.8200      0.039      21.040      0.000      0.744
0.896
ar.S.L12   0.0829      0.017      4.856      0.000      0.049
0.116
ma.S.L12   -0.3089     0.024     -12.926      0.000     -0.356
-0.262
sigma2     1.574e+06    1.25e+05     12.583      0.000     1.33e+06
1.82e+06
=====
```



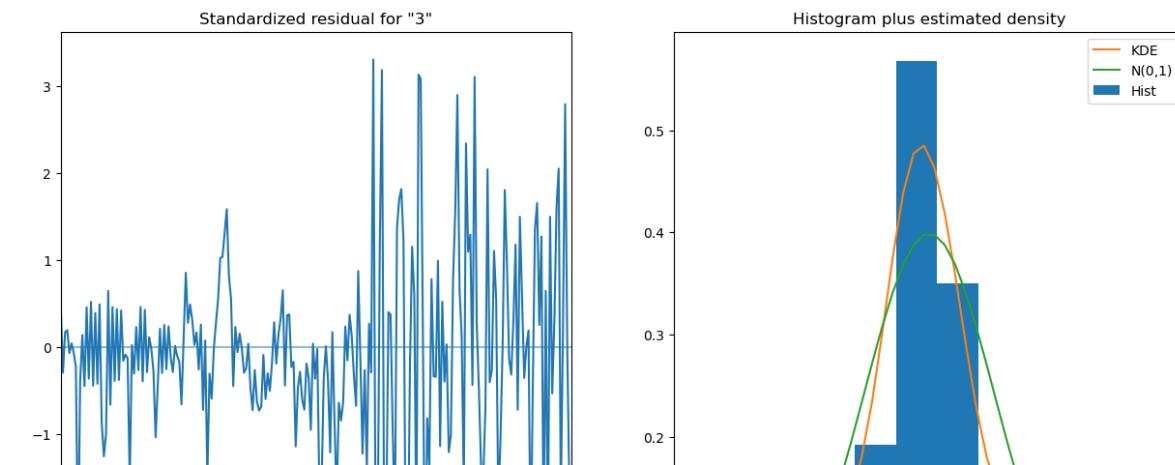
```
['35222']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.8322      0.029      28.728      0.000      0.775
0.889
ma.L1      0.6149      0.033      18.823      0.000      0.551
0.679
ma.S.L12   -0.7776      0.045     -17.446      0.000     -0.865
-0.690
sigma2     1.534e+05    1.06e+04     14.487      0.000     1.33e+05
1.74e+05
=====
```



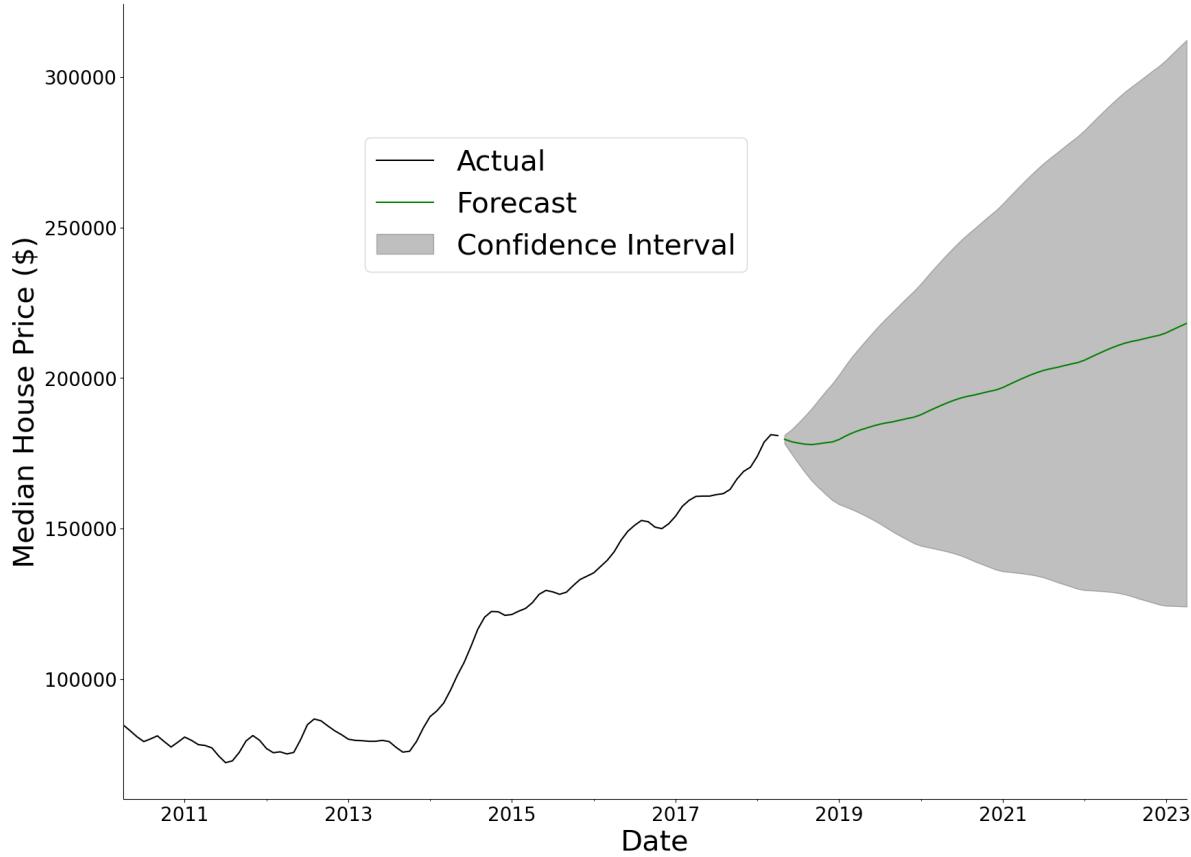
```
[ '32204' ]
```

(1, 1, 1) (1, 1, 1, 12)

	coef	std err	z	P> z	[0.025
0.975]					
ar.L1	0.7482	0.039	19.425	0.000	0.673
0.824					
ma.L1	0.7681	0.029	26.890	0.000	0.712
0.824					
ar.S.L12	0.0770	0.048	1.616	0.106	-0.016
0.170					
ma.S.L12	-0.9257	0.056	-16.407	0.000	-1.036
-0.815					
sigma2	5.603e+05	4.88e+04	11.494	0.000	4.65e+05
6.56e+05					



Five Year Median House Price Forecast

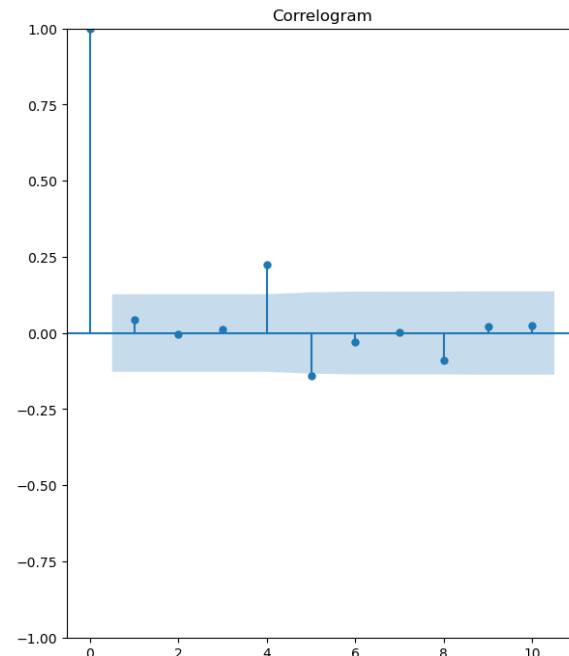
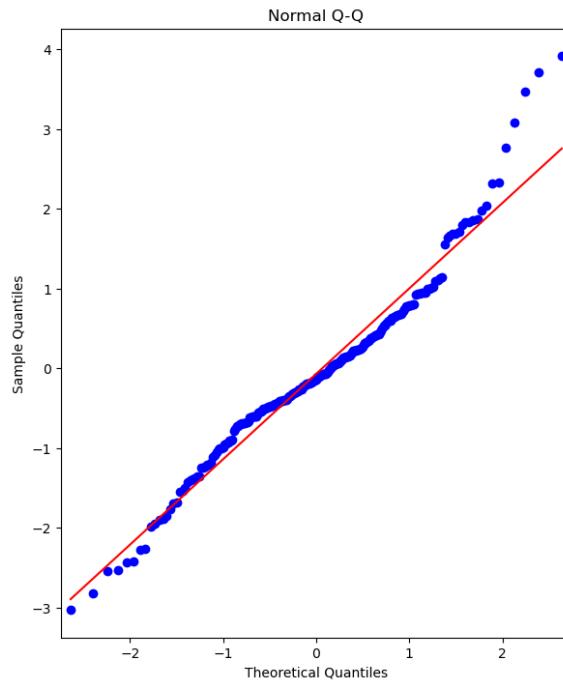
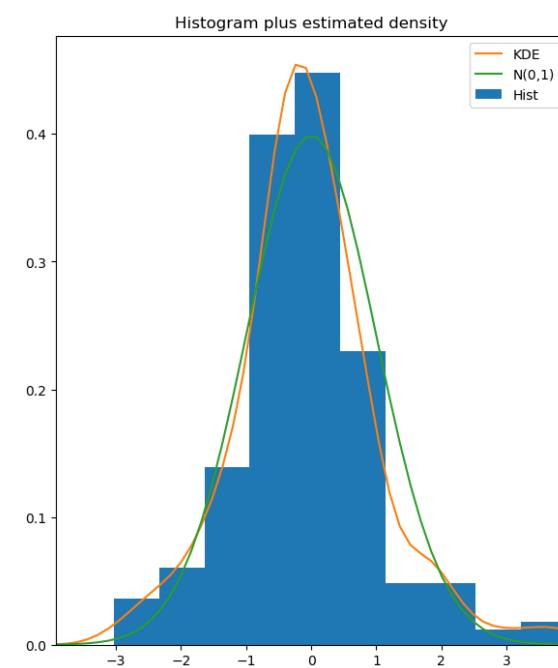
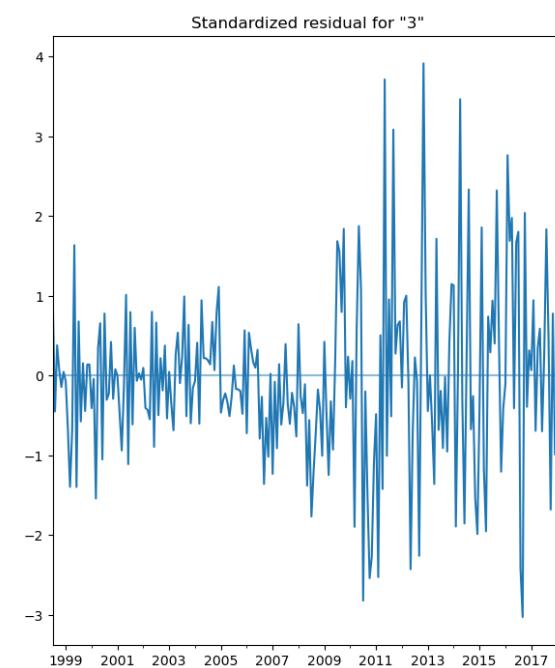


```

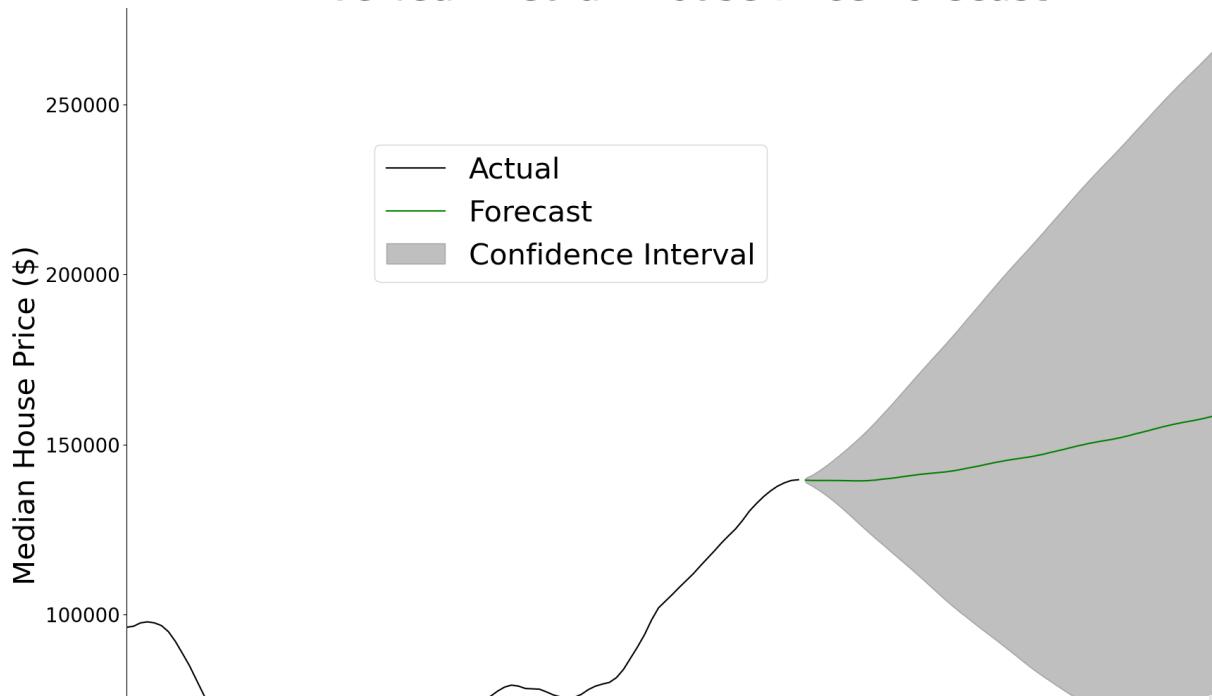
['32211']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
      coef    std err        z     P>|z|      [0.025
0.975]
-----
-----
ar.L1      0.9416    0.016   58.691      0.000      0.910
0.973
ma.L1      0.5652    0.043   13.280      0.000      0.482
0.649
ma.S.L12   -0.9980   1.560   -0.640      0.522     -4.056

```

2.060
 σ^2 9.894e+04 1.55e+05 0.639 0.523 -2.05e+05
4.02e+05
=====



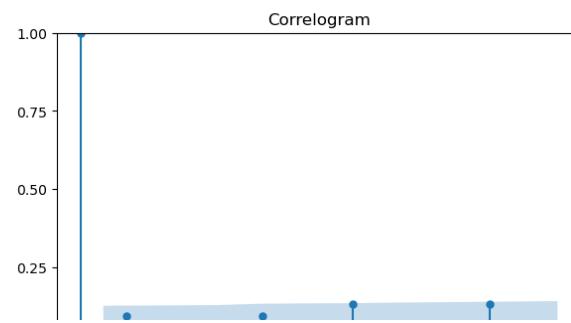
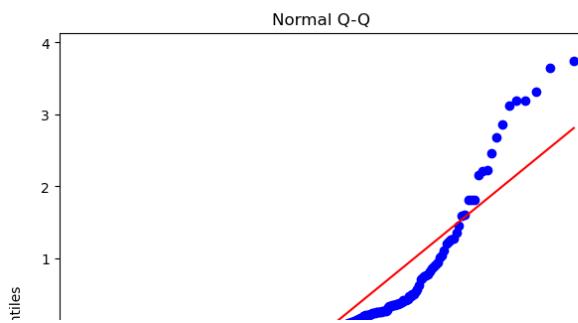
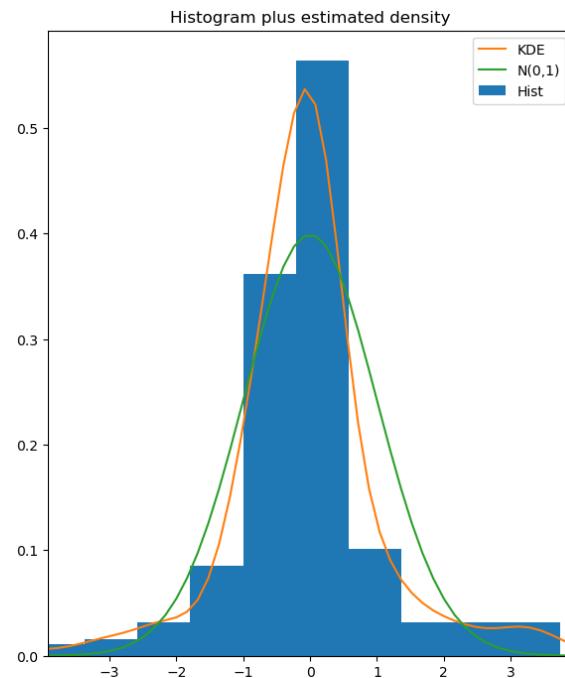
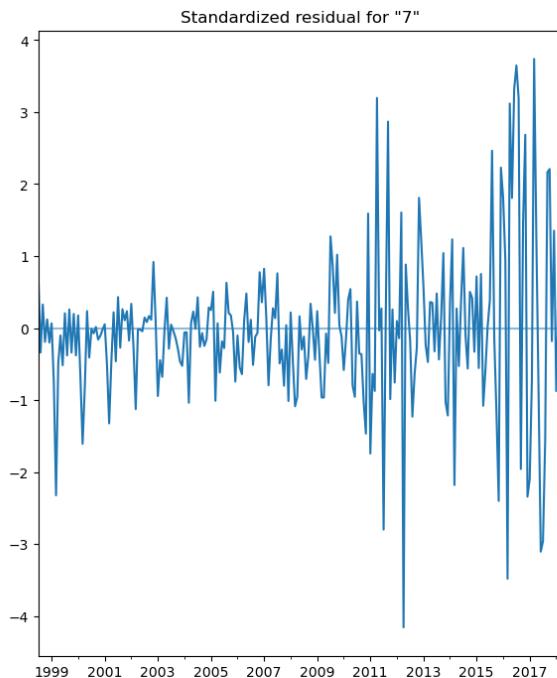
Five Year Median House Price Forecast



```
['73103']
(1, 1, 1) (1, 1, 1, 12)
=====
```

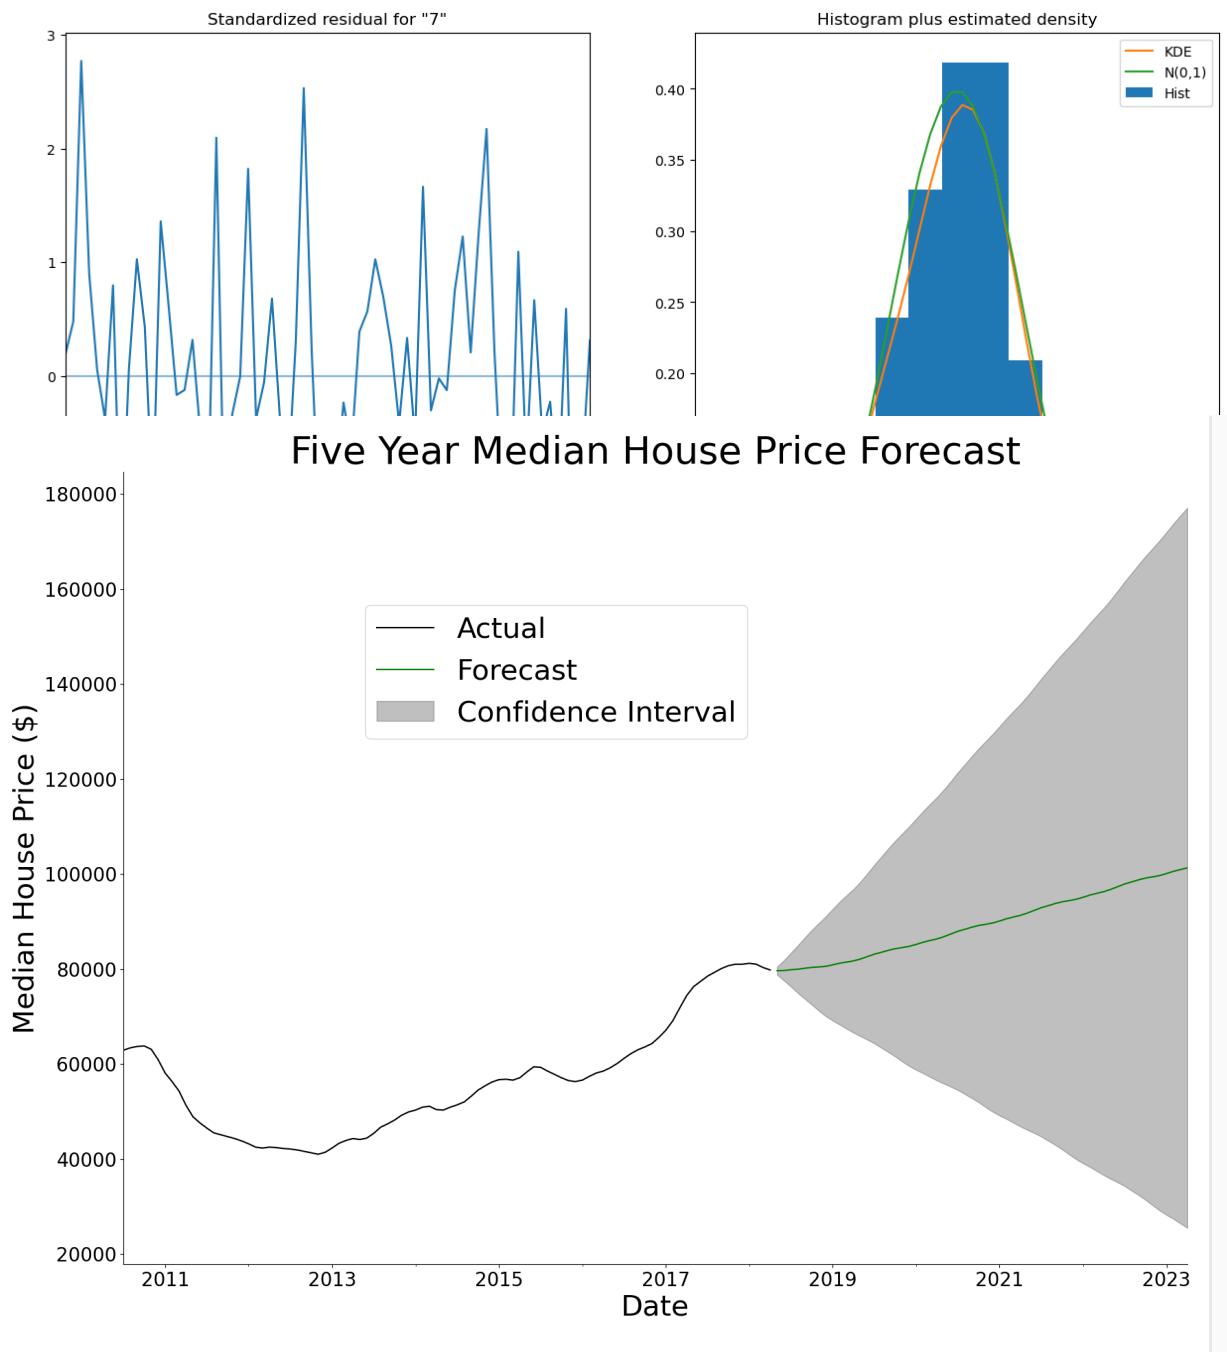
	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.7872	0.020	39.148	0.000	0.748
0.827					
ma.L1	0.6972	0.026	26.317	0.000	0.645
0.749					
ar.S.L12	0.0641	0.048	1.324	0.185	-0.031
0.159					
ma.S.L12	-0.9775	0.120	-8.173	0.000	-1.212
-0.743					
sigma2	3.537e+05	4.15e+04	8.527	0.000	2.72e+05
4.35e+05					
=====					
=====					



Five Year Median House Price Forecast

```
['73114']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.7765      0.099      7.846      0.000      0.583
0.970
ma.L1      0.5774      0.113      5.120      0.000      0.356
0.799
ar.S.L12   0.0694      0.029      2.364      0.018      0.012
0.127
ma.S.L12   -0.6903     0.135     -5.095      0.000     -0.956
-0.425
sigma2    1.827e+05  3.32e+04      5.505      0.000    1.18e+05
2.48e+05
=====
```

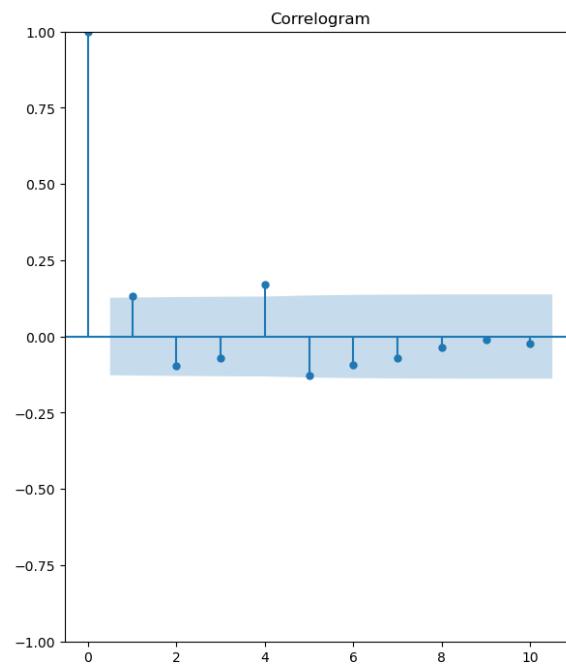
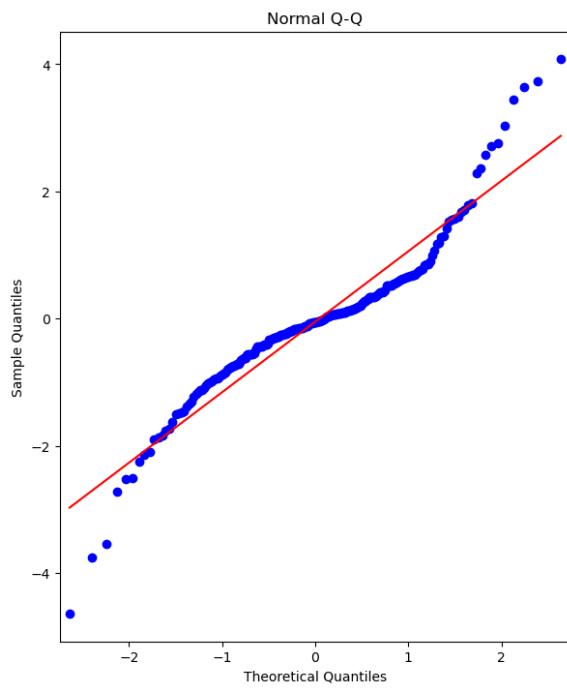
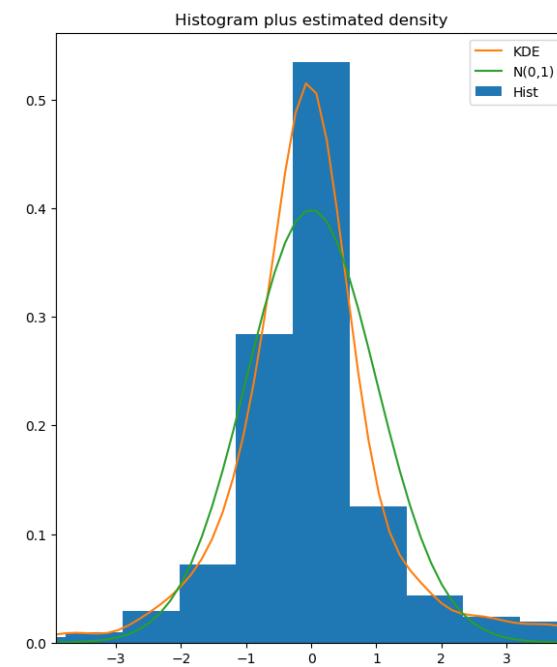
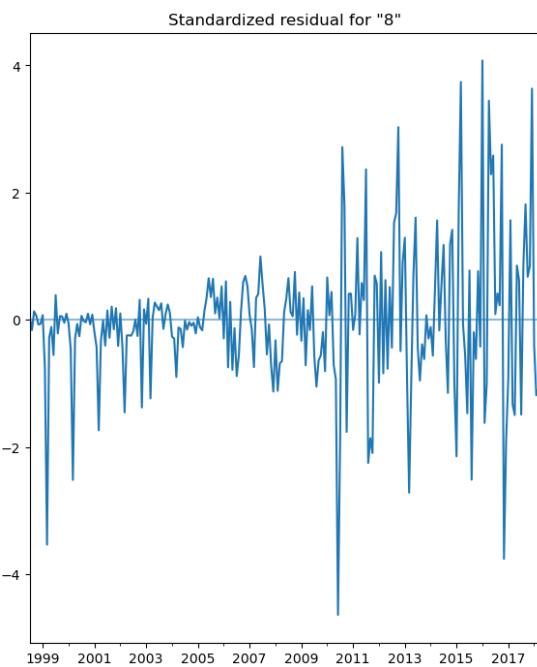


```
['84101']
(1, 1, 1) (1, 1, 1, 12)
=====
```

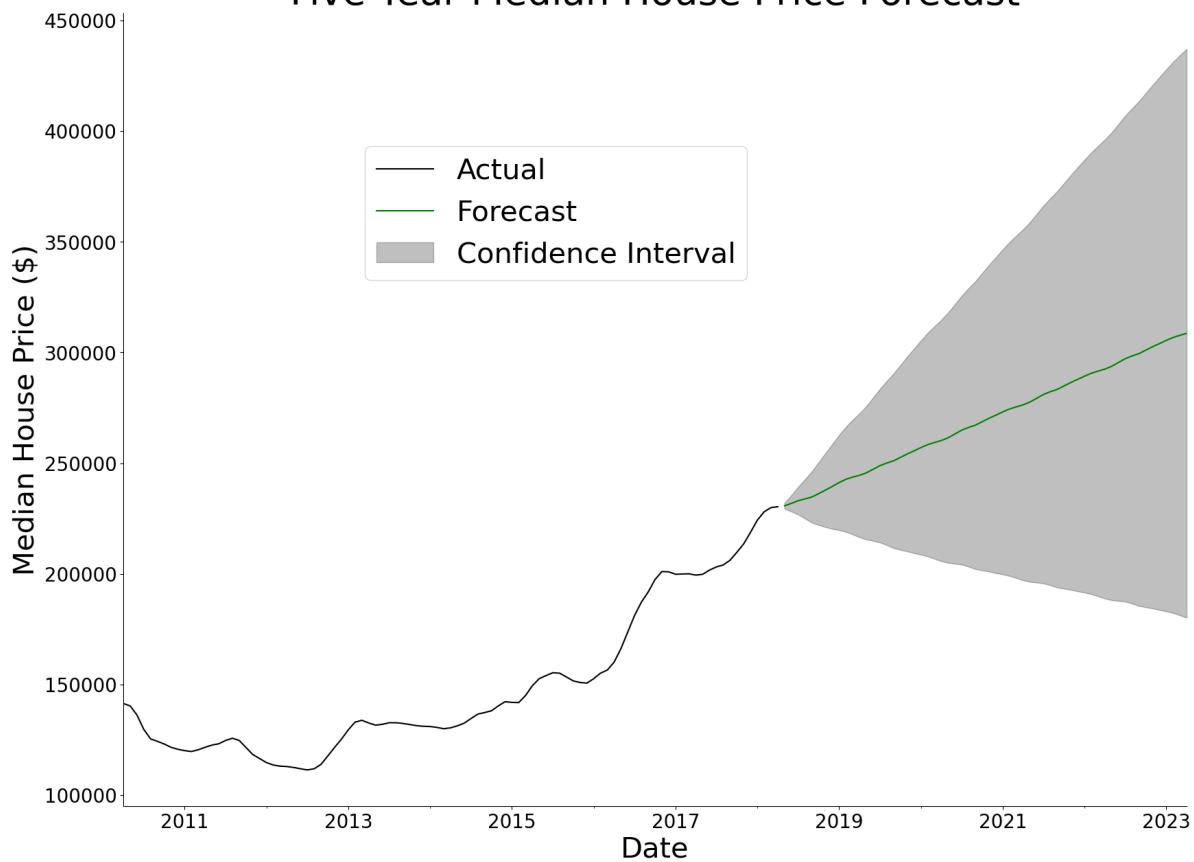
	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.7946	0.024	32.434	0.000	0.747
0.843					
ma.L1	0.6922	0.030	23.236	0.000	0.634
0.751					
ar.S.L12	0.1060	0.049	2.168	0.030	0.010
0.202					

ma.S.L12	-0.7942	0.050	-15.779	0.000	-0.893
-0.696					
sigma2	4.942e+05	2.9e+04	17.030	0.000	4.37e+05
5.51e+05					



Five Year Median House Price Forecast



```
['84111']
(1, 1, 1) (1, 1, 1, 12)
=====
```

```
=====
=====
```

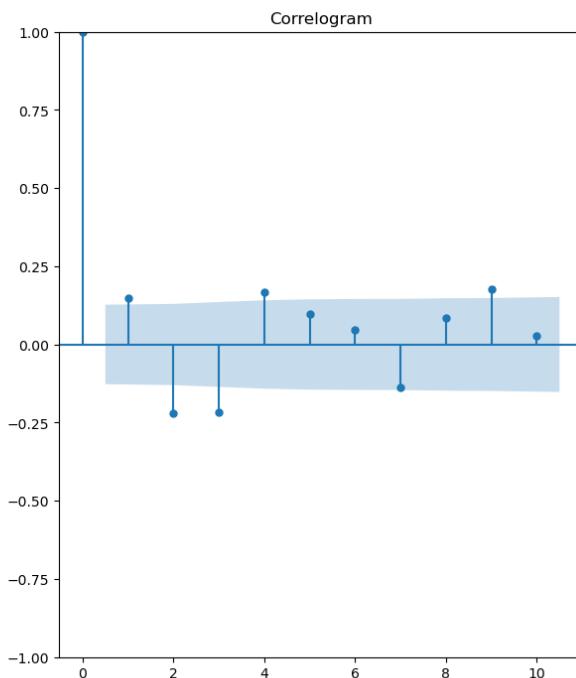
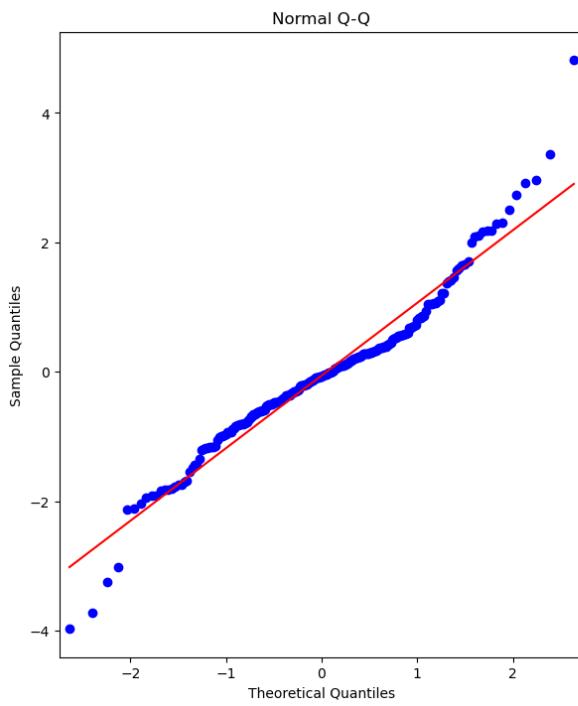
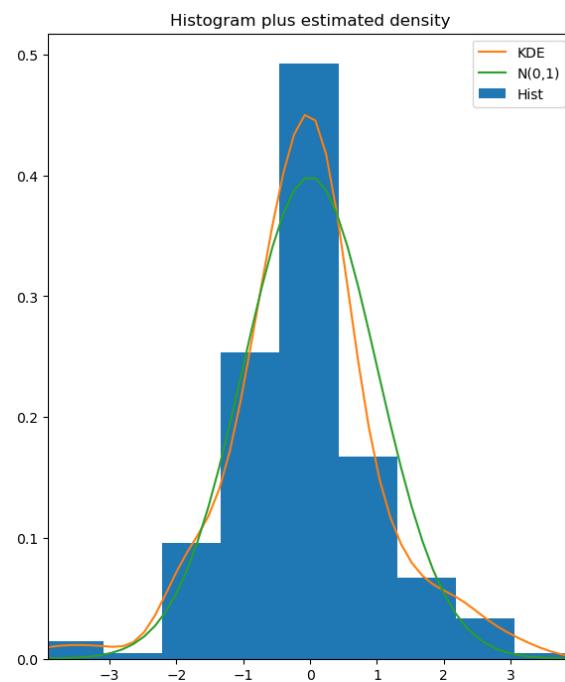
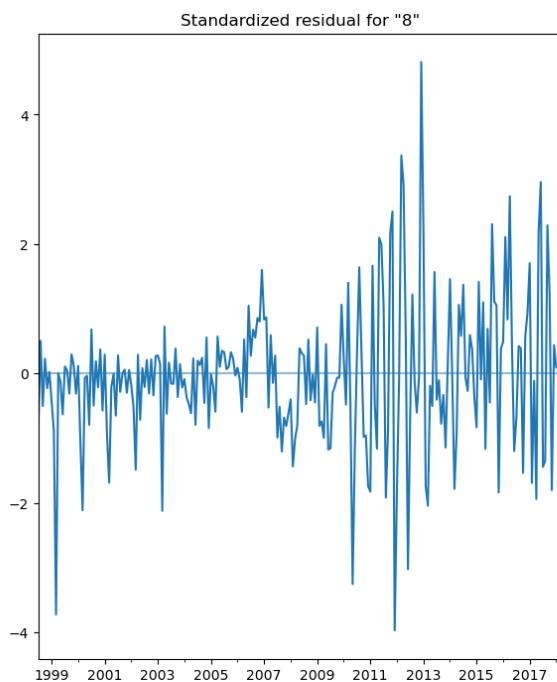
	coef	std err	z	P> z	[0.025
0.975]					

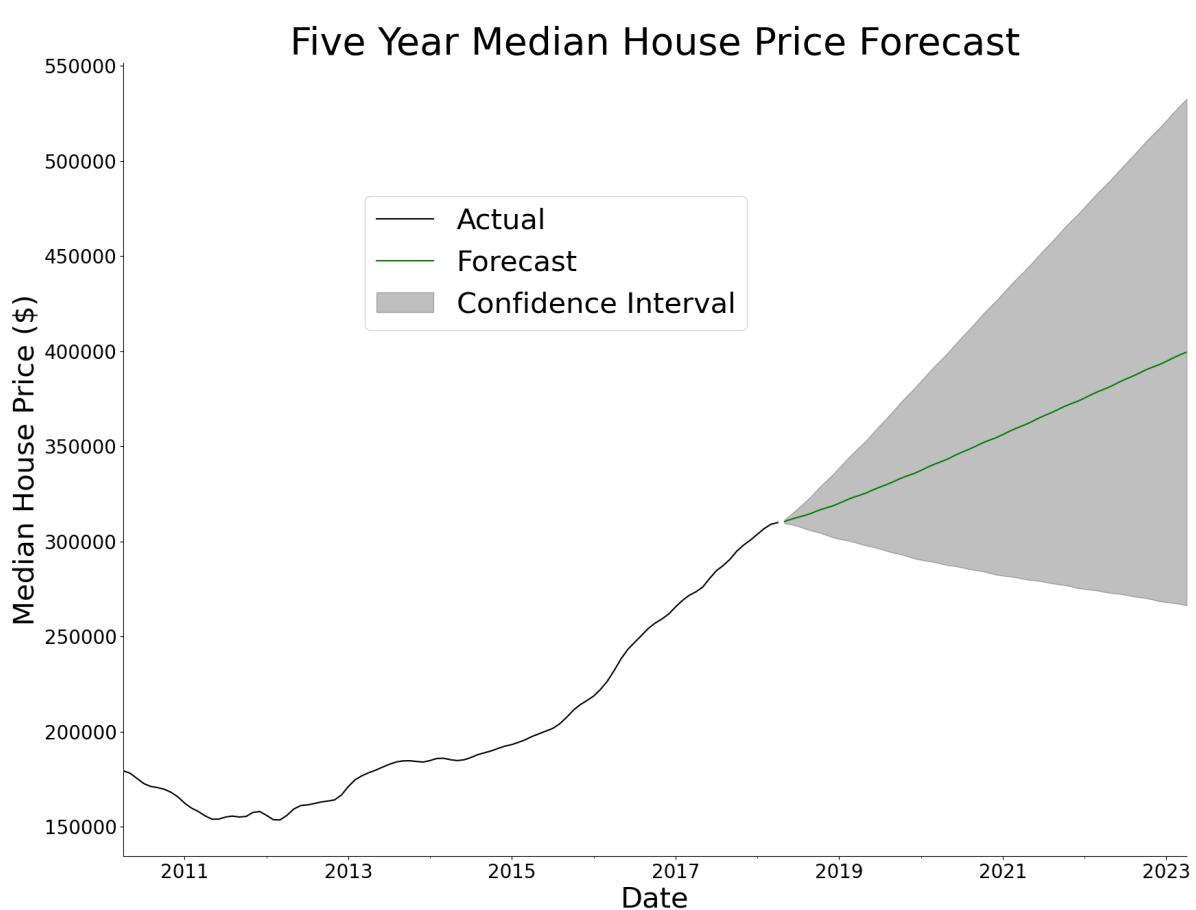
```
=====
-----
```

	ar.L1	0.8690	0.024	35.940	0.000	0.822
0.916						
ma.L1	0.6710	0.033	20.342	0.000	0.606	
0.736						
ar.S.L12	0.0808	0.021	3.808	0.000	0.039	

```

0.122
ma.S.L12      -0.8098     0.041    -19.724     0.000    -0.890
-0.729
sigma2        2.669e+05   1.75e+04   15.256     0.000    2.33e+05
3.01e+05
=====
=====
```





In [52]:

```

1 # Create column for predicted price after 5 years
2 midsize_cities['5_Year_Price'] = pred_list_3
3
4 # Calculate Percent Return from time of investment
5 midsize_cities['Rate_Return'] = ((
6     midsize_cities['5_Year_Price']-
7     midsize_cities['2018-04'])/
8     midsize_cities['2018-04'])*100
9
10 midsize_cities
11

```

executed in 50ms, finished 13:49:06 2022-05-03

Out[52]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05
4121	66102	Kansas City	KS	Kansas City	Wyandotte	4122	38800.0	38800.0
4293	66104	Kansas City	KS	Kansas City	Wyandotte	4294	41300.0	41200.0
4502	48009	Birmingham	MI	Detroit	Oakland	4503	242700.0	242600.0
8459	35222	Birmingham	AL	Birmingham	Jefferson	8460	103400.0	104200.0
8259	32204	Jacksonville	FL	Jacksonville	Duval	8260	41400.0	41400.0
2596	32211	Jacksonville	FL	Jacksonville	Duval	2597	62000.0	61700.0
9540	73103	Oklahoma City	OK	Oklahoma City	Oklahoma	9541	46800.0	47000.0
5702	73114	Oklahoma City	OK	Oklahoma City	Oklahoma	5703	NaN	NaN
9058	84101	Salt Lake City	UT	Salt Lake City	Salt Lake	9059	85100.0	85400.0
6557	84111	Salt Lake City	UT	Salt Lake City	Salt Lake	6558	116700.0	116900.0

10 rows × 273 columns

5.2.2 Bar Chart: Rate of Return

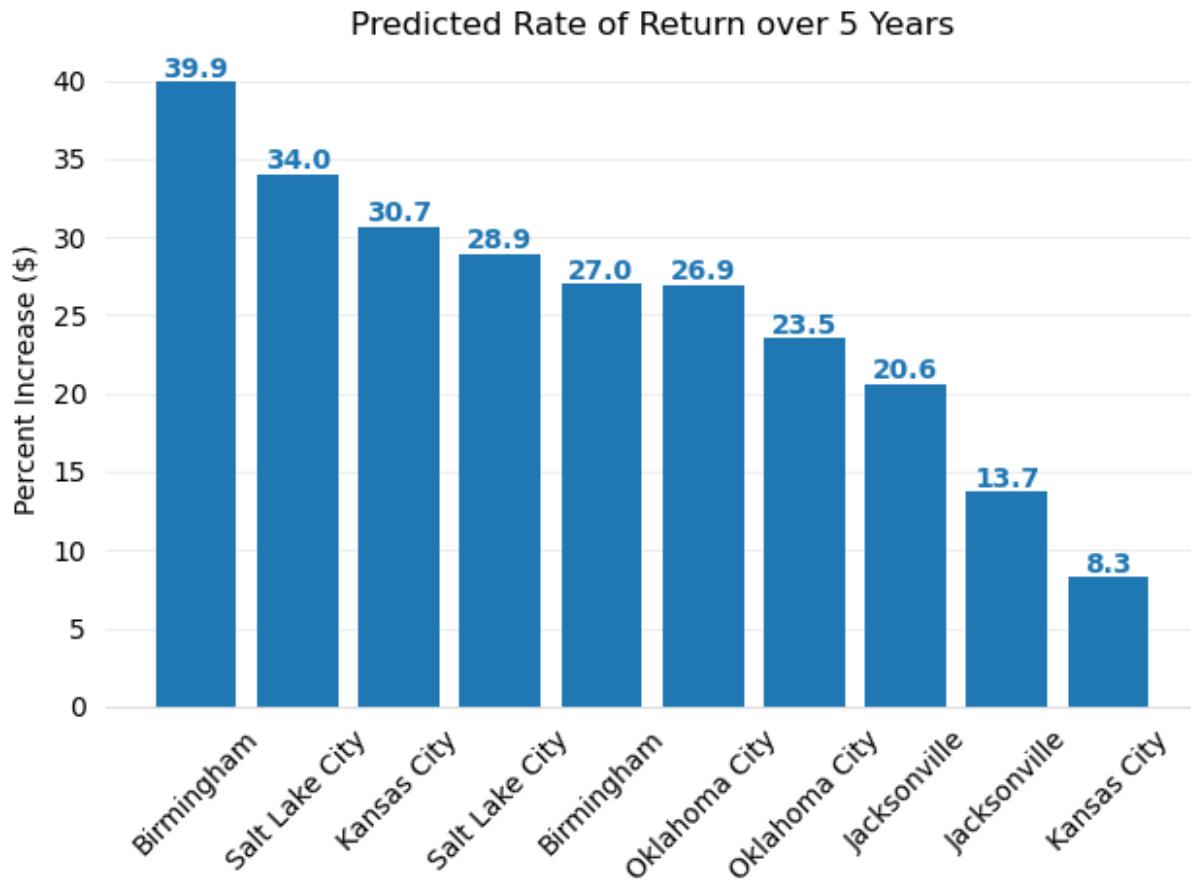
In [70]:

```

1 # create df and sort by Rate_Return in Descending Order
2 rate_3 = midsize_cities[['Rate_Return','City']]
3 #df.sort_values('importances',inplace=True)
4 rate_3.sort_values('Rate_Return',inplace=True,ascending=False)
5
6 fig, ax = plt.subplots()
7
8 # Save the chart so we can loop through the bars below.
9 bars = ax.bar(
10     x=np.arange(rate_3['Rate_Return'].size),
11     height=rate_3['Rate_Return'],
12     tick_label=rate_3['City']
13 )
14
15 # Axis formatting.
16 ax.spines['top'].set_visible(False)
17 ax.spines['right'].set_visible(False)
18 ax.spines['left'].set_visible(False)
19 ax.spines['bottom'].set_color('#DDDDDD')
20 ax.tick_params(bottom=False, left=False)
21 ax.set_axisbelow(True)
22 ax.yaxis.grid(True, color='#EEEEEE')
23 ax.xaxis.grid(False)
24
25 # Grab the color of the bars so we can make the
26 # text the same color.
27 bar_color = bars[0].get_facecolor()
28
29 # Add text annotations to the top of the bars.
30 # Note, you'll have to adjust this slightly (the 0.3)
31 # with different data.
32 for bar in bars:
33     ax.text(
34         bar.get_x() + bar.get_width() / 2,
35         bar.get_height() + 0.3,
36         round(bar.get_height(), 1),
37         horizontalalignment='center',
38         color=bar_color,
39         weight='bold'
40     )
41
42 plt.xticks(rotation = 45)
43 plt.title('Predicted Rate of Return over 5 Years')
44 plt.ylabel('Percent Increase ($)')
45
46 fig.tight_layout()

```

executed in 392ms, finished 13:56:55 2022-05-03



5.2.3 Summary:

- Mid-Size Cities do not have a great rate of return, however a few zip codes could be worth a closer look.
- Birmingham, Michigan 48009 and Salt Lake City, Utah 84101 both have return rates over 34%, but they also have higher risk than the top_prospects zip codes

6 SARIMA Model

- For Top Prospective Zip Codes

In [12]:

```
1 # Define the p, d and q parameters to take any value between 0 and 2
2 p = d = q = range(0, 2)
3
4 # Generate all different combinations of p, q and q triplets
5 pdq = list(itertools.product(p, d, q))
6
7 # Generate all different combinations of seasonal p, q and q triplets
8 pdqs = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
```

executed in 11ms, finished 10:46:33 2022-05-03

6.1 GridSearch optimal parameters, Model, Predict, calculate RMSE

- Run for all zip codes in prospects

In [13]:

```

1 # Create list to add rmse values during iteration
2 rmse_list = []
3
4
5
6 # iterate over prospects to extract optimal SARIMA values,
7 # instantiate SARIMAX model, make predictions
8 # and calculate the RMSE for each zip code
9 for zipcode in prospects['ZipCode']:
10     zips = prospects.loc[prospects['ZipCode'] == zipcode]
11     m1 = melt_data(zips)
12     m1[str(zipcode)] = m1['value']
13     m1 = m1.drop('value', axis=1)
14
15     # create train and test sets
16     train = m1.iloc[:238]
17     test = m1.iloc[238:]
18
19     # Run a grid with pdq and seasonal pdq parameters calculated above and ge
20     ans = []
21     outputs = []
22     comb_list = []
23     combs_list = []
24
25     for comb in pdq:
26         for combs in pdqs:
27             try:
28                 mod = sm.tsa.statespace.SARIMAX(train,
29                                         order=comb,
30                                         seasonal_order=combs,
31                                         enforce_stationarity=False,
32                                         enforce_invertibility=False)
33
34                 output = mod.fit()
35                 ans.append([comb, combs, output.aic])
36                 outputs.append([output.aic])
37                 comb_list.append([comb])
38                 combs_list.append([combs])
39                 AIC_df = pd.DataFrame(list(outputs, comb_list, combs_list))
40                 #print('ARIMA {} x {}12 : AIC Calculated ={}'.format(comb, co
41             except:
42                 continue
43
44     # Print the lowest AIC and it's parameters for SARIMA
45     AIC_df[0] = AIC_df[0].str.get(0)
46     best_AIC = AIC_df.copy()
47     lowest_AIC = min(best_AIC[0])
48     lowest_AIC_row = AIC_df.loc[AIC_df[0] == lowest_AIC]
49     #print(lowest_AIC_row)
50
51     # call value for comb and combs for model input
52     lowest_AIC_row[1] = lowest_AIC_row[1].str.get(0)

```

```

53     lowest_AIC_row[2] = lowest_AIC_row[2].str.get(0)
54     arima = lowest_AIC_row.iat[0,1]
55     s = lowest_AIC_row.iat[0,2]
56     print(str(m1.columns.values))
57     print(arima, s)
58
59     # Plug the optimal parameter values into a new SARIMAX model
60     ARIMA_MODEL = sm.tsa.statespace.SARIMAX(train,
61                                              order=arima,
62                                              seasonal_order=s,
63                                              enforce_stationarity=False,
64                                              enforce_invertibility=False)
65
66     # Fit the model and print results
67     output = ARIMA_MODEL.fit()
68
69     print(output.summary().tables[1])
70
71     # Call plot_diagnostics() on the results calculated above
72     output.plot_diagnostics(figsize=(15, 18))
73     plt.show()
74
75     # Get actual results in dataframe
76     price_truth = test['2016-02-01':]
77
78     # Get predictions and forecast in dataframe
79     prediction = output.get_forecast(steps=27)
80     pred = prediction.predicted_mean
81     forecast_result = output.get_forecast(steps=75)
82     forecast = forecast_result.predicted_mean
83
84     # Get confidence intervals of forecasts
85     #pred_conf = forecast.conf_int()
86
87     # Get the real and predicted values
88     prediction = pred[0:]
89     price_truth = price_truth[str(zipcode)]
90
91     # Compute the mean square error
92     mse = ((prediction - price_truth) ** 2).mean()
93     rmse = round(np.sqrt(mse))
94     rmse_list.append(rmse)
95     print('The Root Mean Squared Error of our forecasts is {}'.format(
96         round(rmse, 2)))
97     print()
98     print()
99
100    # Print list of RMSE
101    print(rmse_list)

```

executed in 8m 58s, finished 10:55:31 2022-05-03

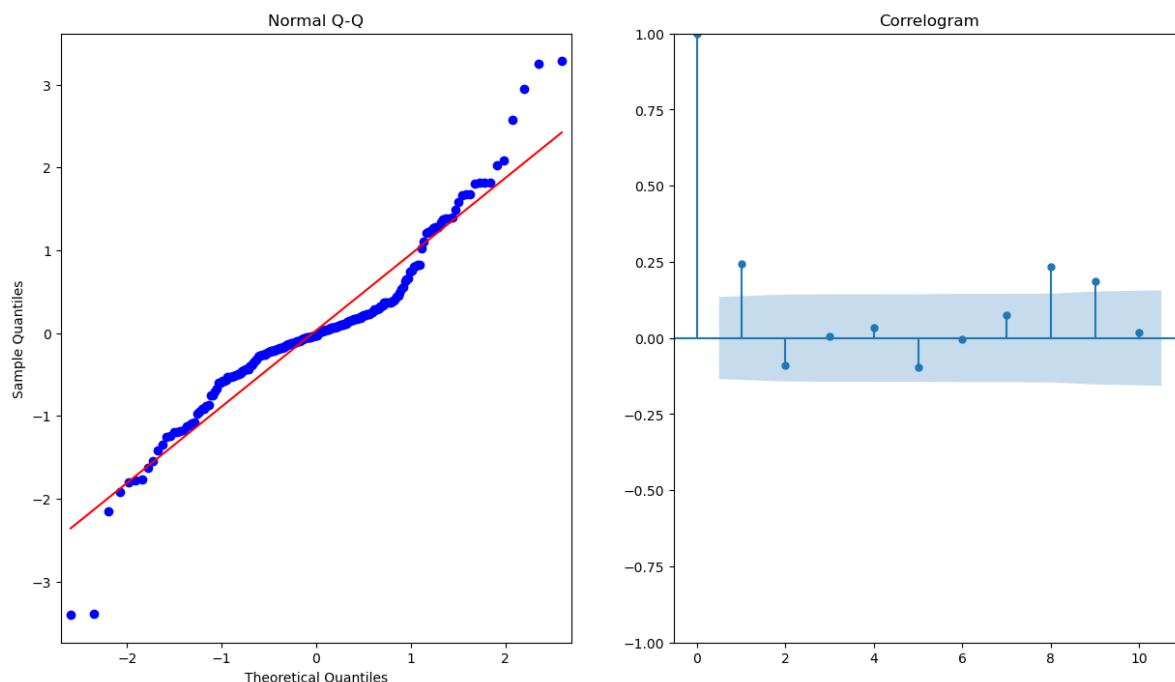
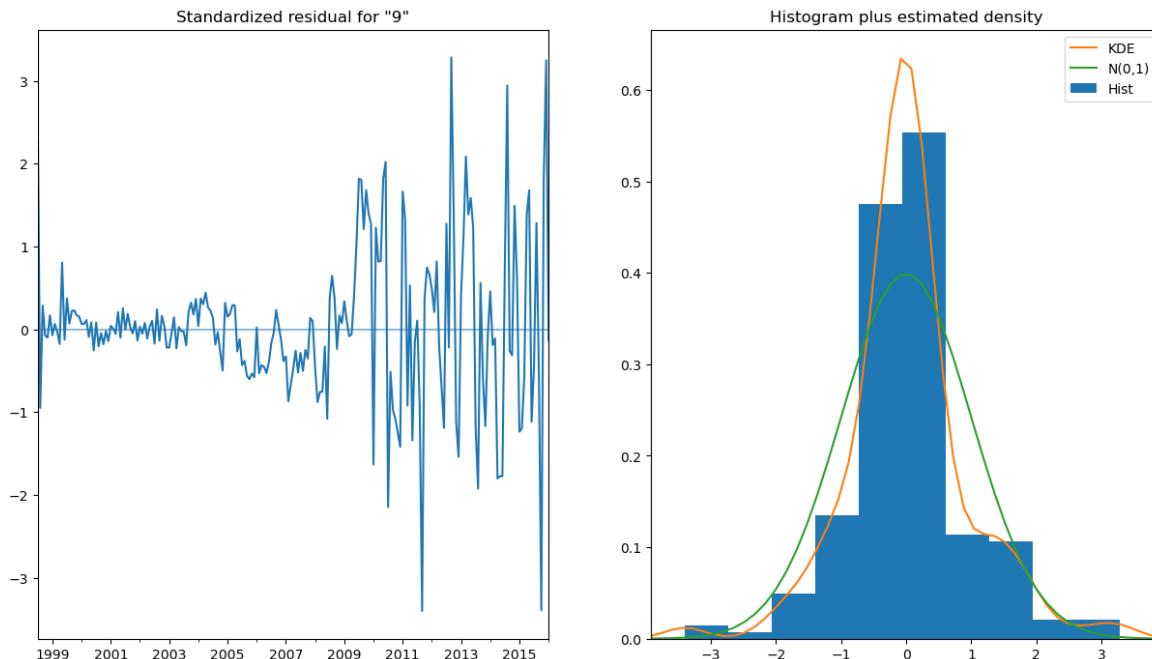
```

['94601']
(1, 1, 1) (1, 1, 1, 12)
=====

```

	coef	std err	z	P> z	[0 . 025
0.975]					

ar.L1	0.9596	0.022	44.037	0.000	0.917
1.002					
ma.L1	0.3172	0.024	13.207	0.000	0.270
0.364					
ar.S.L12	-0.6059	0.067	-9.009	0.000	-0.738
-0.474					
ma.S.L12	-0.0231	0.044	-0.520	0.603	-0.110
0.064					
sigma2	1.631e+06	1.27e+05	12.846	0.000	1.38e+06
1.88e+06					
=====					
=====					

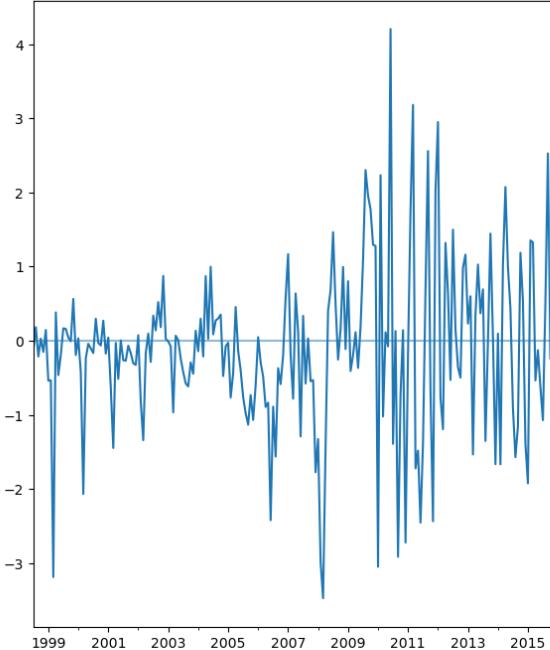


The Root Mean Squared Error of our forecasts is 58117.0

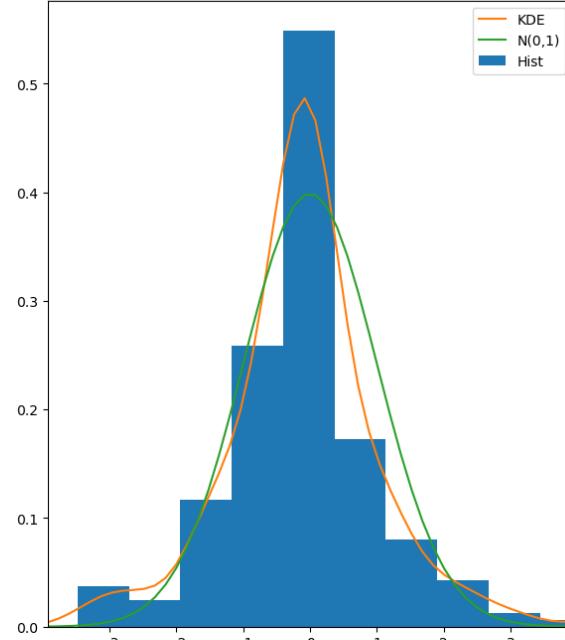
	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ar.L1	0.9712	0.016	61.356	0.000	0.940
1.002					
ma.L1	0.5069	0.039	13.013	0.000	0.431

0.583					
ar.S.L12	0.1627	0.032	5.014	0.000	0.099
0.226					
ma.S.L12	-0.8593	0.046	-18.616	0.000	-0.950
-0.769					
sigma2	4.474e+05	3.39e+04	13.214	0.000	3.81e+05
5.14e+05					

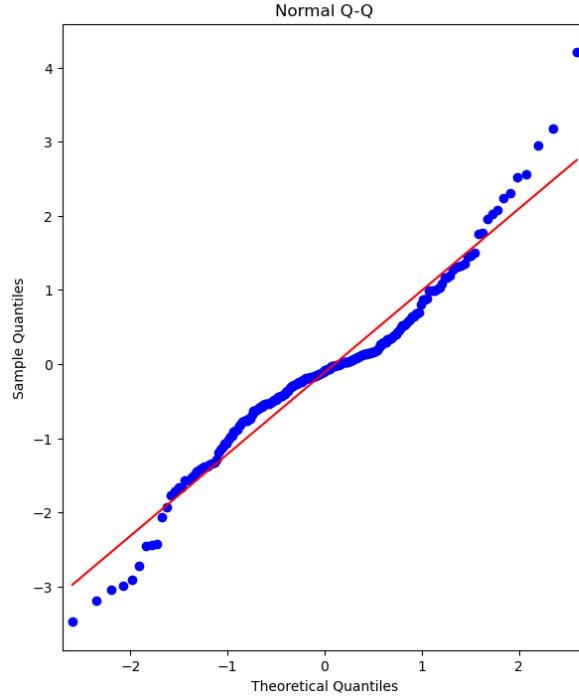
Standardized residual for "9"



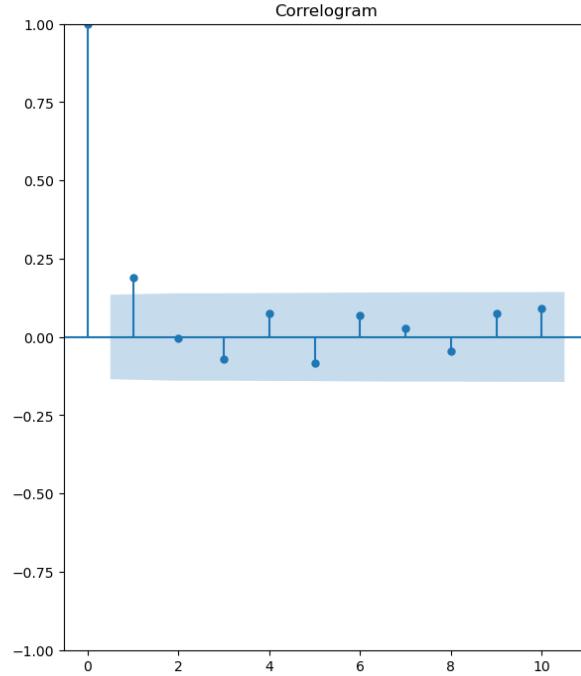
Histogram plus estimated density



Normal Q-Q

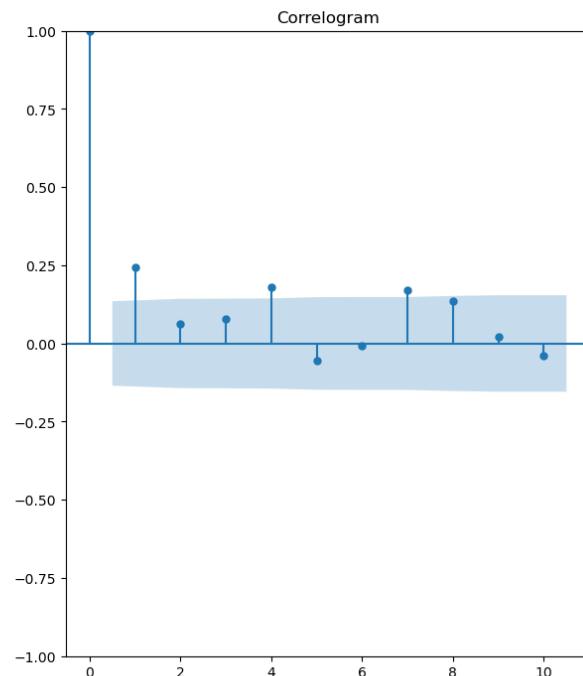
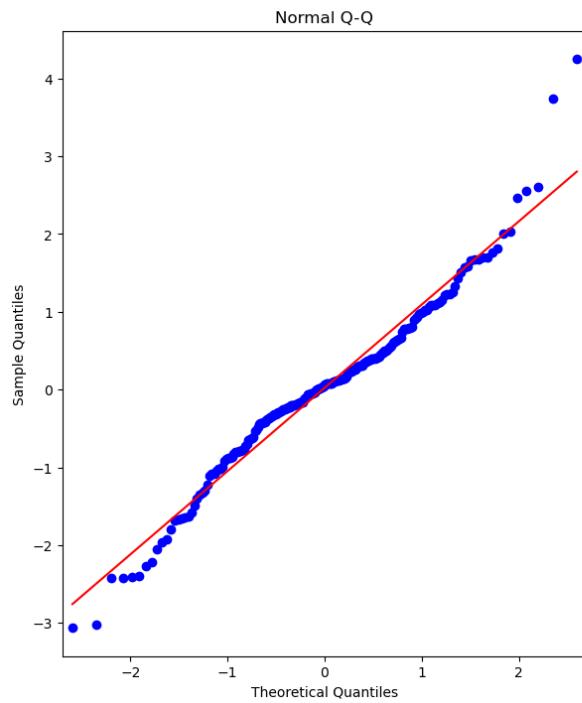
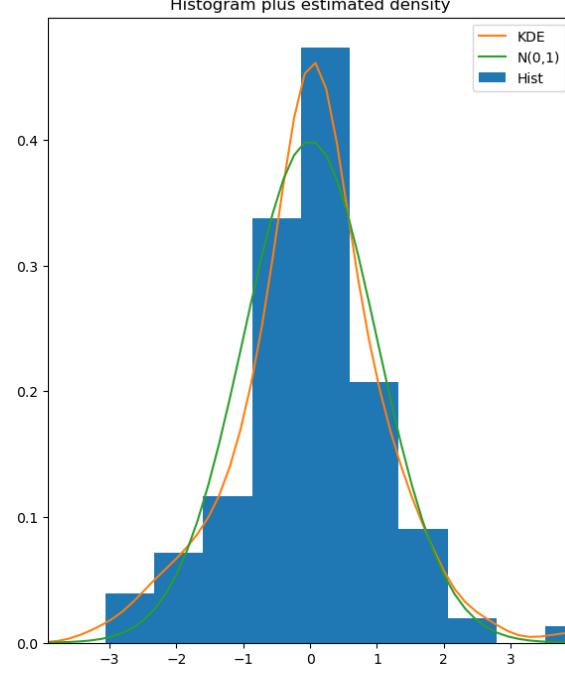
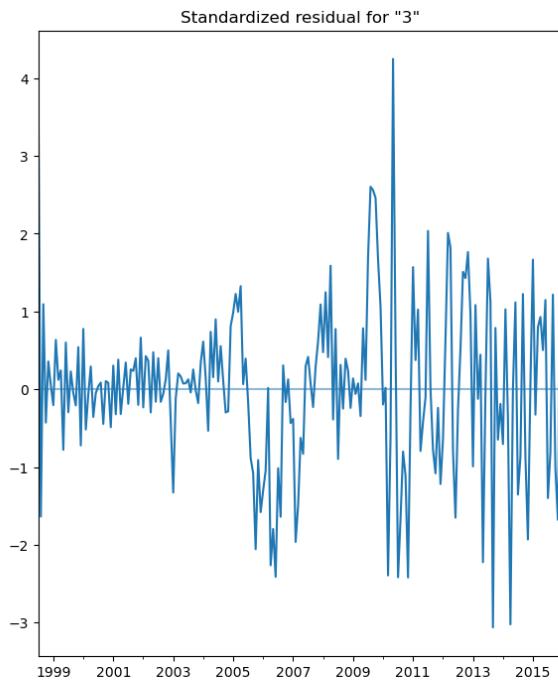


Correlogram



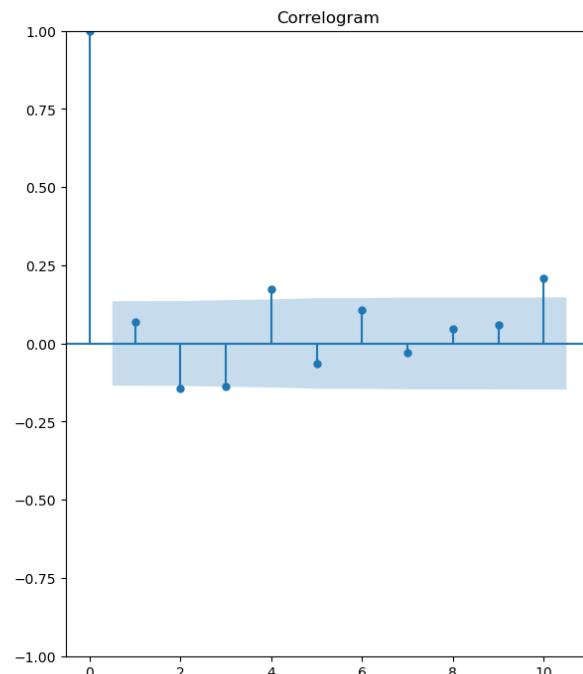
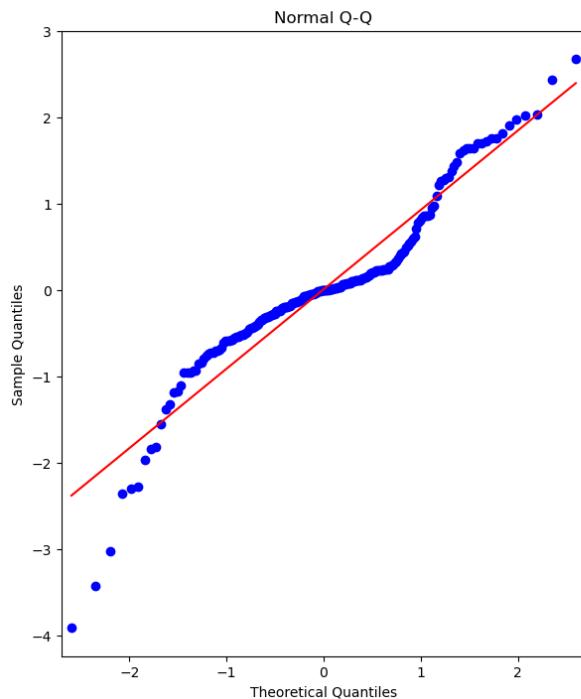
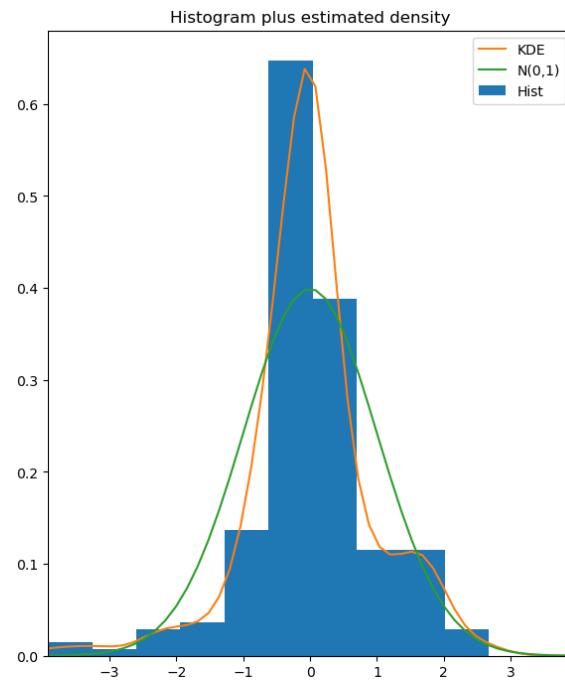
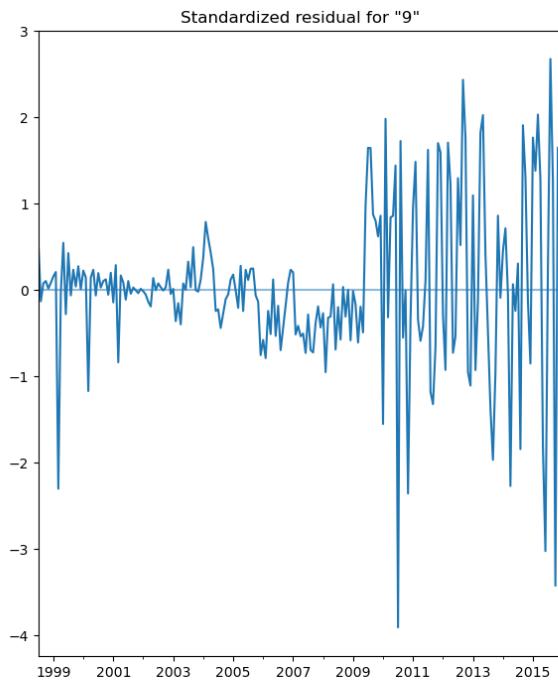
The Root Mean Squared Error of our forecasts is 9986.0

	coef	std err	z	P> z	[0.025
0.975]					
-----	-----	-----	-----	-----	-----
ar.L1	0.9778	0.014	68.825	0.000	0.950
1.006					
ma.L1	0.5459	0.024	22.808	0.000	0.499
0.593					
ar.S.L12	-0.6195	0.063	-9.809	0.000	-0.743
-0.496					
ma.S.L12	19.6087	27.280	0.719	0.472	-33.859
73.076					
sigma2	403.6864	1126.099	0.358	0.720	-1803.427
2610.799					
=====	=====	=====	=====	=====	=====
=====	=====	=====	=====	=====	=====



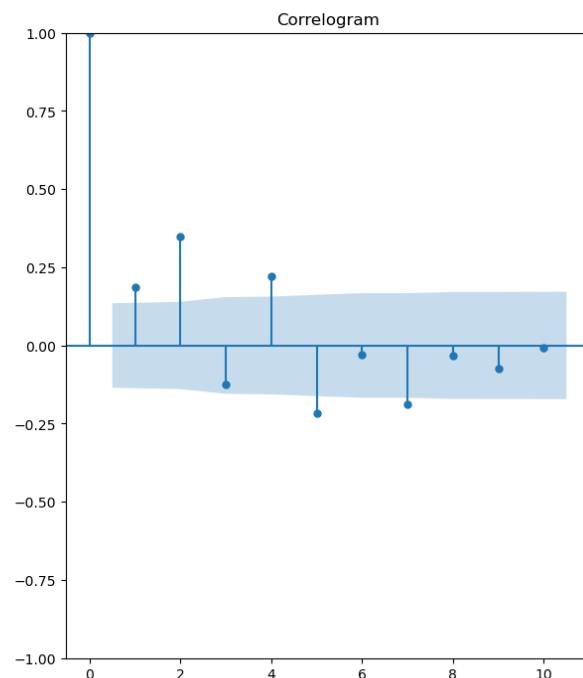
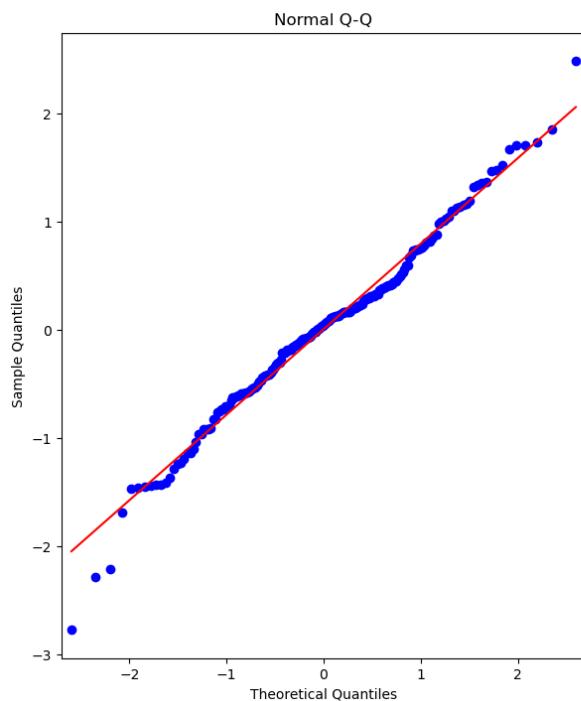
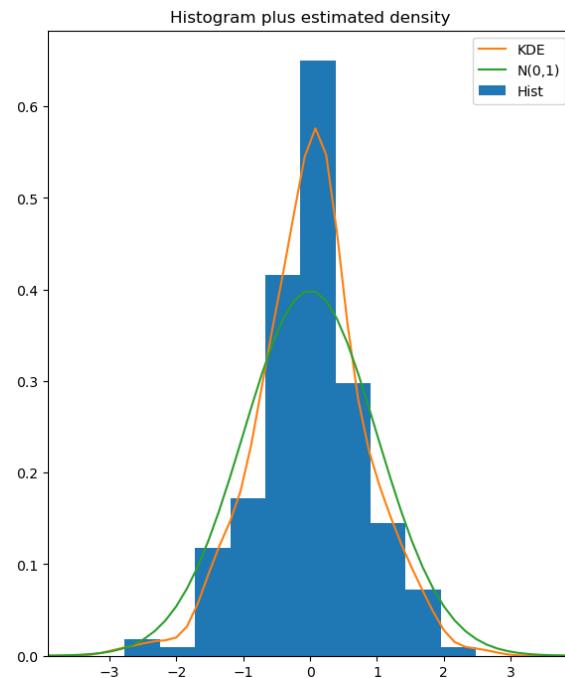
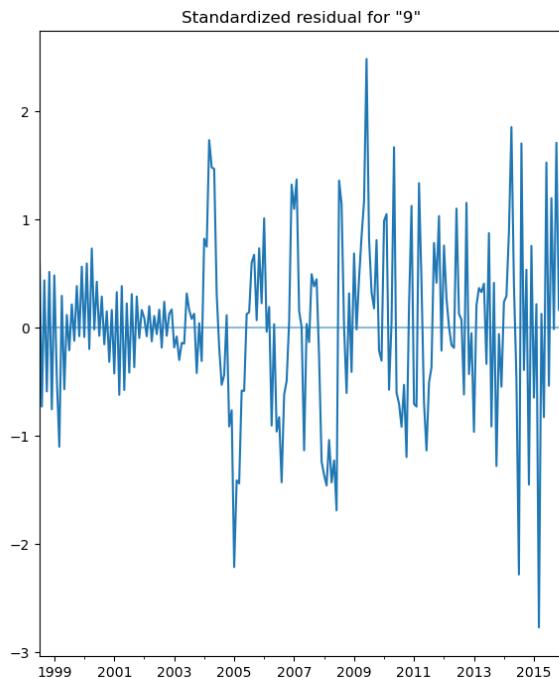
The Root Mean Squared Error of our forecasts is 2775.0

```
['94606']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [0.025
0.975]
-----
ar.L1        0.8734      0.034      25.839      0.000      0.807
0.940
ma.L1        0.6692      0.047      14.135      0.000      0.576
0.762
ar.S.L12     0.2098      0.083       2.515      0.012      0.046
0.373
ma.S.L12    -0.5293      0.066      -8.053      0.000     -0.658
-0.400
sigma2      2.408e+06   2.24e+05      10.770      0.000     1.97e+06
2.85e+06
=====
```



The Root Mean Squared Error of our forecasts is 20071.0

['94804']	(1, 1, 1) (1, 1, 1, 12)	=====	=====	=====	=====
0.975]	-----	-----	-----	-----	-----
ar.L1 1.001	0.9260	0.038	24.364	0.000	0.852
ma.L1 1.187	1.1079	0.040	27.458	0.000	1.029
ar.S.L12 0.182	0.1276	0.028	4.592	0.000	0.073
ma.S.L12 -0.135	-0.2794	0.074	-3.784	0.000	-0.424
sigma2 3.09e+06	2.364e+06	3.69e+05	6.404	0.000	1.64e+06

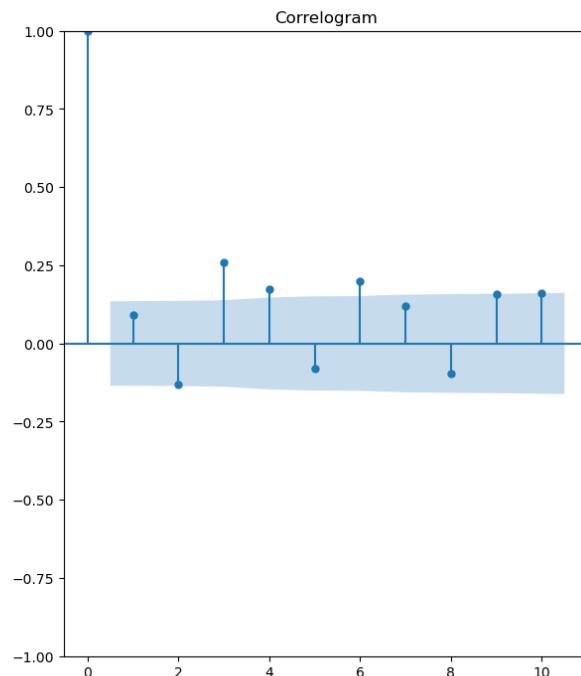
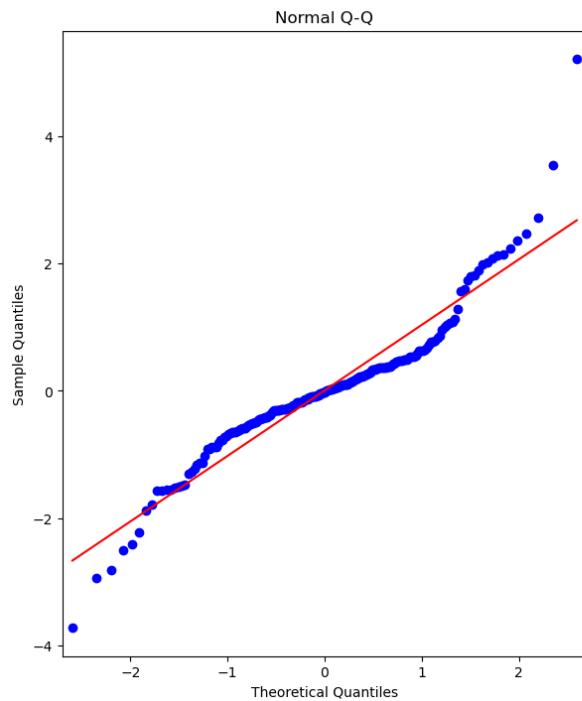
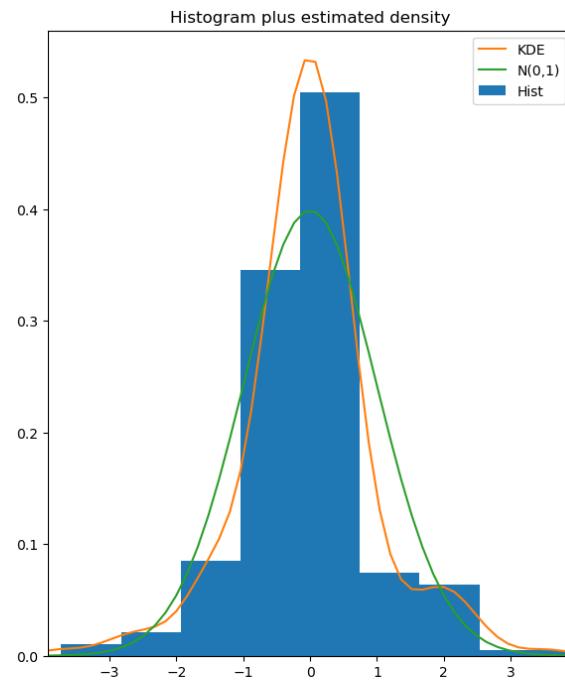
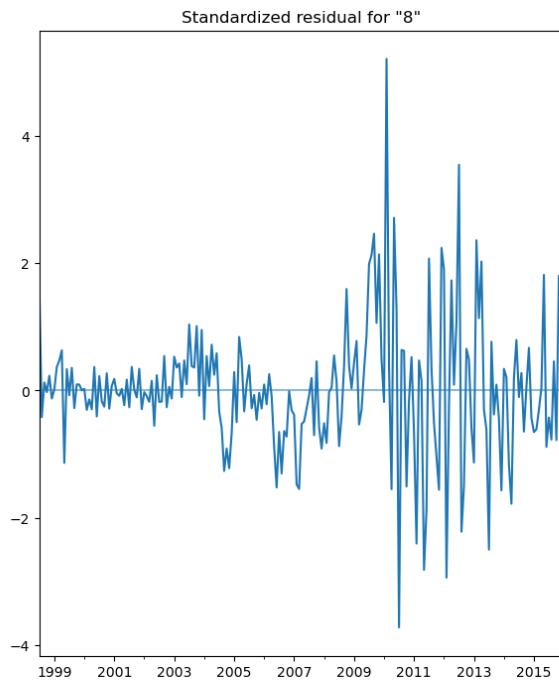


The Root Mean Squared Error of our forecasts is 44237.0

```
[ '89104' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
=====
```

	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.9780	0.013	74.511	0.000	0.952
1.004					
ma.L1	0.2738	0.037	7.446	0.000	0.202
0.346					
ar.S.L12	-0.5424	0.040	-13.727	0.000	-0.620
-0.465					
ma.S.L12	16.2301	8.367	1.940	0.052	-0.170
32.630					
sigma2	903.7866	935.906	0.966	0.334	-930.555
2738.128					
=====					
=====					

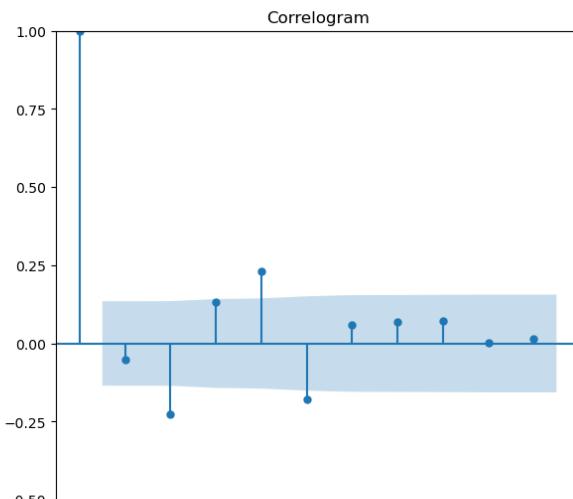
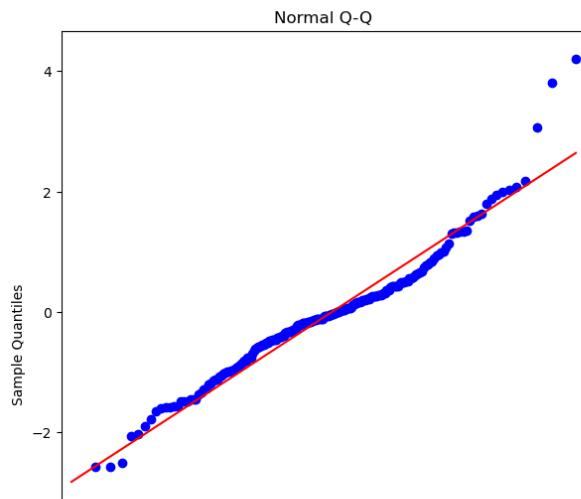
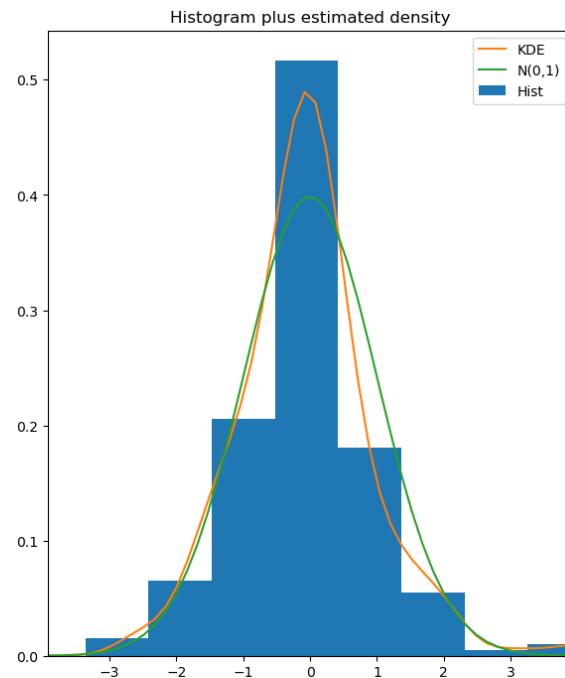
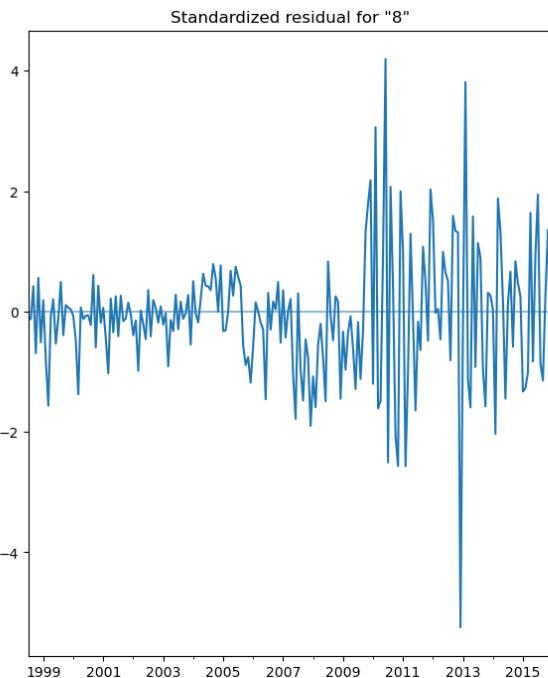


The Root Mean Squared Error of our forecasts is 12709.0

```
['85008']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
```

coef	std err	z	P> z	[0.025
0.975]				

```
-----  
ar.L1      0.9468    0.023    41.557    0.000    0.902  
0.991  
ma.L1      0.5435    0.039    14.114    0.000    0.468  
0.619  
ar.S.L12    0.0930    0.038    2.423     0.015    0.018  
0.168  
ma.S.L12    -0.9883   0.411    -2.403    0.016    -1.794  
-0.182  
sigma2     2.325e+05  9.64e+04  2.412     0.016    4.35e+04  
4.21e+05  
=====
```



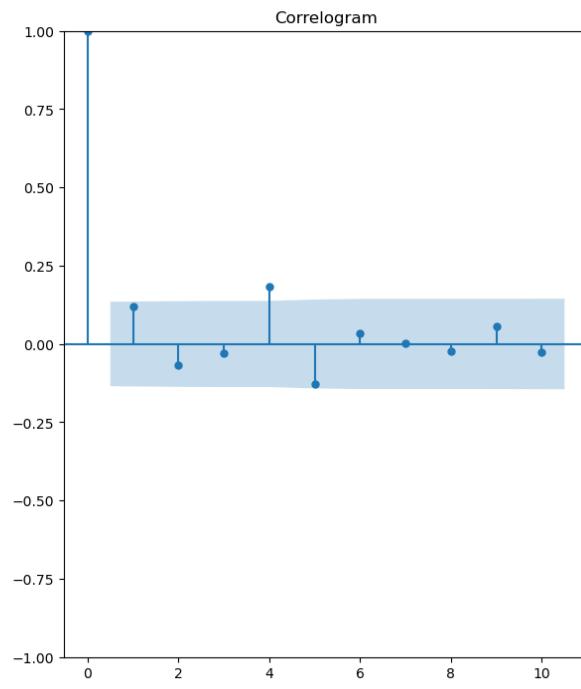
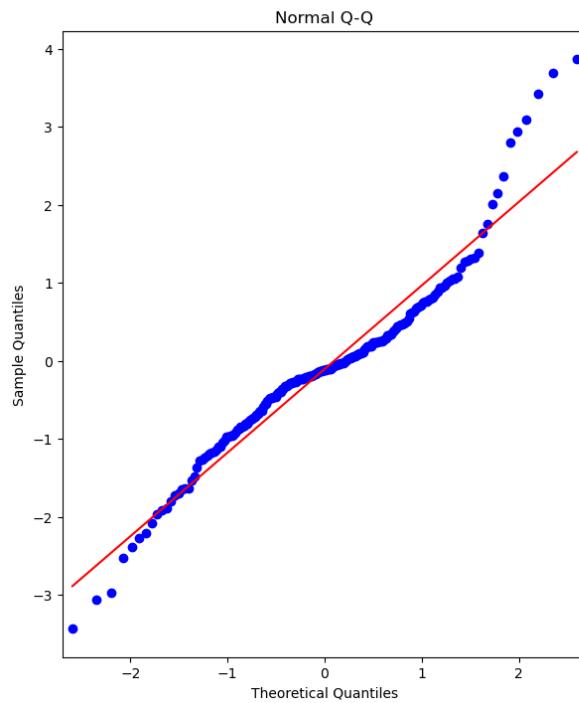
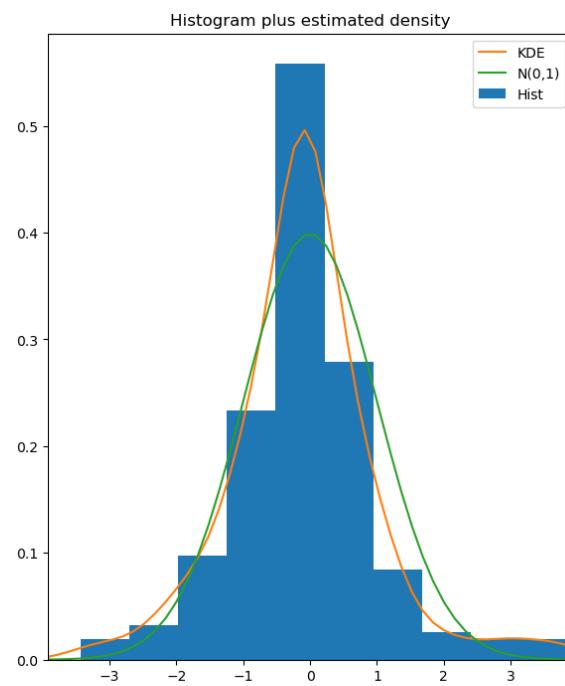
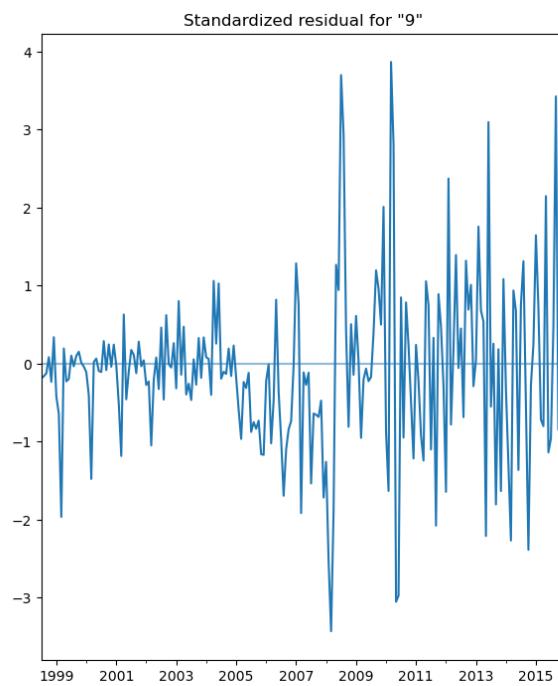
The Root Mean Squared Error of our forecasts is 12072.0

```
[ '95824' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
          coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.9656      0.012     82.031      0.000      0.943
0.989
ma.L1      0.6079      0.040     15.376      0.000      0.530
0.685
ar.S.L12   0.1057      0.030      3.530      0.000      0.047
0.164
ma.S.L12  -0.9670      0.112     -8.656      0.000     -1.186
-0.748
```

σ^2	2.552e+05	3.45e+04	7.398	0.000	1.88e+05
------------	-----------	----------	-------	-------	----------

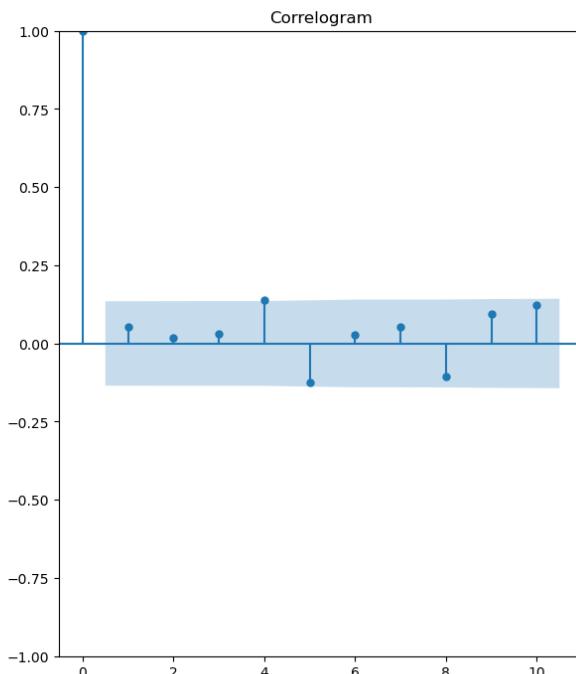
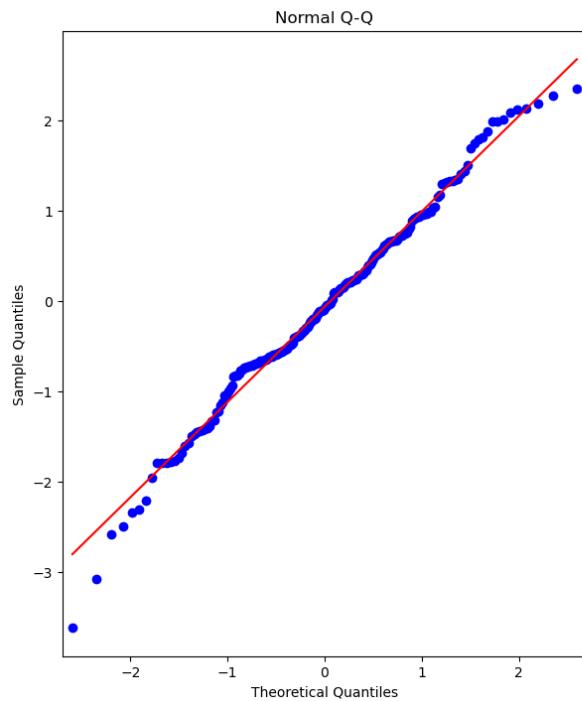
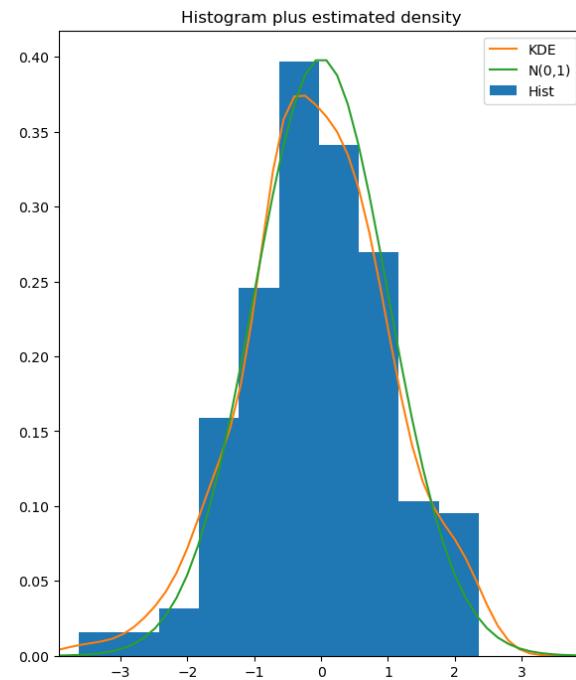
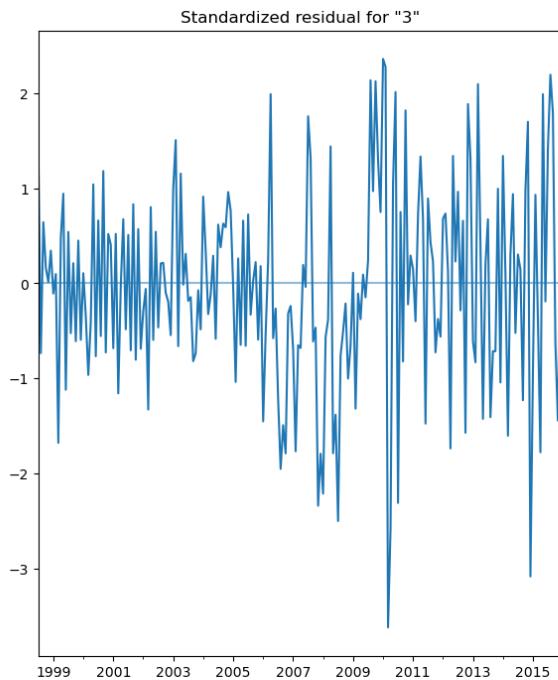
3.23e+05	=====
----------	-------

=====



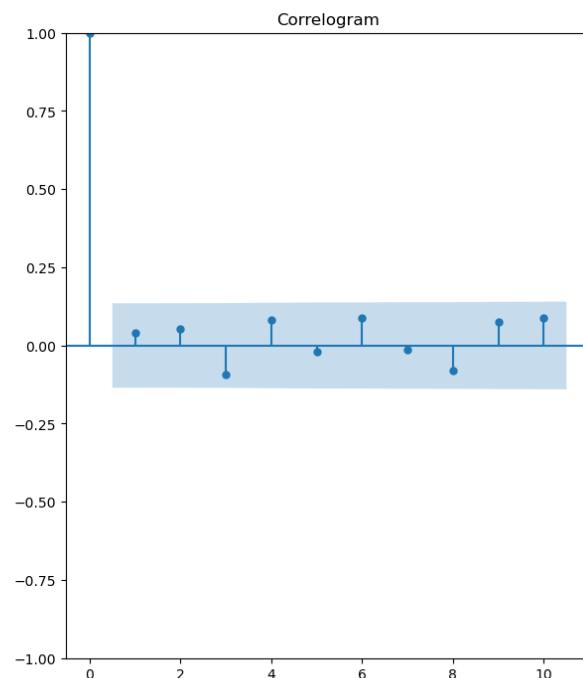
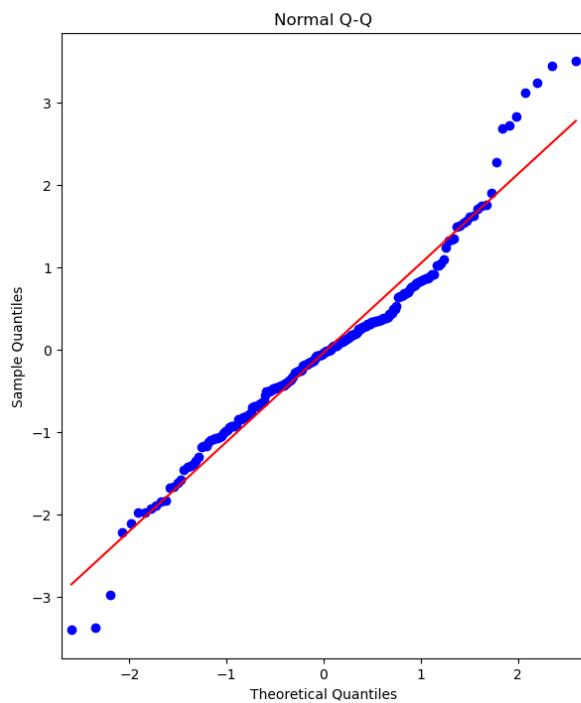
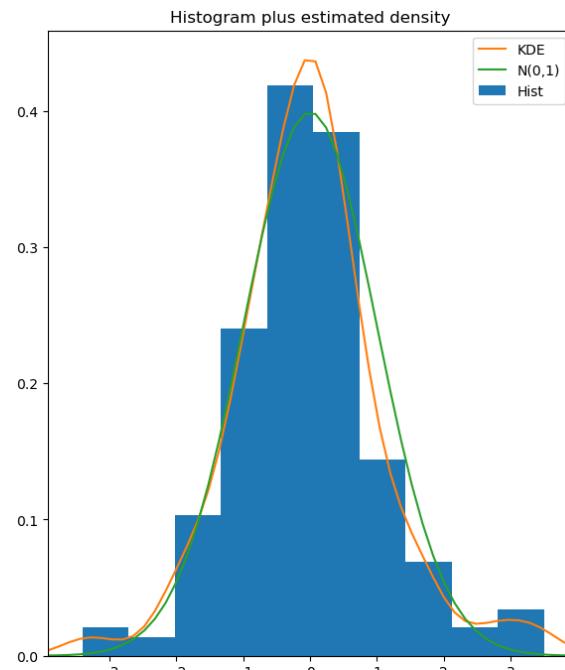
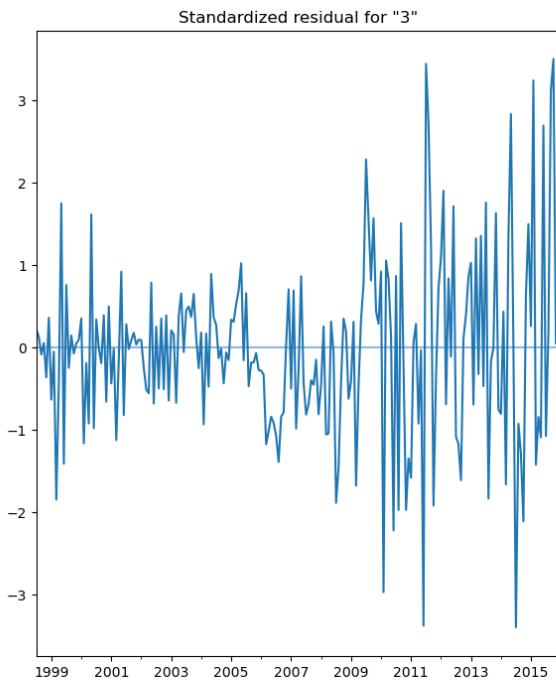
The Root Mean Squared Error of our forecasts is 5100.0

```
['33404']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [0.025
0.975]
-----
ar.L1      0.9743      0.015      65.986      0.000      0.945
1.003
ma.L1      0.5756      0.045      12.882      0.000      0.488
0.663
ma.S.L12   -0.9976     1.736      -0.575      0.566      -4.401
2.406
sigma2     8.182e+04    1.4e+05      0.583      0.560      -1.93e+05
3.57e+05
=====
=====
```



The Root Mean Squared Error of our forecasts is 20612.0

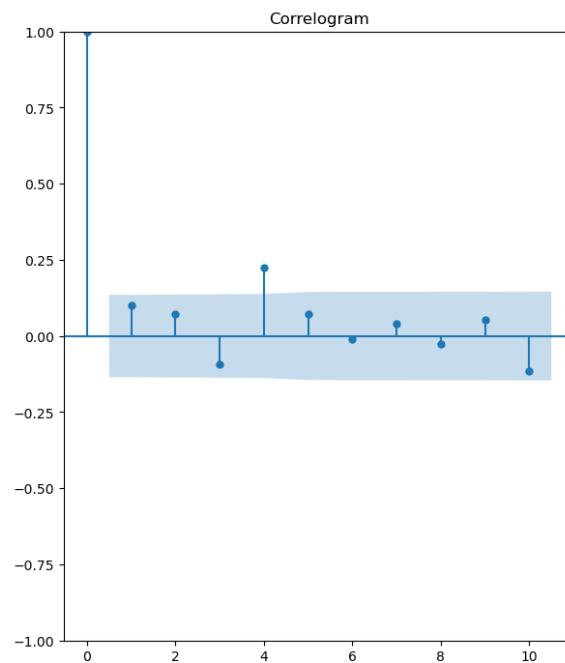
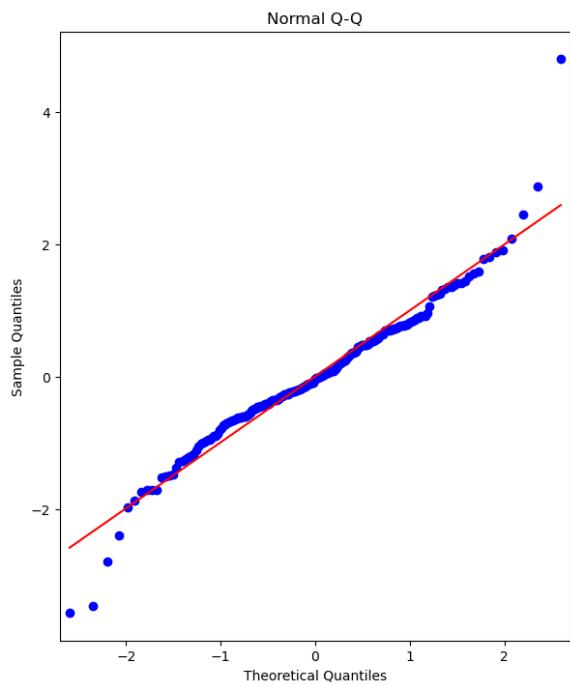
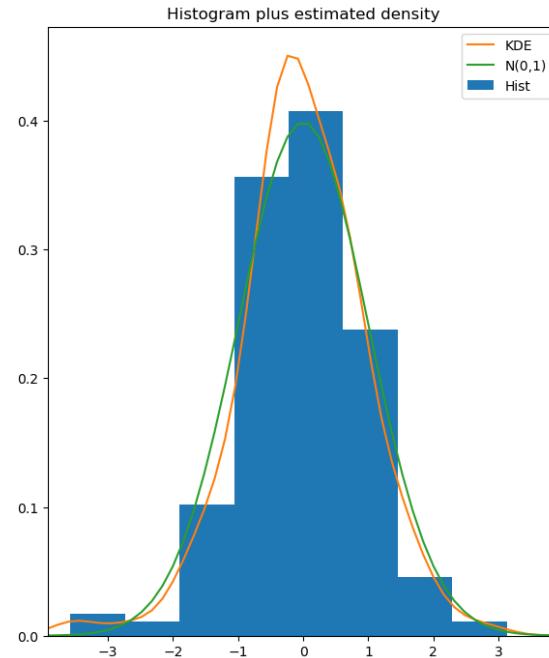
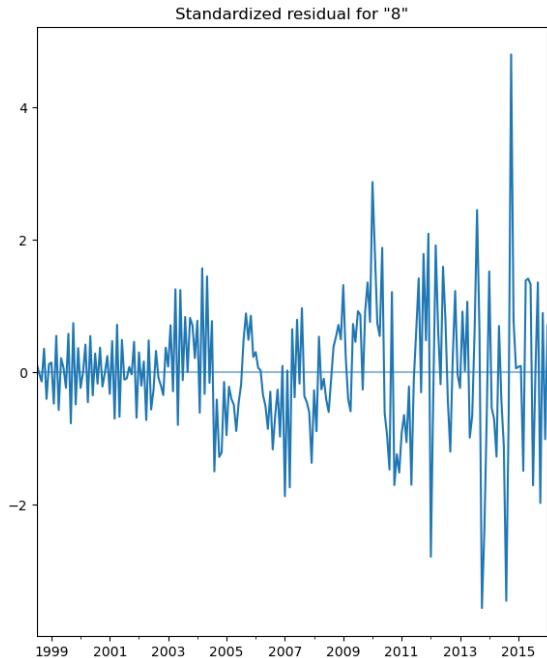
```
['33705']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.9582      0.020      48.155      0.000      0.919
0.997
ma.L1      0.5818      0.035      16.674      0.000      0.513
0.650
ma.S.L12   -0.9439      0.065     -14.448      0.000     -1.072
-0.816
sigma2     1.483e+05    1.21e+04     12.218      0.000     1.25e+05
1.72e+05
=====
=====
```



The Root Mean Squared Error of our forecasts is 19611.0

	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.9414	0.020	47.806	0.000	0.903
0.980					
ma.L1	0.6926	0.047	14.810	0.000	0.601
0.784					
ma.S.L12	32.9930	60.954	0.541	0.588	-86.474
152.460					
sigma2	219.4013	820.043	0.268	0.789	-1387.853
1826.655					
=====					
=====					

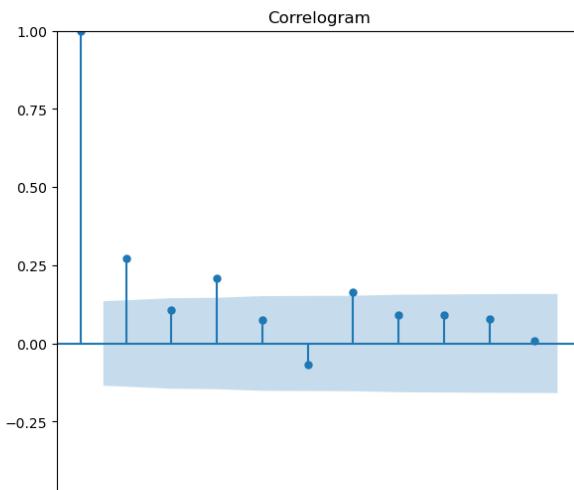
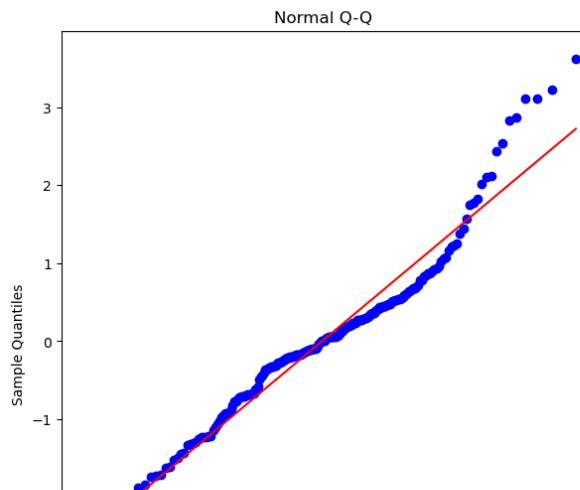
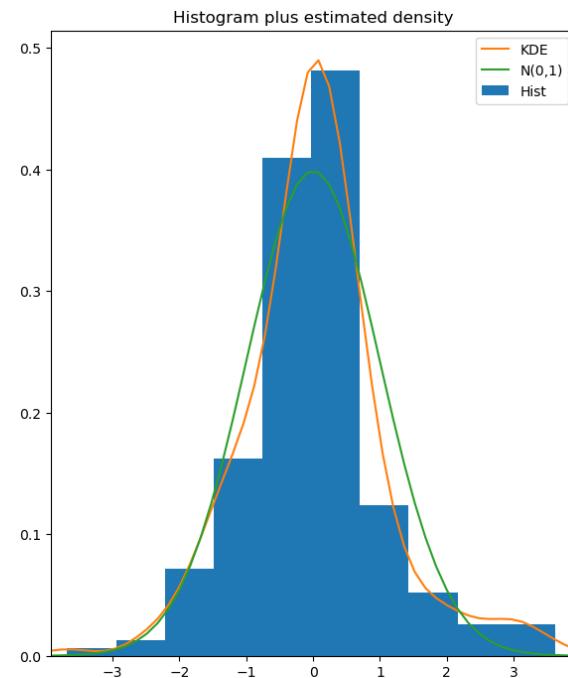
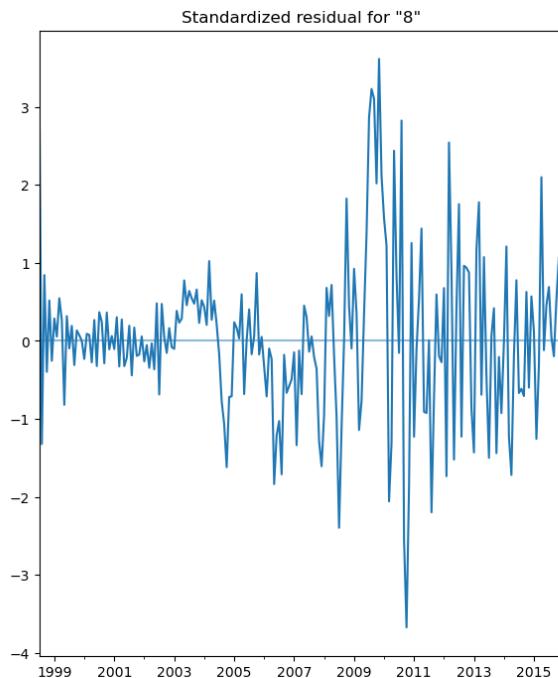


The Root Mean Squared Error of our forecasts is 5250.0

```
[ '89115' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
```

coef	std err	z	P> z	[0.025
0.975]				

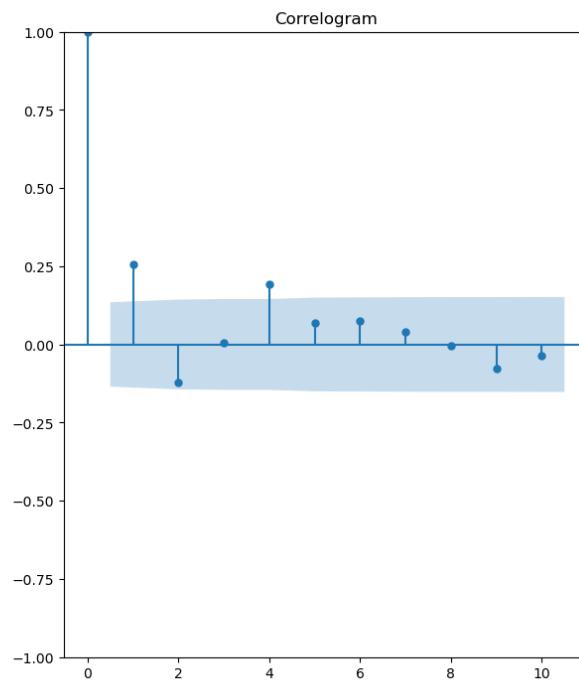
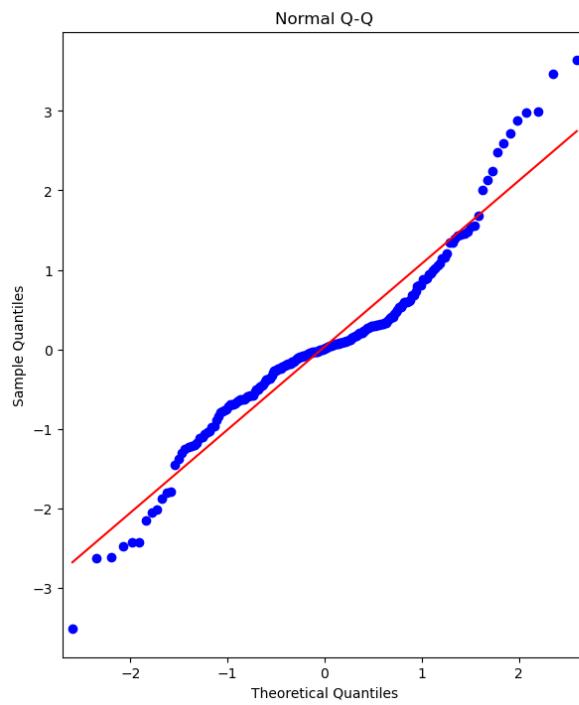
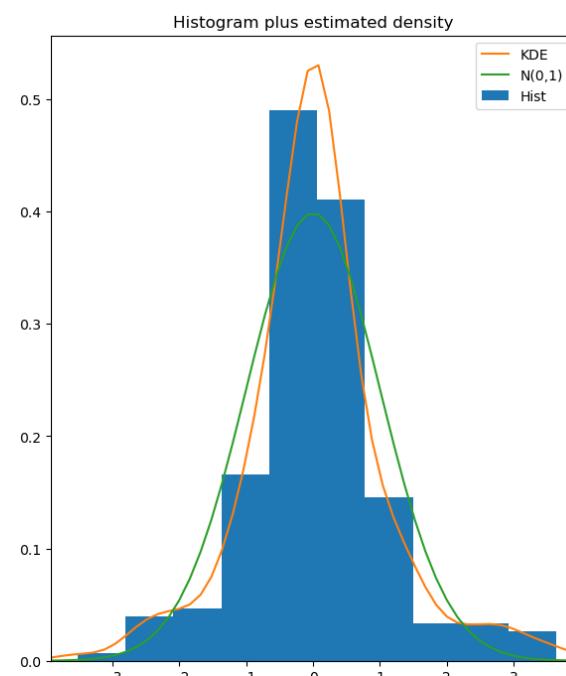
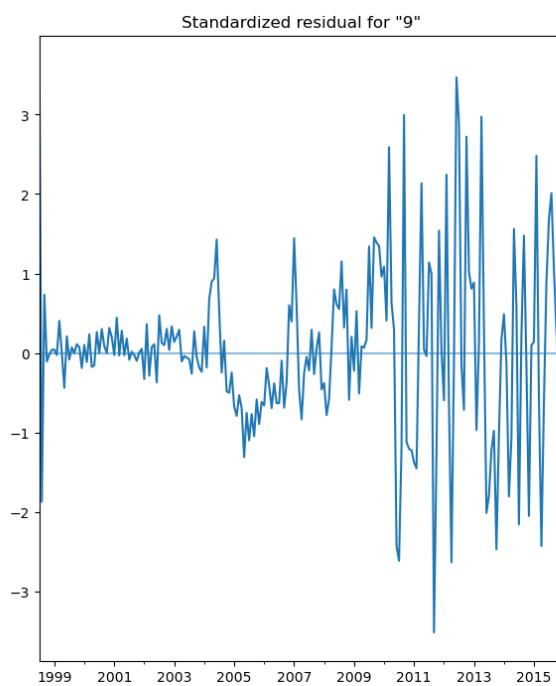
```
-----  
ar.L1      0.9762    0.012    78.954    0.000    0.952  
1.000  
ma.L1      0.3515    0.019    18.566    0.000    0.314  
0.389  
ar.S.L12   -0.4913   0.065   -7.536    0.000   -0.619  
-0.364  
ma.S.L12   19.4042   17.505   1.108    0.268   -14.905  
53.714  
sigma2     445.5423  807.393   0.552    0.581  -1136.920  
2028.004  
=====
```



The Root Mean Squared Error of our forecasts is 3654.0

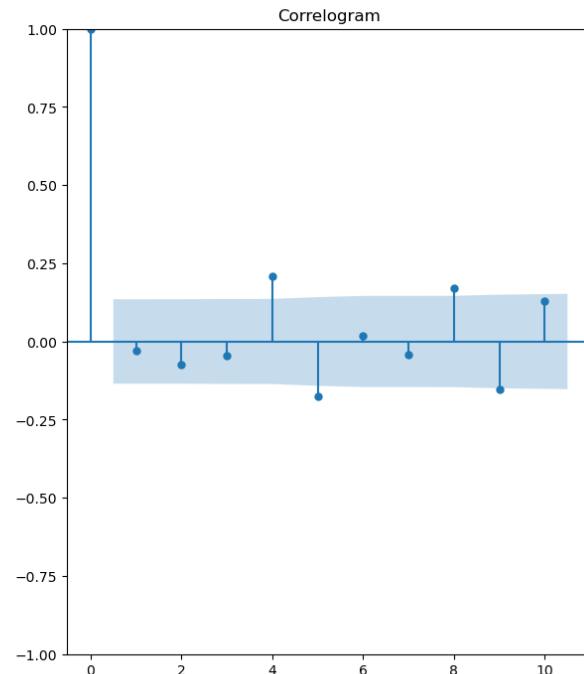
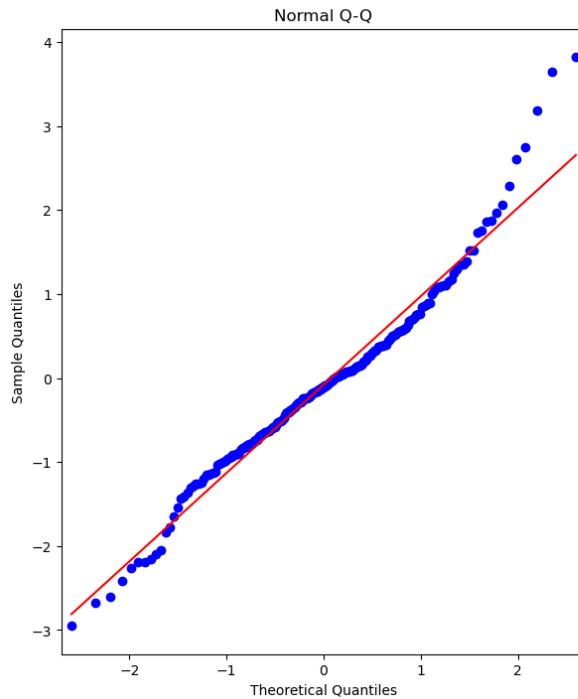
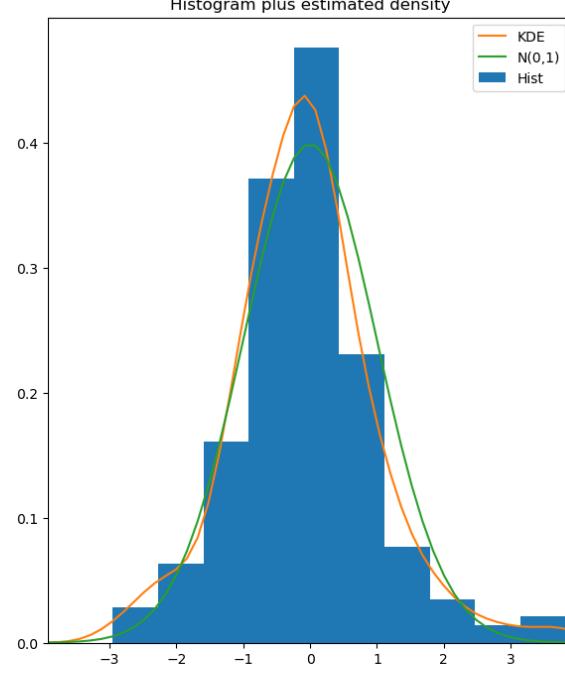
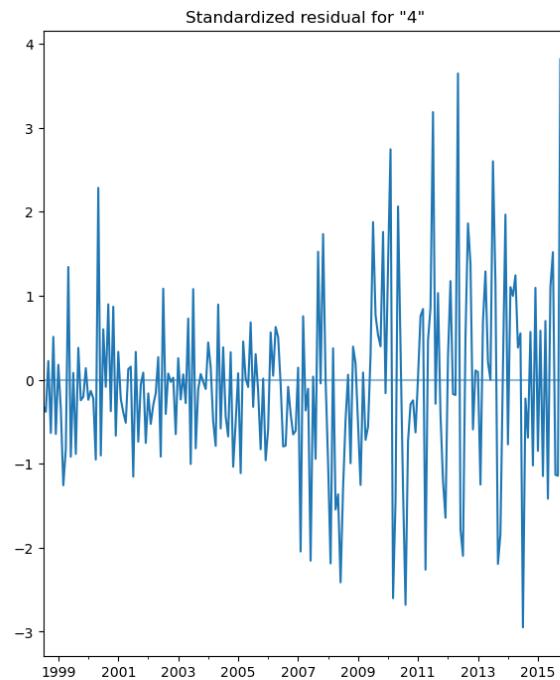
```
[ '95820' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
          coef      std err       z     P>|z|      [ 0.025
0.975]
-----
-----
ar.L1      0.9642      0.017    55.430      0.000      0.930
0.998
ma.L1      0.4264      0.019    22.201      0.000      0.389
0.464
ar.S.L12   -0.5107      0.073   -6.988      0.000     -0.654
-0.367
ma.S.L12   21.1874    38.151      0.555      0.579    -53.588
95.963
```

sigma2	650.8028	2366.818	0.275	0.783	-3988.075
5289.681					



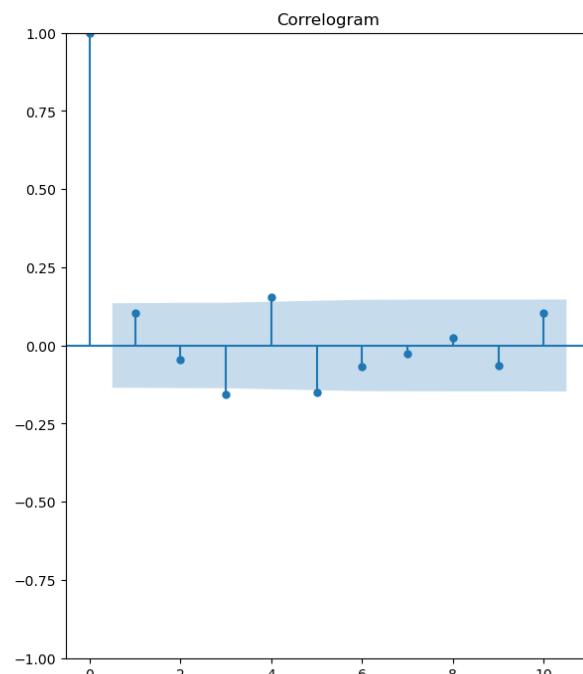
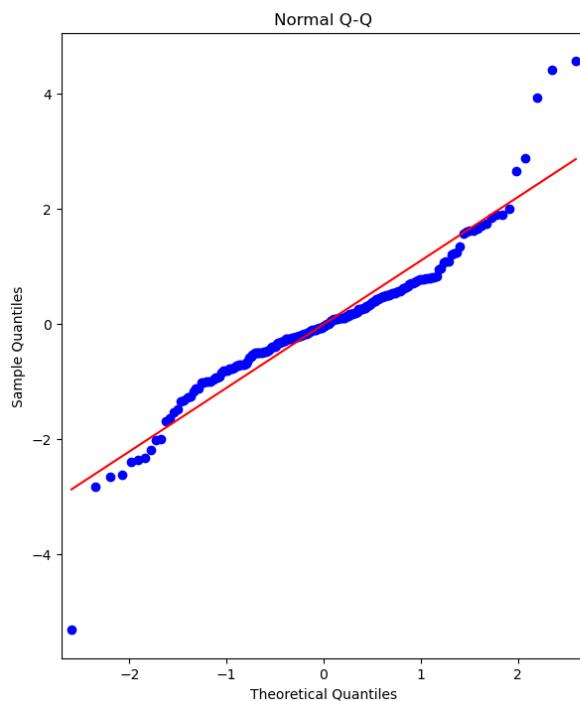
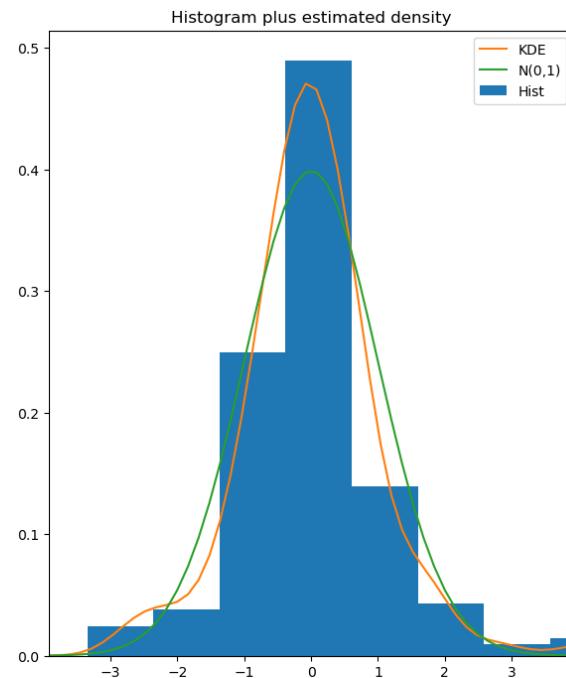
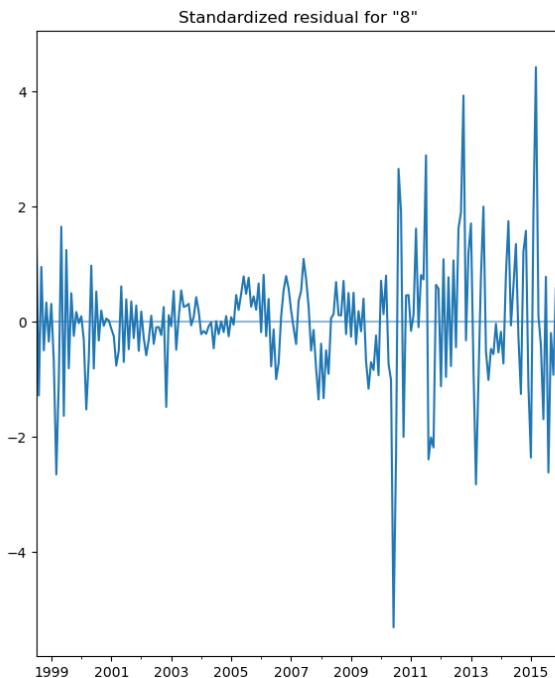
The Root Mean Squared Error of our forecasts is 19365.0

```
['48240']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [0.025
0.975]
-----
ar.L1      0.9596      0.021      45.757      0.000      0.918
1.001
ma.L1      0.4983      0.053      9.352      0.000      0.394
0.603
ma.S.L12   -0.9980     1.989     -0.502      0.616     -4.896
2.900
sigma2     4.774e+04    9.38e+04      0.509      0.611     -1.36e+05
2.31e+05
=====
```



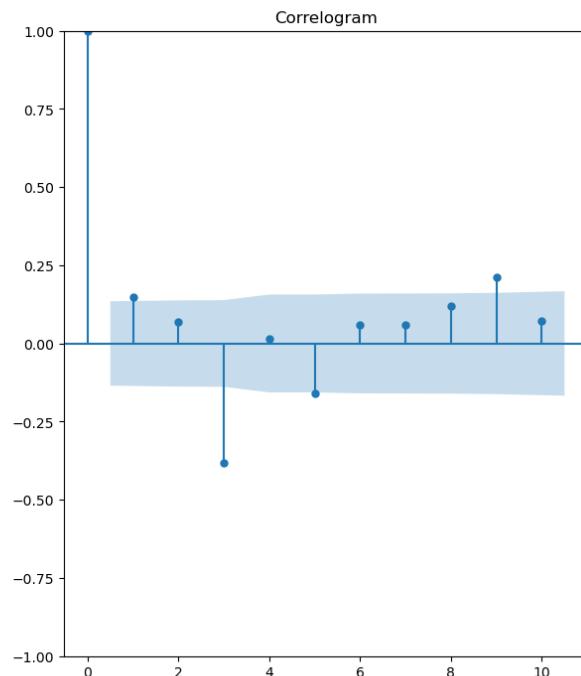
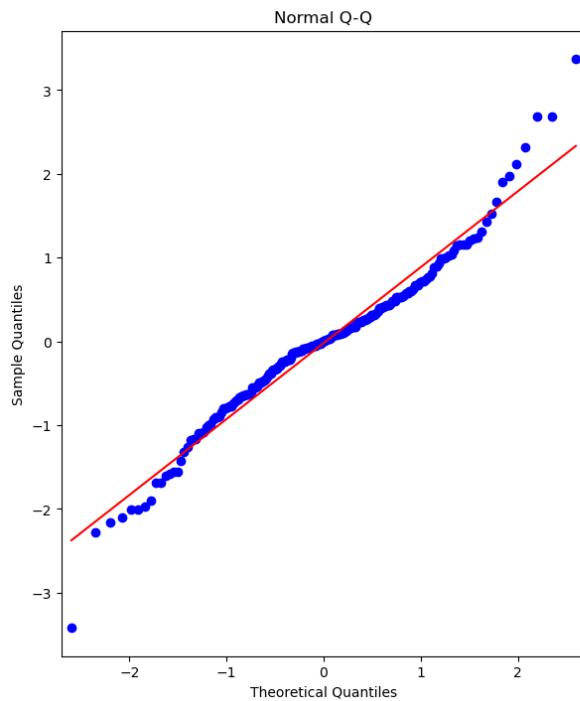
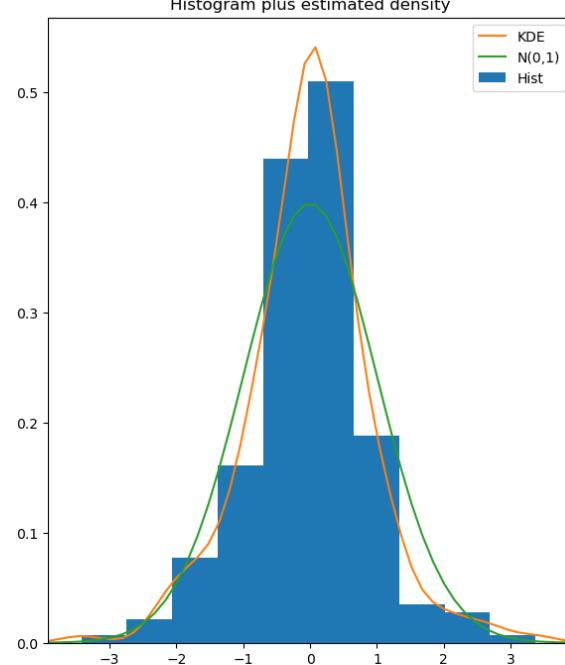
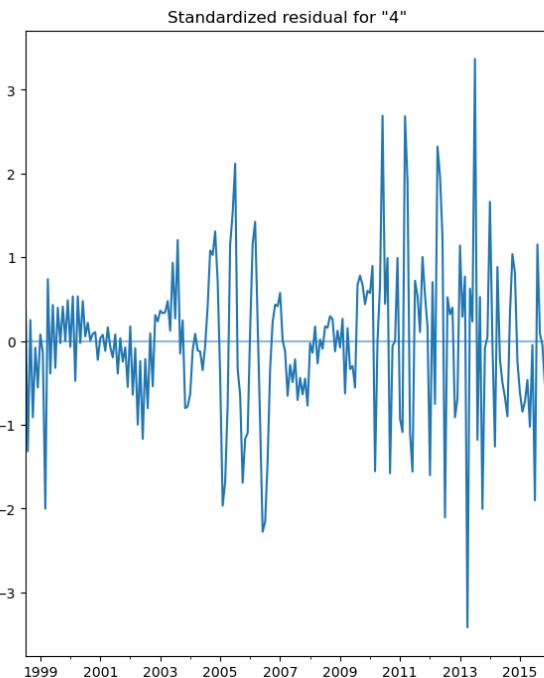
The Root Mean Squared Error of our forecasts is 2626.0

```
['84101']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
          coef    std err        z     P>|z|      [0.025
0.975]
-----
ar.L1      0.7415    0.029    25.721    0.000    0.685
0.798
ma.L1      0.7009    0.033    21.038    0.000    0.636
0.766
ma.S.L12   -0.5911   0.037   -16.026    0.000   -0.663
-0.519
sigma2     4.191e+05  2.19e+04   19.124    0.000   3.76e+05
4.62e+05
=====
=====
```



The Root Mean Squared Error of our forecasts is 24559.0

```
[ '48009' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.8114      0.047     17.207      0.000      0.719
0.904
ma.L1      1.1354      0.042     26.998      0.000      1.053
1.218
ar.S.L12    0.0805      0.013      6.419      0.000      0.056
0.105
ma.S.L12   -0.2772      0.031     -9.055      0.000     -0.337
-0.217
sigma2     1.275e+06    1.6e+05     7.959      0.000     9.61e+05
1.59e+06
=====
=====
```



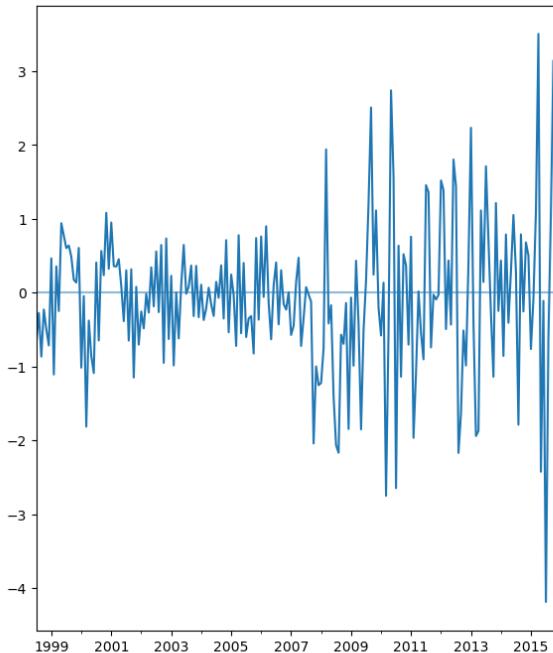
The Root Mean Squared Error of our forecasts is 22105.0

```
[ '66104' ]
(1, 1, 1) (0, 1, 1, 12)
=====
=====
      coef      std err          z      P>|z|      [ 0.025
0.975 ]
```

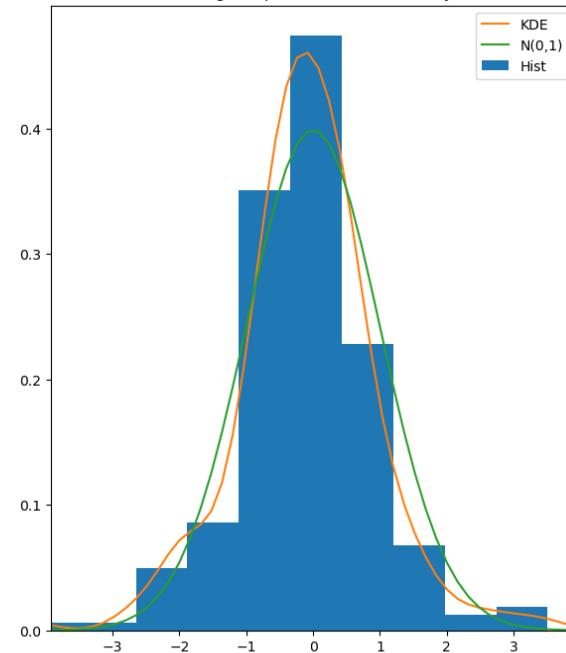
ar.L1	0.8911	0.029	30.434	0.000	0.834
0.948					
ma.L1	0.6211	0.050	12.362	0.000	0.523
0.720					
ma.S.L12	-0.9982	2.608	-0.383	0.702	-6.110
4.114					
sigma2	5.87e+04	1.53e+05	0.385	0.701	-2.41e+05
3.58e+05					

=====

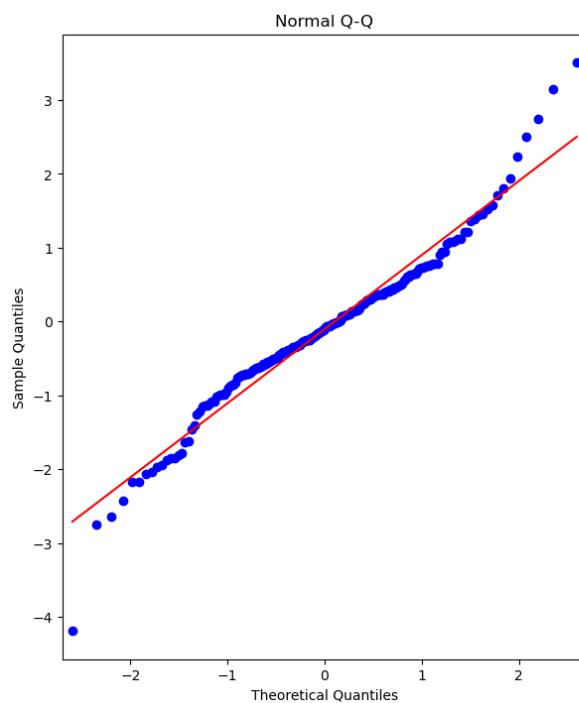
Standardized residual for "6"



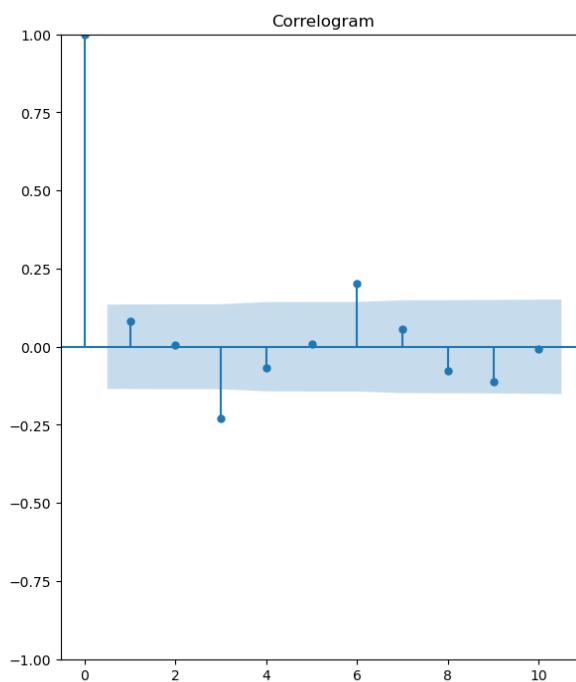
Histogram plus estimated density



Normal Q-Q



Correlogram



The Root Mean Squared Error of our forecasts is 17671.0

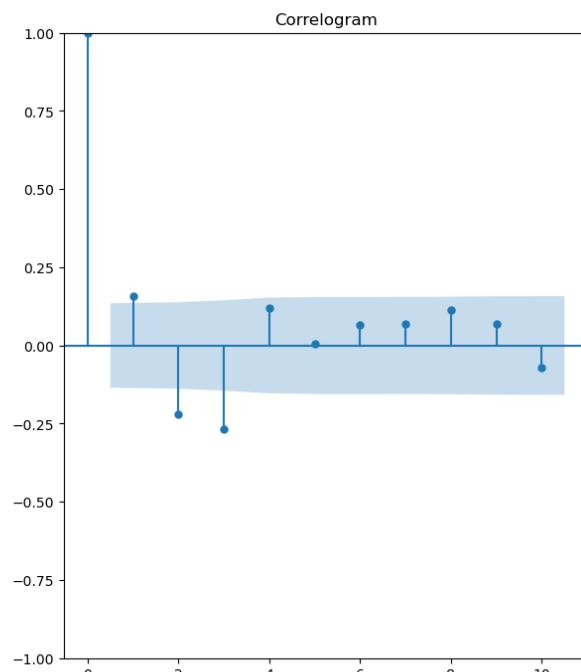
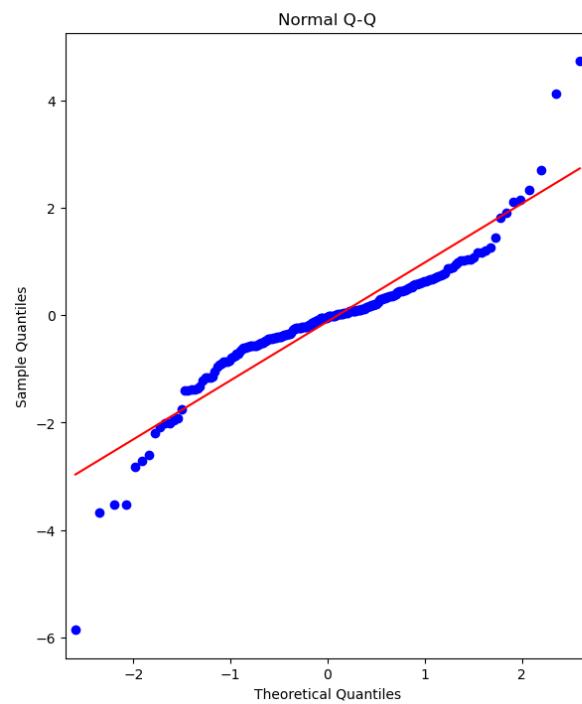
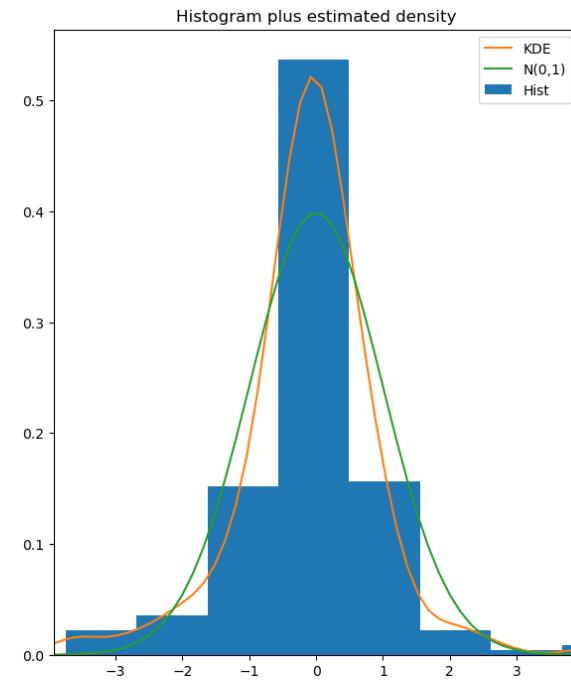
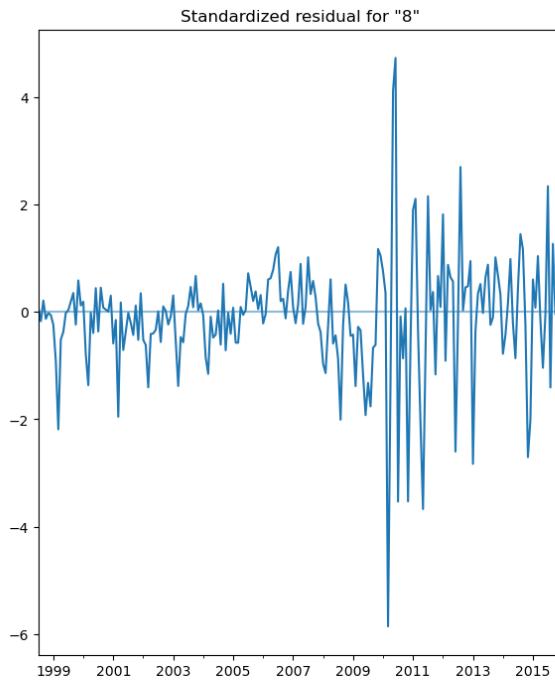
	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.8445	0.033	25.855	0.000	0.781
0.909					
ma.L1	0.3847	0.041	9.321	0.000	0.304
0.466					
ar.S.L12	0.0900	0.016	5.778	0.000	0.059
0.121					
ma.S.L12	-0.9551	0.118	-8.128	0.000	-1.185
-0.725					
sigma2	1.611e+05	1.27e+04	12.648	0.000	1.36e+05

1.86e+05

=====

=====



The Root Mean Squared Error of our forecasts is 31924.0

```
['87506']
(1, 1, 1) (1, 1, 1, 12)
```

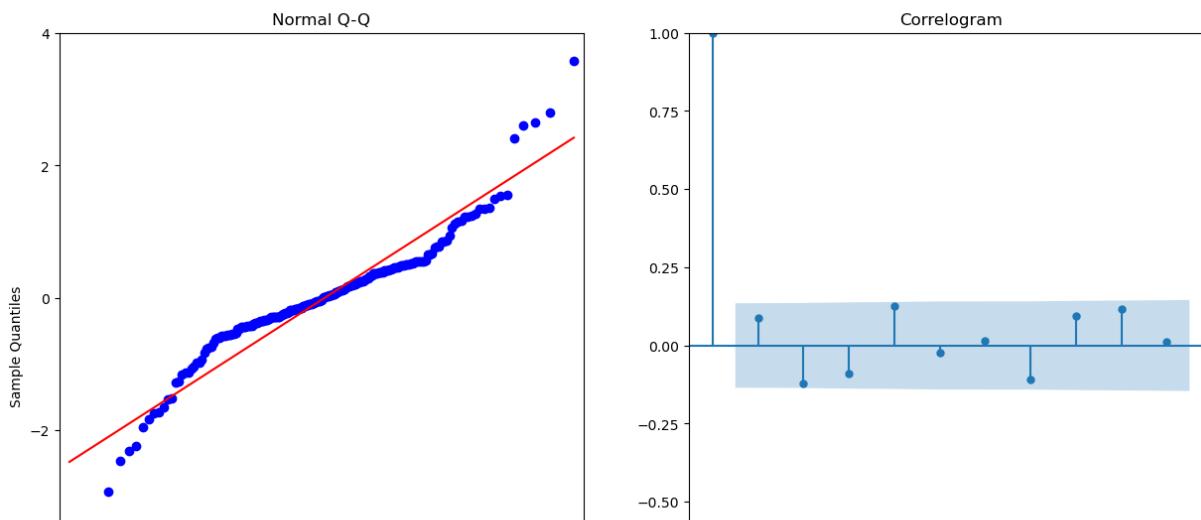
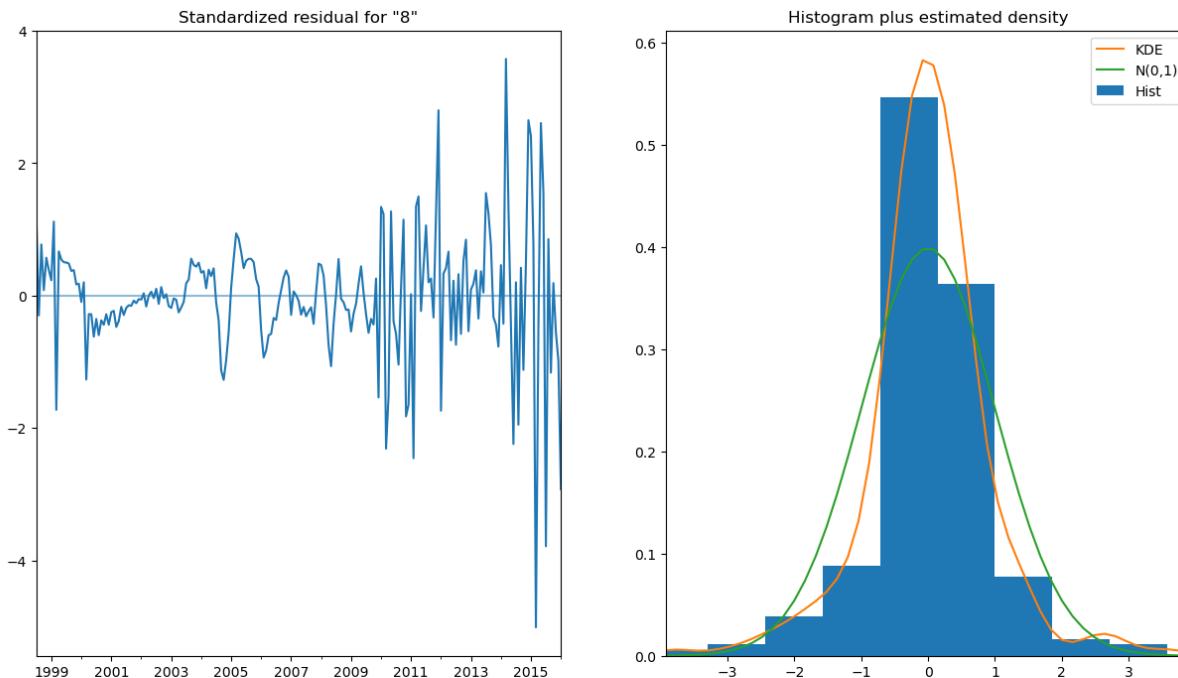
=====

=====

coef	std err	z	P> z	[0.025
------	---------	---	------	---------

0.975]

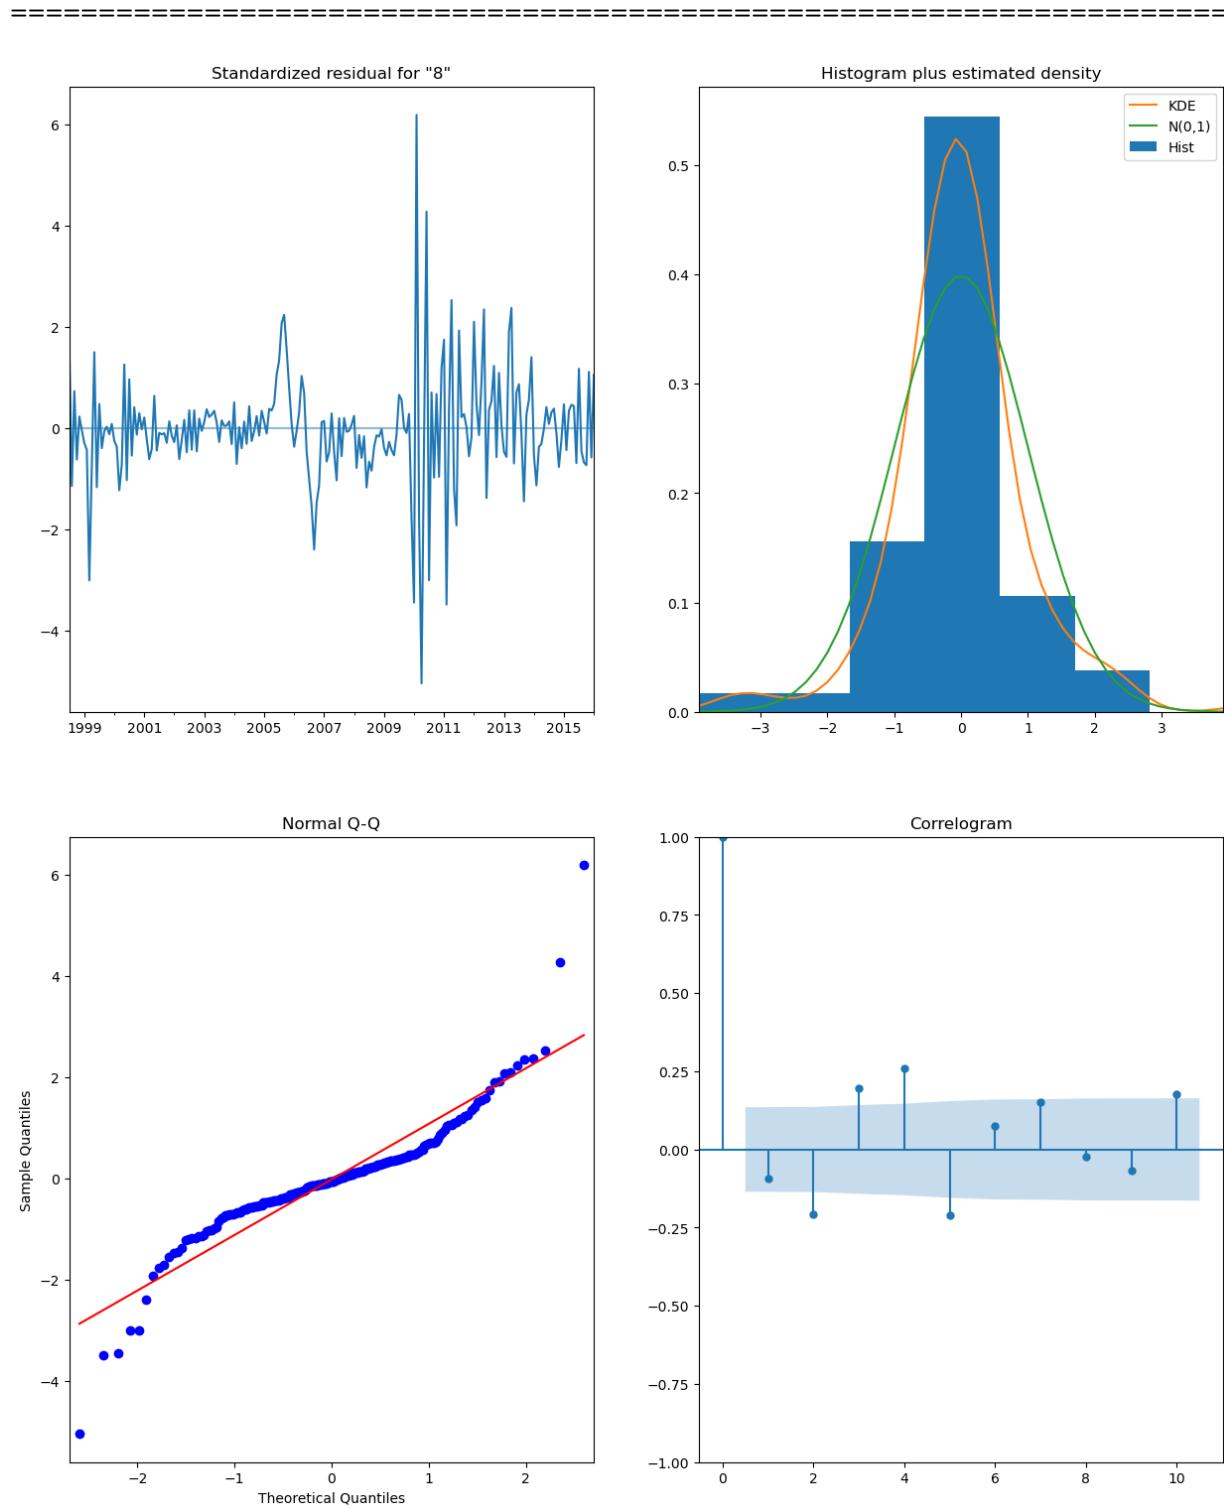
ar.L1	0.7522	0.041	18.431	0.000	0.672
0.832					
ma.L1	0.7636	0.042	18.162	0.000	0.681
0.846					
ar.S.L12	0.2102	0.040	5.236	0.000	0.132
0.289					
ma.S.L12	-0.3916	0.046	-8.568	0.000	-0.481
-0.302					
sigma2	5.411e+06	3.93e+05	13.767	0.000	4.64e+06
6.18e+06					



The Root Mean Squared Error of our forecasts is 137358.0

	coef	std err	z	P> z	[0.025
0.975]					

ar.L1	0.8095	0.033	24.782	0.000	0.746
0.874					
ma.L1	0.5904	0.026	22.479	0.000	0.539
0.642					
ma.S.L12	-0.5333	0.027	-19.633	0.000	-0.587
-0.480					
sigma2	3.998e+05	2.04e+04	19.581	0.000	3.6e+05

4.4e+05

The Root Mean Squared Error of our forecasts is 21177.0

```
[58117.0, 9986.0, 2775.0, 20071.0, 44237.0, 12709.0, 12072.0, 5100.0,  
20612.0, 19611.0, 5250.0, 3654.0, 19365.0, 2626.0, 24559.0, 22105.0,  
17671.0, 31924.0, 137358.0, 21177.0]
```

6.2 RMSE and RMSE Ratio to Median Value

In [14]:

```
1 rmse_list
```

executed in 20ms, finished 10:55:31 2022-05-03

Out[14]:

```
[58117.0,  
 9986.0,  
 2775.0,  
 20071.0,  
 44237.0,  
 12709.0,  
 12072.0,  
 5100.0,  
 20612.0,  
 19611.0,  
 5250.0,  
 3654.0,  
 19365.0,  
 2626.0,  
 24559.0,  
 22105.0,  
 17671.0,  
 31924.0,  
 137358.0,  
 21177.0]
```

In [15]:

```

1 prospects['RMSE'] = rmse_list
2 prospects['RMSE_Ratio'] = prospects['RMSE']/prospects['2018-04']
3 #top_prospects = prospects.loc[[33460,94606,4298,815,5830]]
4 prospects

```

executed in 60ms, finished 10:57:06 2022-05-03

Out[15]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05	1996-06	199
1477	94601	Oakland	CA	San Francisco	Alameda	1478	114600.0	114500.0	114500.0	1144
1239	94590	Vallejo	CA	Vallejo	Solano	1240	108200.0	107800.0	107500.0	1073
2627	33460	Lake Worth	FL	Miami-Fort Lauderdale	Palm Beach	2628	59800.0	59900.0	60000.0	602
1681	94606	Oakland	CA	San Francisco	Alameda	1682	120400.0	120300.0	120300.0	1202
1853	94804	Richmond	CA	San Francisco	Contra Costa	1854	171300.0	170300.0	169400.0	1685
1960	89104	Las Vegas	NV	Las Vegas	Clark	1961	94800.0	94700.0	94700.0	947
474	85008	Phoenix	AZ	Phoenix	Maricopa	475	61600.0	61900.0	62200.0	626

6.3 Create top_prospects df

- Zip Codes Included:
 - Redford, MI 48240
 - Las Vegas, NV 89115
 - Sacramento, CA 95824
 - Oakland, CA 94606
 - Lake Worth, FL 33460

In [44]:

```
1 top_prospects = prospects.loc[[2627,1681,4298,815,5830,5764]]
```

executed in 39ms, finished 13:43:28 2022-05-03

In [45]:

```

1 # drop RMSE and RMSE_Ratio columns
2 top_prospects = top_prospects.drop(columns=['RMSE', 'RMSE_Ratio'])
3 top_prospects

```

executed in 69ms, finished 13:43:29 2022-05-03

Out[45]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05
2627	33460	Lake Worth	FL	Miami-Fort Lauderdale	Palm Beach	2628	59800.0	59900.0
1681	94606	Oakland	CA	San Francisco	Alameda	1682	120400.0	120300.0
4298	95824	Sacramento	CA	Sacramento	Sacramento	4299	73800.0	73400.0
815	89115	Las Vegas	NV	Las Vegas	Clark	816	92500.0	92500.0
5830	48240	Redford	MI	Detroit	Wayne	5831	67800.0	68200.0
5764	83703	Boise	ID	Boise City	Ada	5765	107500.0	107500.0

6 rows × 271 columns

7 Price Forecast - Top Zip Codes

- 5 year investment horizon

In [74]:

```

1 # Create list for final predicted value
2 pred_list_1 = []
3
4 # Construct model from entirety of data - no train/test
5 # iterate over prospects to extract optimal SARIMA values,
6 # instantiate SARIMAX model, make predictions
7 # and produce forecast with confidence intervals
8 for zipcode in top_prospects['ZipCode']:
9     zips = top_prospects.loc[top_prospects['ZipCode'] == zipcode]
10    m1 = melt_data(zips)
11    m1[str(zipcode)] = m1['value']
12    m1 = m1.drop('value', axis=1)
13
14    # do NOT create train and test sets
15    #train = m1.iloc[:238]
16    #test = m1.iloc[238:]
17
18    # Run a grid with pdq and seasonal pdq parameters calculated above and ge
19    ans = []
20    outputs = []
21    comb_list = []
22    combs_list = []
23
24    for comb in pdq:
25        for combs in pdqs:
26            try:
27                mod = sm.tsa.statespace.SARIMAX(m1,
28                                                order=comb,
29                                                seasonal_order=combs,
30                                                enforce_stationarity=False,
31                                                enforce_invertibility=False)
32
33                output = mod.fit()
34                ans.append([comb, combs, output.aic])
35                outputs.append([output.aic])
36                comb_list.append([comb])
37                combs_list.append([combs])
38                AIC_df = pd.DataFrame(list(outputs, comb_list, combs_list))
39                #print('ARIMA {} x {}12 : AIC Calculated ={}'.format(comb, co
40            except:
41                continue
42
43    # Print the lowest AIC and it's parameters for SARIMA
44    AIC_df[0] = AIC_df[0].str.get(0)
45    best_AIC = AIC_df.copy()
46    lowest_AIC = min(best_AIC[0])
47    lowest_AIC_row = AIC_df.loc[AIC_df[0] == lowest_AIC]
48    #print(lowest_AIC_row)
49
50    # call value for comb and combs for model input
51    lowest_AIC_row[1] = lowest_AIC_row[1].str.get(0)
52    lowest_AIC_row[2] = lowest_AIC_row[2].str.get(0)

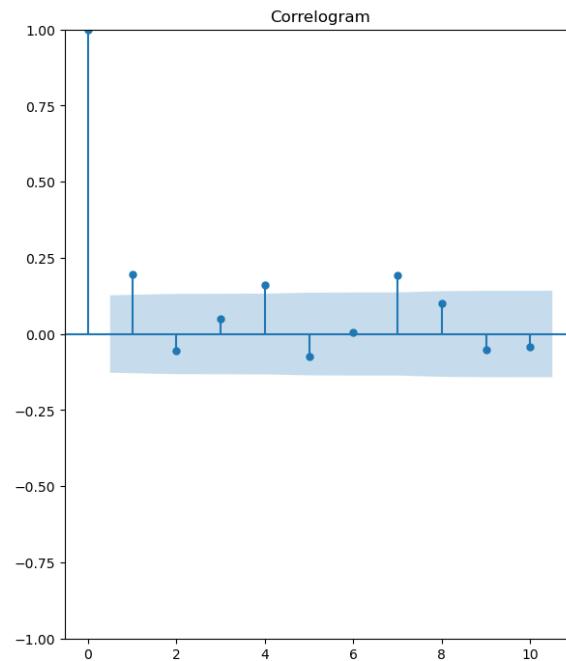
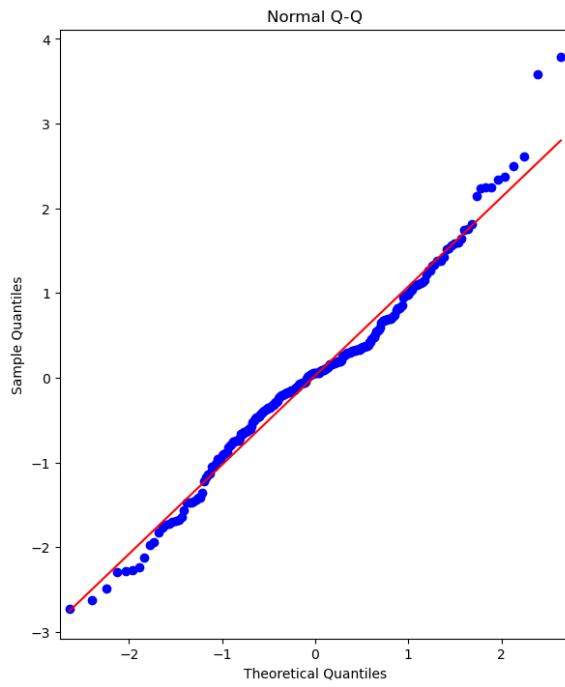
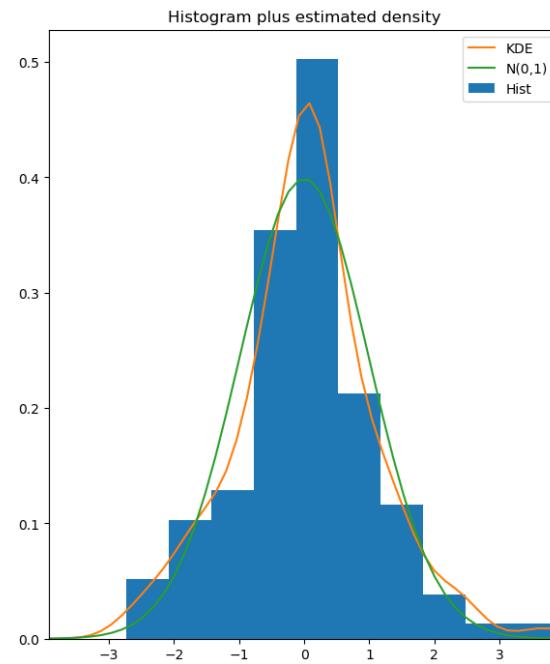
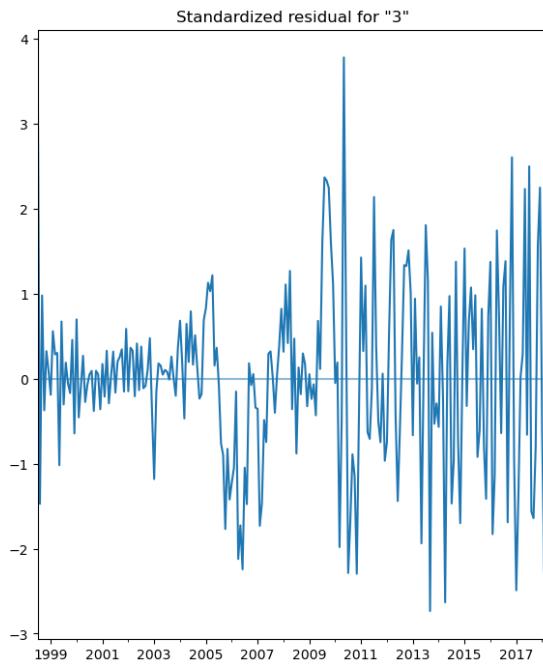
```

```
53     arima = lowest_AIC_row.iat[0,1]
54     s = lowest_AIC_row.iat[0,2]
55     print(str(m1.columns.values))
56     print(arima, s)
57
58     # Plug the optimal parameter values into a new SARIMAX model
59     ARIMA_MODEL = sm.tsa.statespace.SARIMAX(m1,
60                                              order=arima,
61                                              seasonal_order=s,
62                                              enforce_stationarity=False,
63                                              enforce_invertibility=False)
64
65     # Fit the model and print results
66     output = ARIMA_MODEL.fit()
67
68     print(output.summary().tables[1])
69
70     # Call plot_diagnostics() on the results calculated above
71     output.plot_diagnostics(figsize=(15, 18))
72     plt.show()
73
74     # Get actual results in dataframe
75     price_truth = m1['2018-04-01':]
76
77     # Get forecast 500 steps ahead in future
78     prediction = output.get_forecast(steps=60)
79
80     # Create predictions df
81     pred = prediction.predicted_mean
82
83     # Add prediction for 5 year forecasted value to list
84     pred_list_1.append(pred['2023-04-01'])
85
86     # Get confidence intervals of forecasts
87     pred_conf = prediction.conf_int()
88
89     # Only use data from April 2010 to improve visual of forecast
90     plot_vals = m1.loc['2010-04-01':]
91
92     # Plot future predictions with confidence intervals
93     ax = plot_vals.plot(label='observed', figsize=(20, 15), color='black')
94     prediction.predicted_mean.plot(ax=ax, label='Forecast', color='green')
95     ax.fill_between(pred_conf.index,
96                      pred_conf.iloc[:, 0],
97                      pred_conf.iloc[:, 1], color='k', alpha=0.25)
98     ax.set_xlabel('Date').set_fontsize(30)
99     ax.set_ylabel('Median House Price ($)').set_fontsize(30)
100    ax.set_title('Five Year Median House Price Forecast').set_fontsize(40)
101
102   for item in (ax.get_xticklabels() + ax.get_yticklabels()):
103       item.set_fontsize(20)
104
105   right_side = ax.spines["right"]
106   right_side.set_visible(False)
```

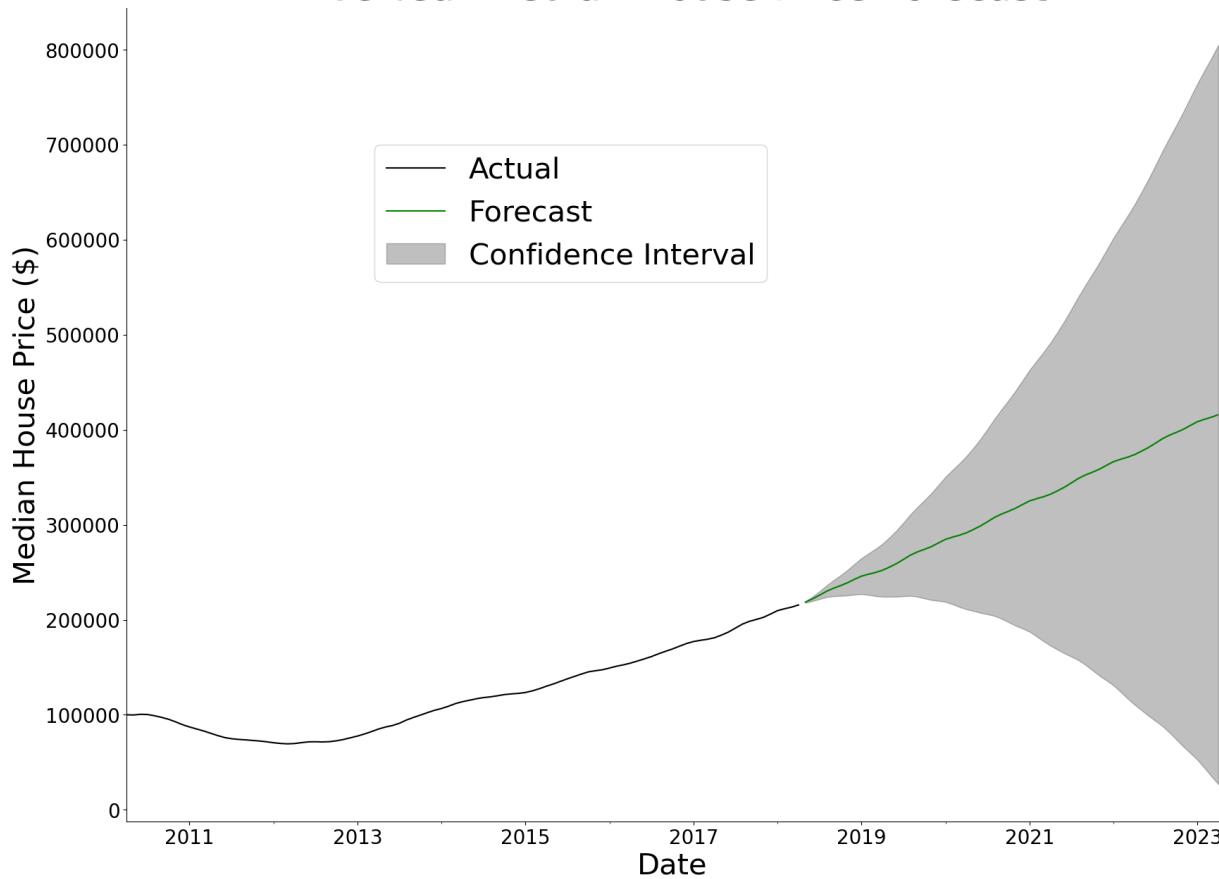
```
107
108     top_side = ax.spines["top"]
109     top_side.set_visible(False)
110
111     mylabels = 'Actual', 'Forecast', 'Confidence Interval'
112     plt.legend(fontsize=30, labels=mylabels, bbox_to_anchor=(0.6,0.85))
113     plt.show()
114     print()
115     print()
116
117
```

executed in 3m 10s, finished 14:01:19 2022-05-03

```
['33460']
(1, 1, 1) (1, 1, 1, 12)
=====
=====
          coef      std err           z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.9646      0.017      57.776      0.000      0.932
0.997
ma.L1      0.5143      0.022      23.638      0.000      0.472
0.557
ar.S.L12   -0.5424      0.065     -8.383      0.000     -0.669
-0.416
ma.S.L12   14.5237     12.200      1.190      0.234     -9.389
38.436
sigma2     925.0061    1559.321      0.593      0.553    -2131.207
3981.219
=====
```

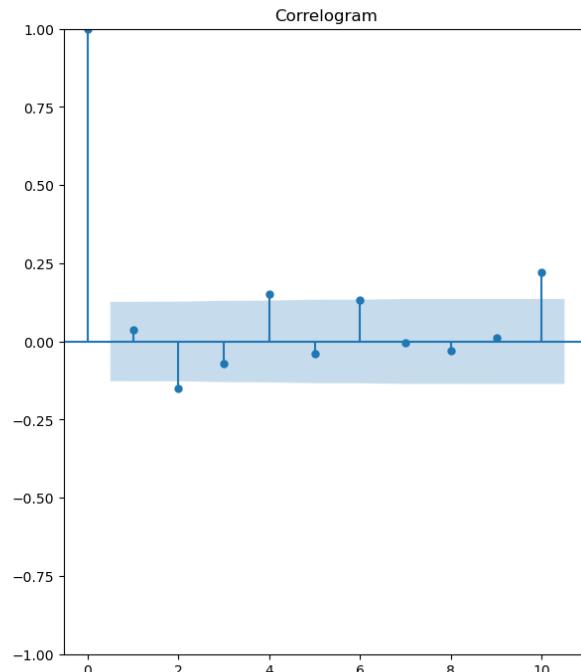
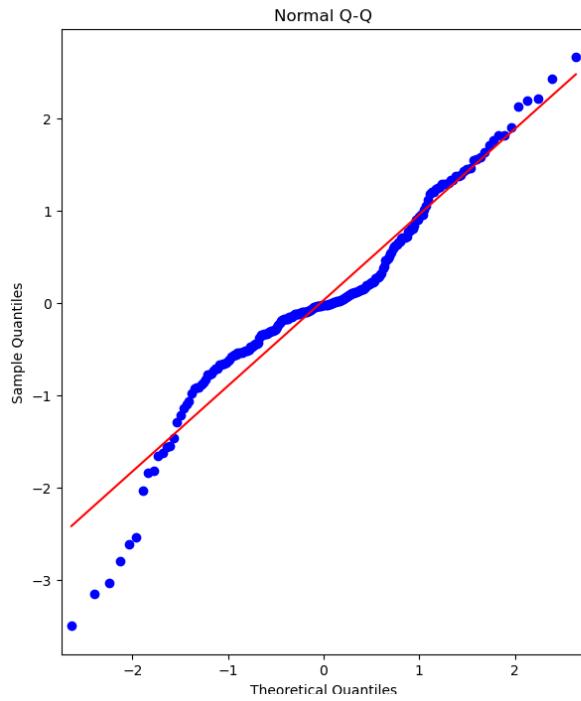
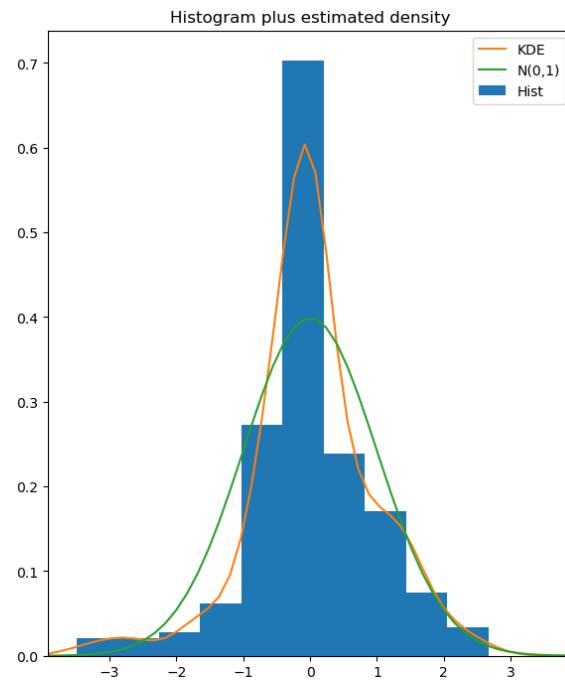
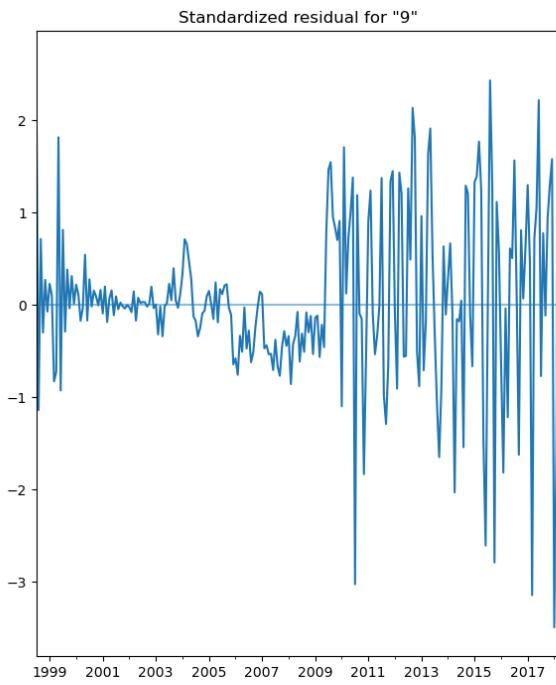


Five Year Median House Price Forecast

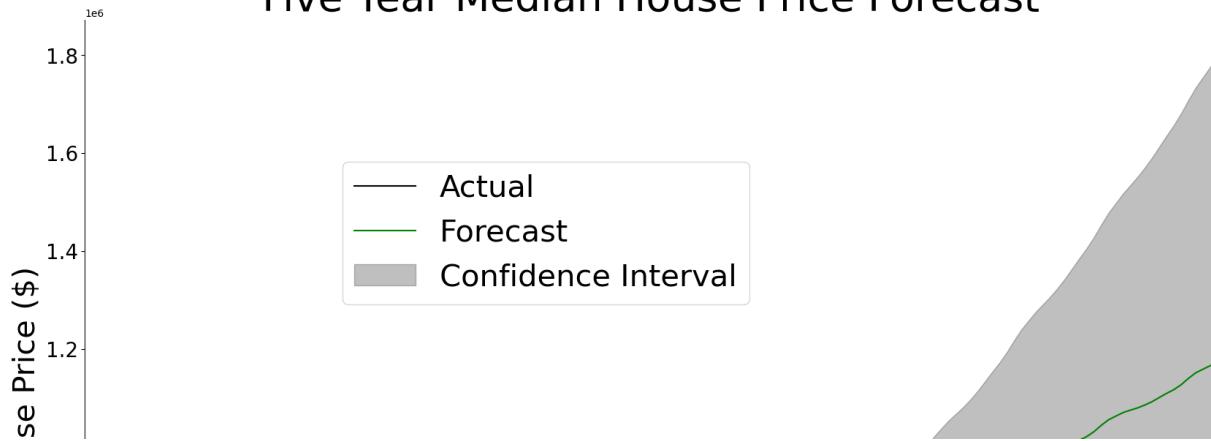


```

['94606']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [0.025
0.975]
-----
ar.L1        0.8335      0.039     21.543      0.000      0.758
0.909
ma.L1        0.5934      0.048     12.337      0.000      0.499
0.688
ma.S.L12     -0.2920      0.037     -7.855      0.000     -0.365
-0.219
sigma2       3.391e+06   3.21e+05    10.559      0.000     2.76e+06
4.02e+06
=====
=====
```

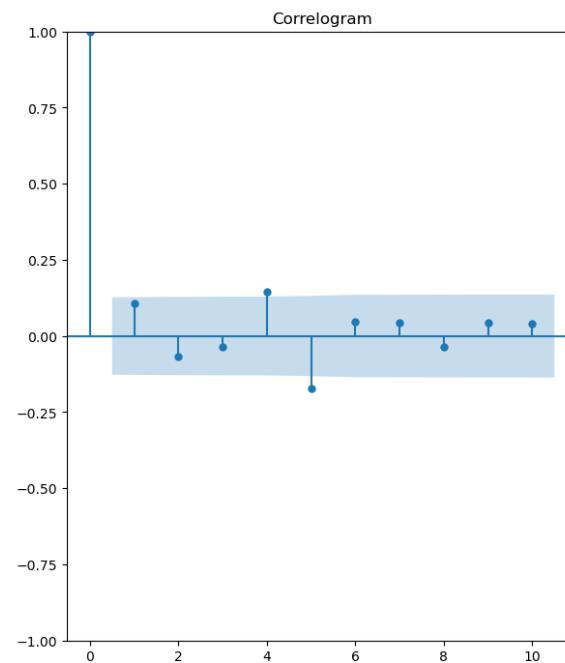
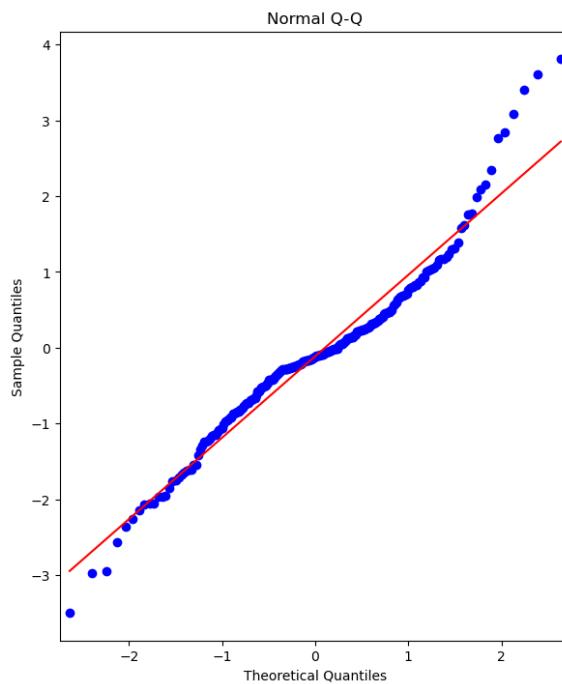
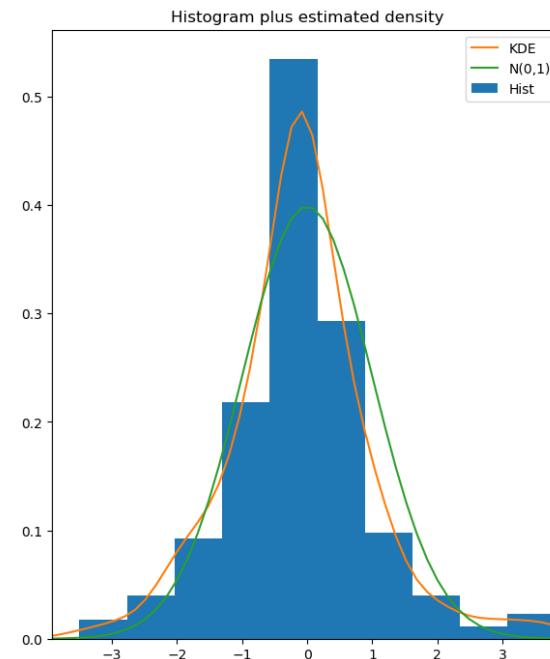
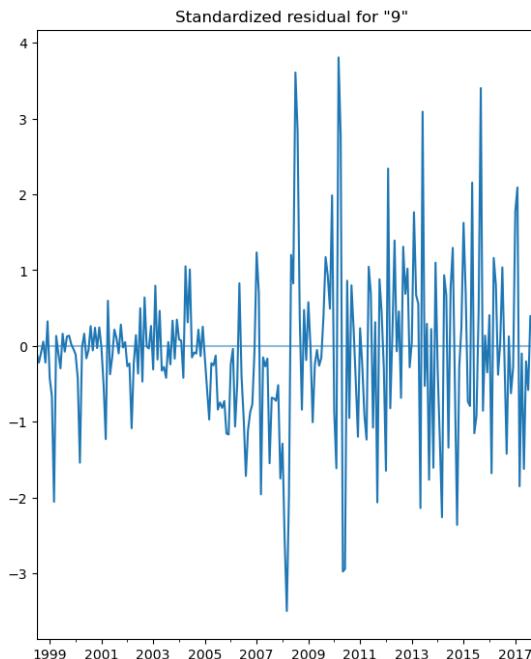


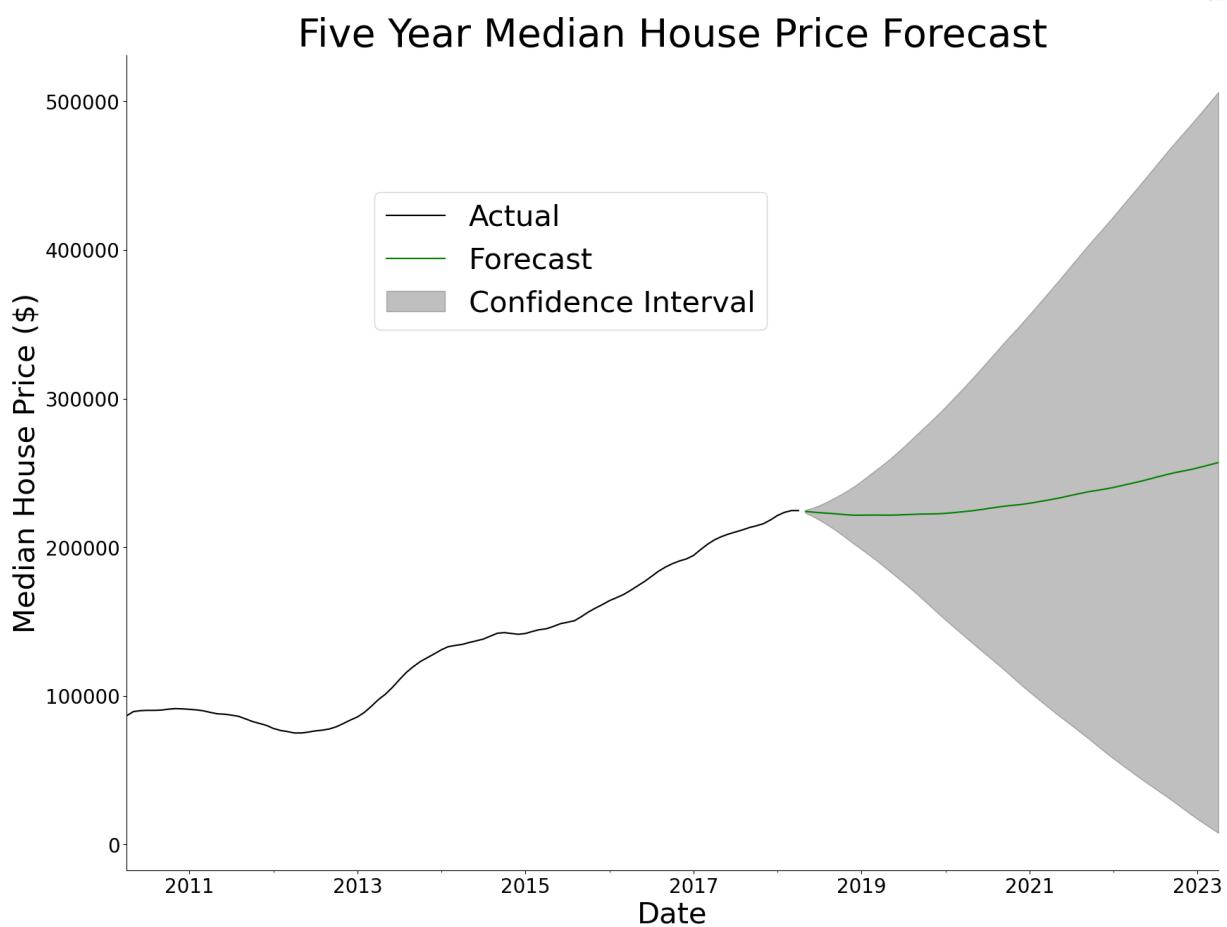
Five Year Median House Price Forecast



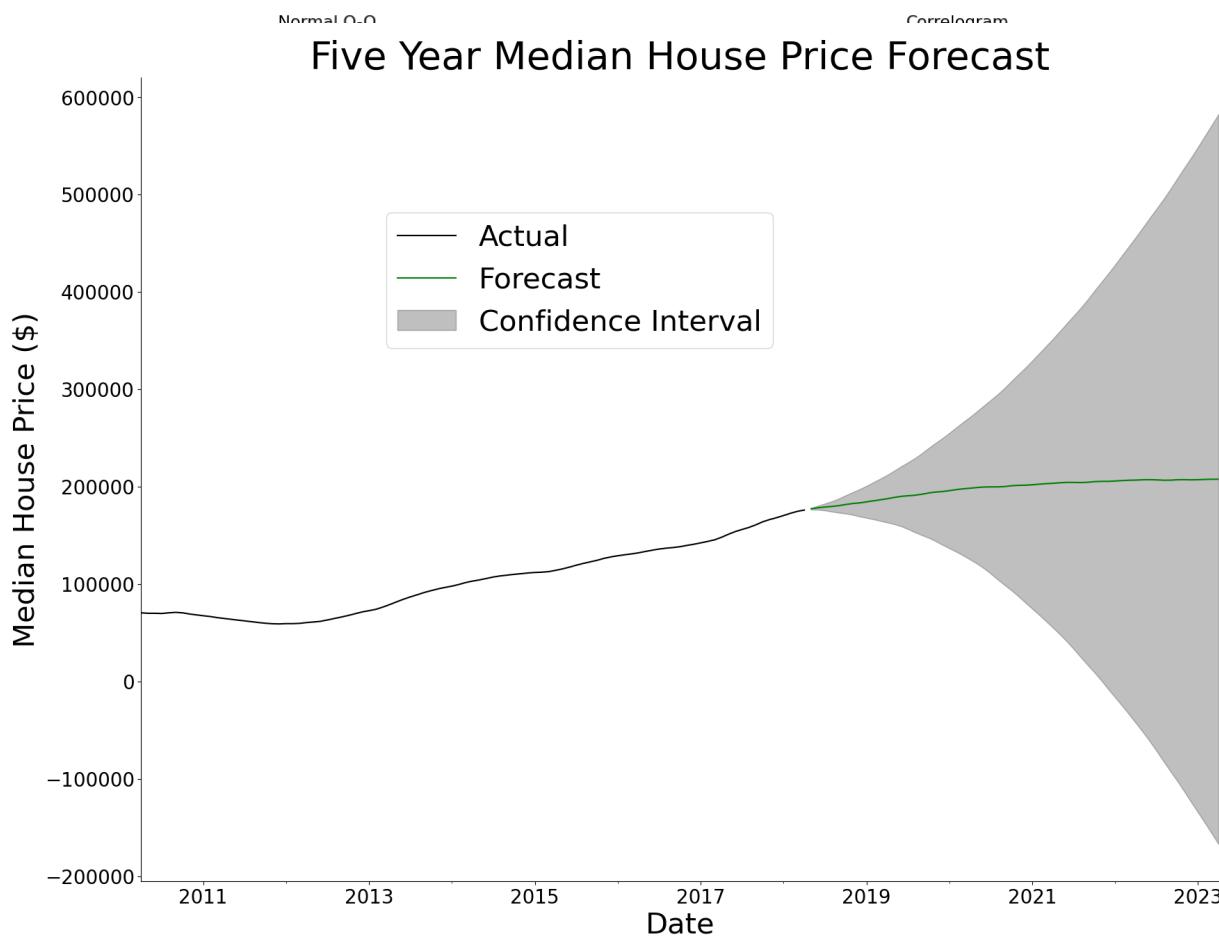
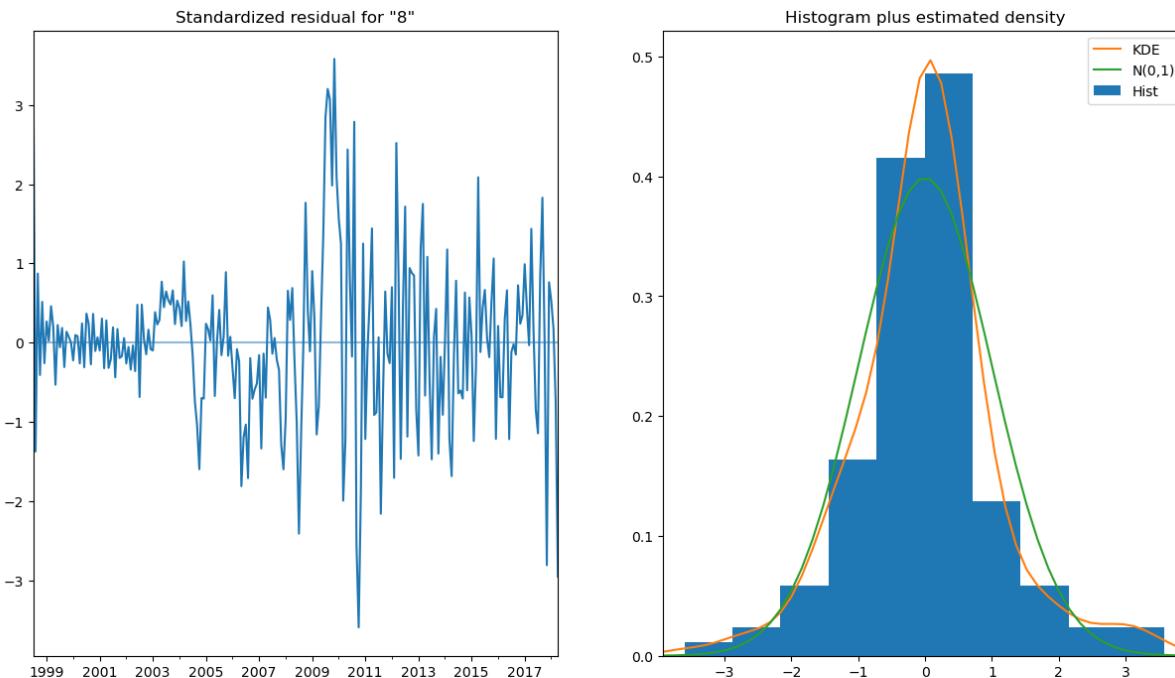
```
['95824']
(1, 1, 1) (1, 1, 1, 12)
=====
```

	coef	std err	z	P> z	[0.025
0.975]					
ar.L1	0.9601	0.012	82.851	0.000	0.937
0.983					
ma.L1	0.6136	0.038	15.984	0.000	0.538
0.689					
ar.S.L12	0.1000	0.030	3.377	0.001	0.042
0.158					
ma.S.L12	-0.9692	0.104	-9.328	0.000	-1.173
-0.766					
sigma2	2.585e+05	3.27e+04	7.913	0.000	1.94e+05
3.23e+05					

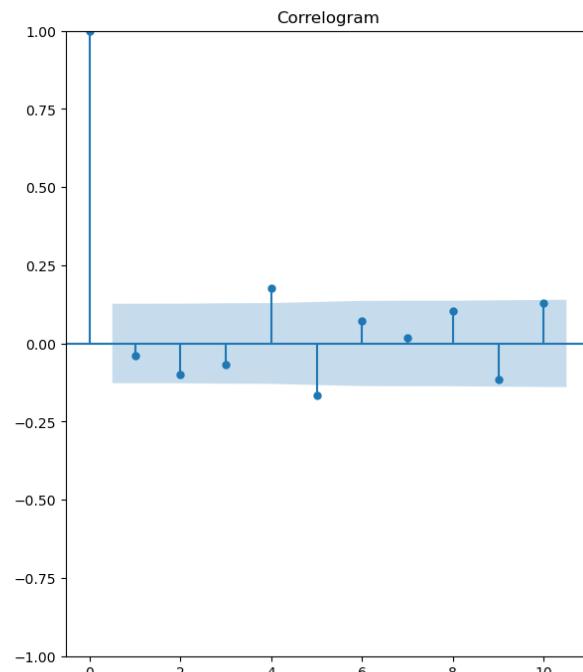
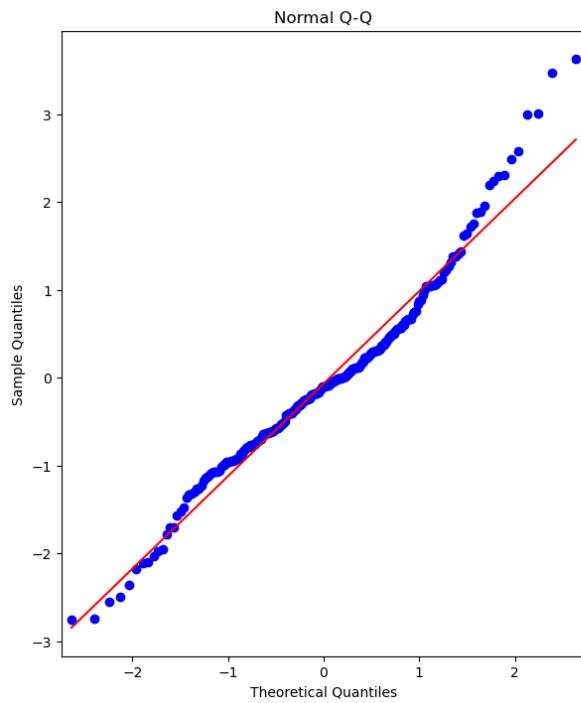
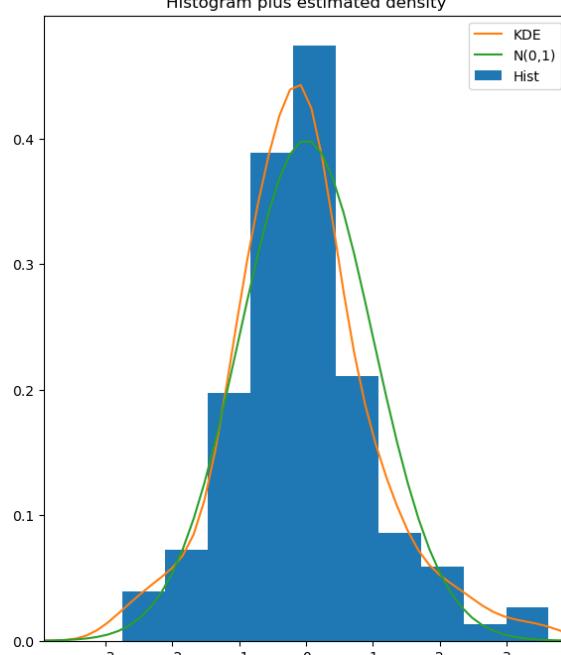
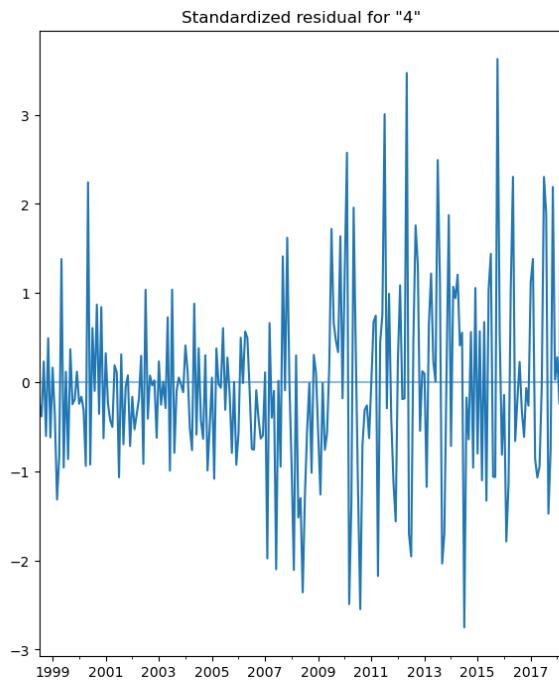




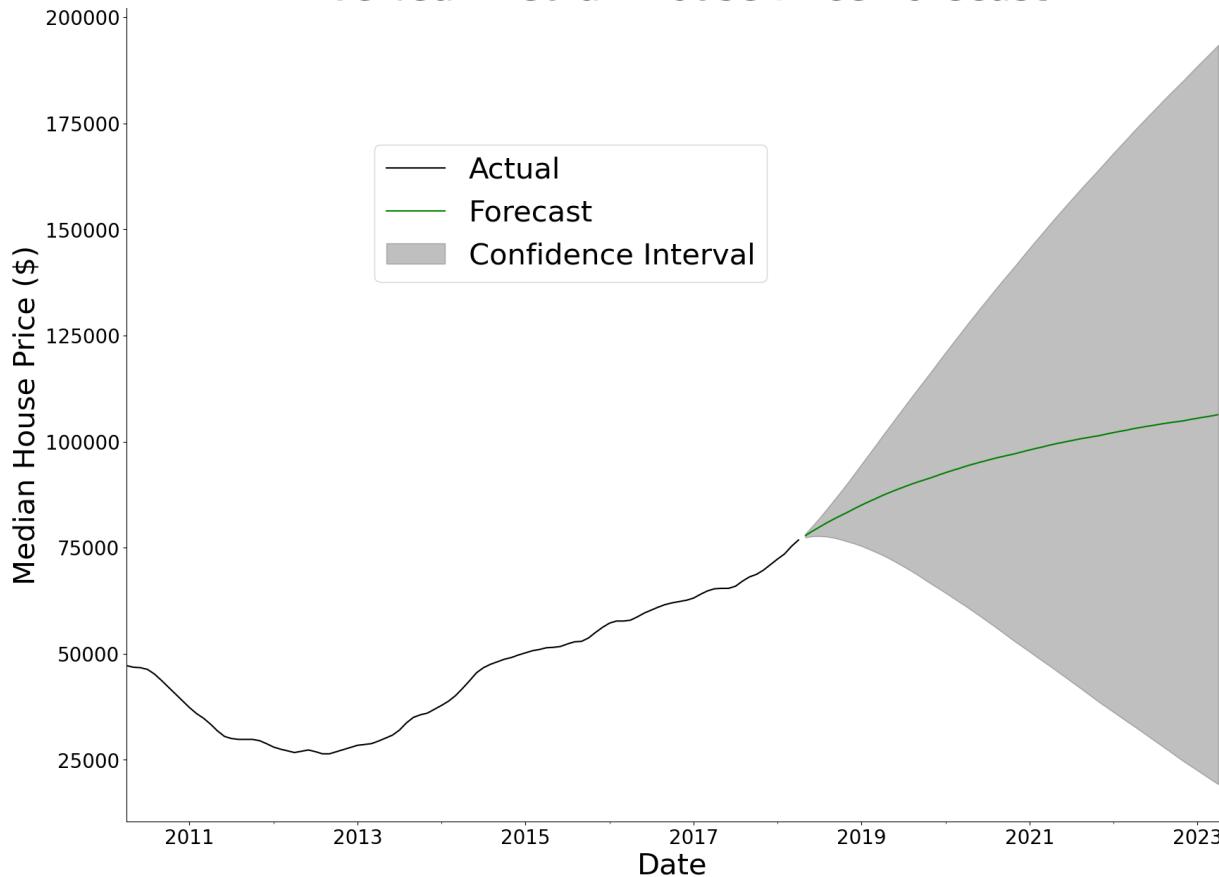
```
[ '89115' ]
(1, 1, 1) (1, 1, 1, 12)
=====
=====
          coef      std err       z     P>|z|      [ 0.025
0.975]
-----
ar.L1      0.9733    0.012    77.911    0.000    0.949
0.998
ma.L1      0.3587    0.019    19.009    0.000    0.322
0.396
ar.S.L12   -0.4685   0.074    -6.347    0.000   -0.613
-0.324
ma.S.L12   25.3835   42.137    0.602    0.547   -57.203
107.970
sigma2     264.1590  880.163    0.300    0.764  -1460.928
1989.246
=====
```



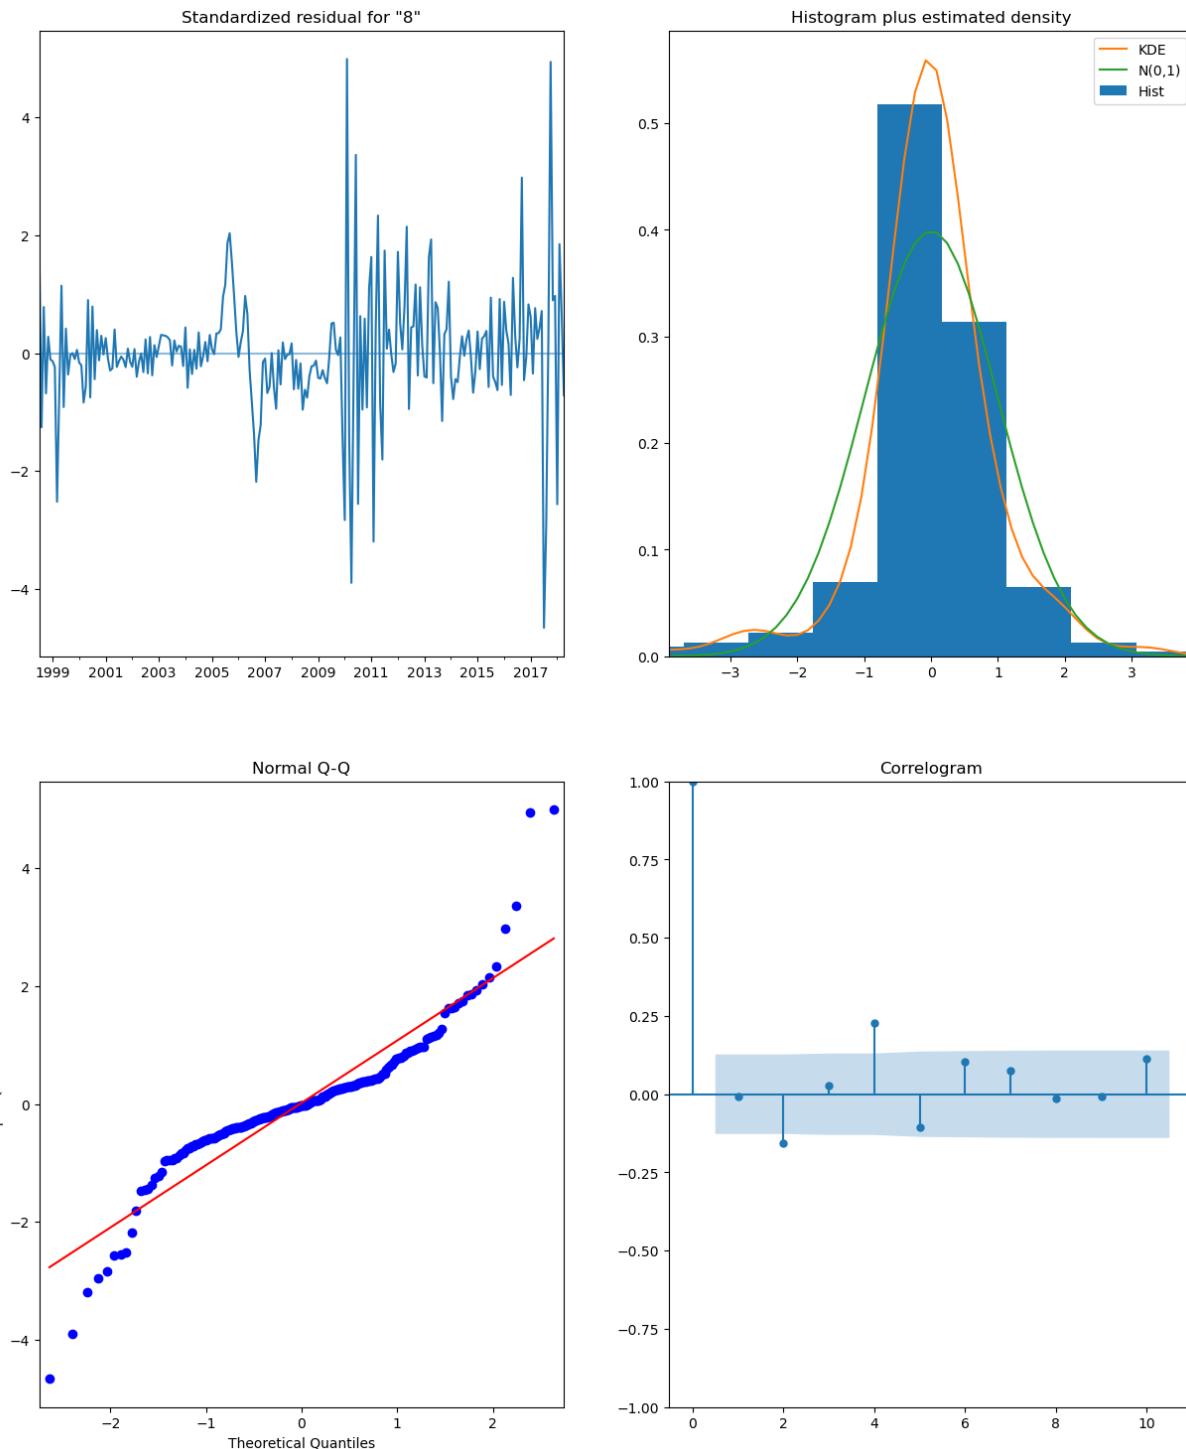
```
['48240']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
            coef      std err          z      P>|z|      [0.025
0.975]
-----
ar.L1        0.9519      0.020      47.106      0.000      0.912
0.991
ma.L1        0.5049      0.050      10.011      0.000      0.406
0.604
ma.S.L12     -0.9968      1.096      -0.910      0.363     -3.144
1.150
sigma2       5.313e+04    5.72e+04      0.929      0.353     -5.9e+04
1.65e+05
=====
```

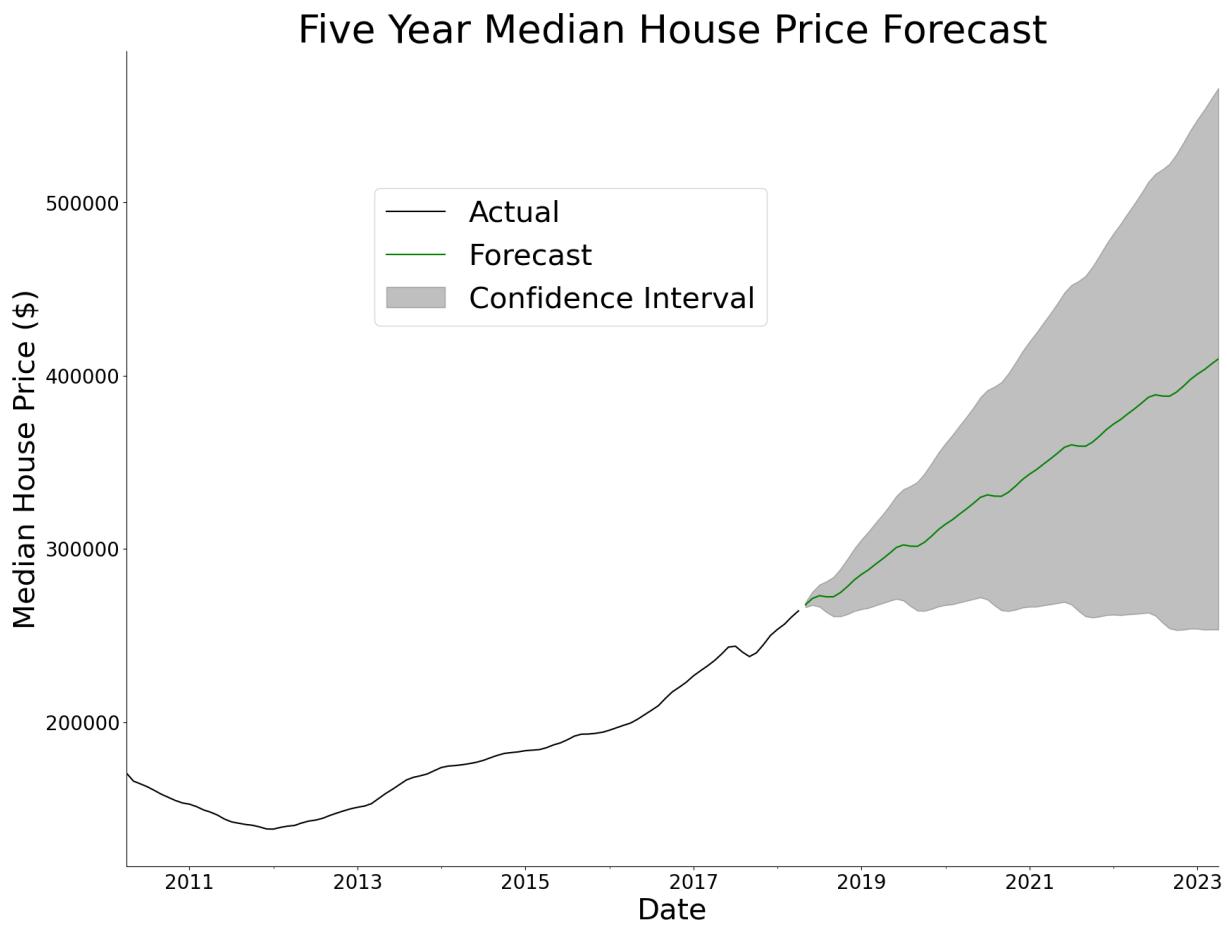


Five Year Median House Price Forecast



```
[ '83703']
(1, 1, 1) (0, 1, 1, 12)
=====
=====
          coef      std err          z      P>|z|      [ 0.025
0.975]
-----
ar.L1      0.7428      0.038     19.304      0.000      0.667
0.818
ma.L1      0.6268      0.033     18.860      0.000      0.562
0.692
ma.S.L12   -0.4391      0.032    -13.887      0.000     -0.501
-0.377
sigma2     5.895e+05  3.31e+04     17.792      0.000     5.25e+05
6.54e+05
```





In [75]:

```
1 #pred_list.append(pred['2023-04-01'])
2 pred_list_1
```

executed in 32ms, finished 14:02:41 2022-05-03

Out[75]:

```
[416016.39007328794,
 1171061.0561964822,
 257003.9336061536,
 207692.41758930407,
 106363.34240621062,
 409711.22703533754]
```

7.1 Calculate predicted rate of return after 5 years

In [76]:

```

1 # Create column for predicted price after 5 years
2 top_prospects['5_Year_Price'] = pred_list_1
3
4 # Calculate Percent Return from time of investment
5 top_prospects['Rate_Return'] = ((
6     top_prospects['5_Year_Price']-
7     top_prospects['2018-04'])/
8     top_prospects['2018-04'])*100
9
10 # Bar Chart of Rate of Return by Zip Code
11 #plt.bahr(top_prospects['Rate_Return'], align='center')
12

```

executed in 27ms, finished 14:02:43 2022-05-03

In [77]:

```
1 top_prospects
```

executed in 59ms, finished 14:02:44 2022-05-03

Out[77]:

	ZipCode	City	State	Metro	CountyName	SizeRank	1996-04	1996-05
2627	33460	Lake Worth	FL	Miami-Fort Lauderdale	Palm Beach	2628	59800.0	59900.0
1681	94606	Oakland	CA	San Francisco	Alameda	1682	120400.0	120300.0
4298	95824	Sacramento	CA	Sacramento	Sacramento	4299	73800.0	73400.0
815	89115	Las Vegas	NV	Las Vegas	Clark	816	92500.0	92500.0
5830	48240	Redford	MI	Detroit	Wayne	5831	67800.0	68200.0
5764	83703	Boise	ID	Boise City	Ada	5765	107500.0	107500.0

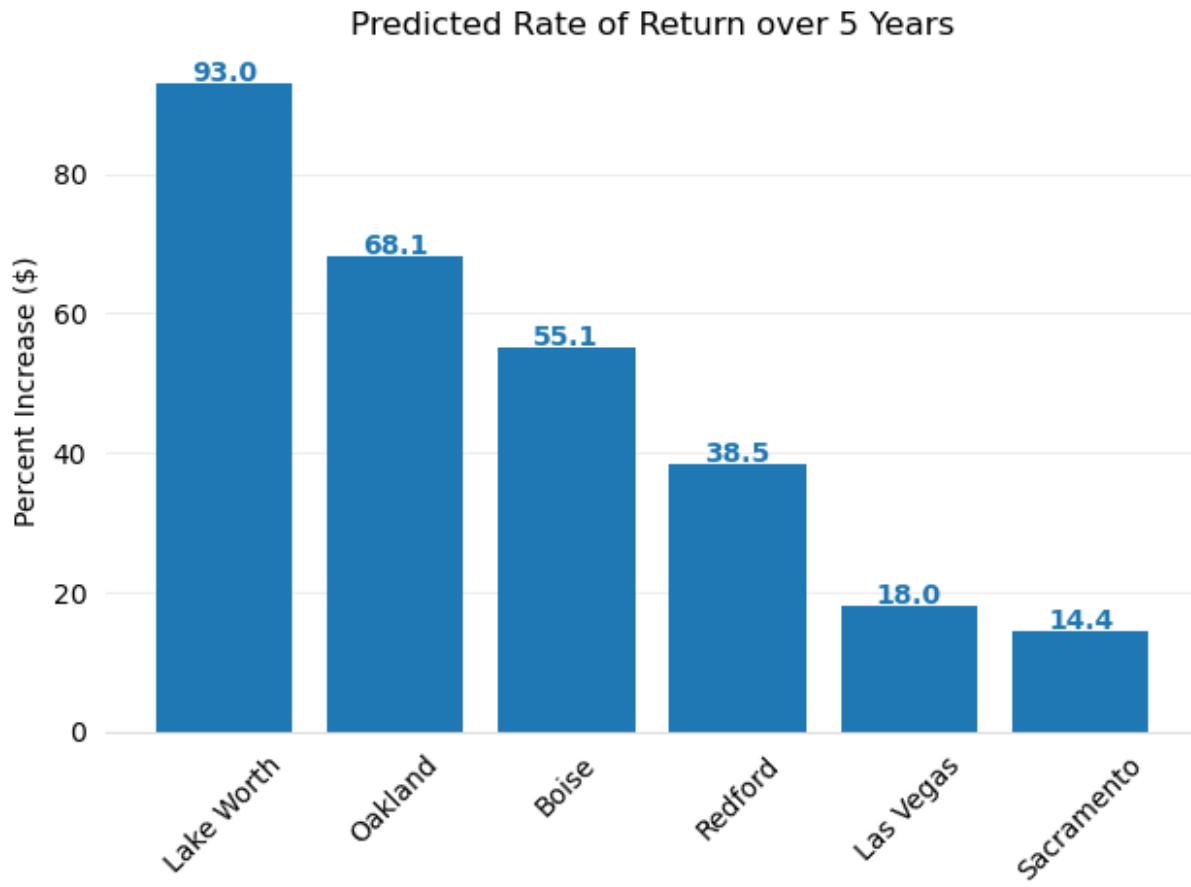
6 rows × 273 columns

7.2 Bar Chart: Rate of Return

In [79]:

```
1 # create df and sort by Rate_Return in Descending Order
2 rate_1 = top_prospects[['Rate_Return','City']]
3 #df.sort_values('importances',inplace=True)
4 rate_1.sort_values('Rate_Return',inplace=True,ascending=False)
5
6 fig, ax = plt.subplots()
7
8 # Save the chart so we can loop through the bars below.
9 bars = ax.bar(
10     x=np.arange(rate_1['Rate_Return'].size),
11     height=rate_1['Rate_Return'],
12     tick_label=rate_1['City']
13 )
14
15 # Axis formatting.
16 ax.spines['top'].set_visible(False)
17 ax.spines['right'].set_visible(False)
18 ax.spines['left'].set_visible(False)
19 ax.spines['bottom'].set_color('#DDDDDD')
20 ax.tick_params(bottom=False, left=False)
21 ax.set_axisbelow(True)
22 ax.yaxis.grid(True, color='#EEEEEE')
23 ax.xaxis.grid(False)
24
25 # Grab the color of the bars so we can make the
26 # text the same color.
27 bar_color = bars[0].get_facecolor()
28
29 # Add text annotations to the top of the bars.
30 # Note, you'll have to adjust this slightly (the 0.3)
31 # with different data.
32 for bar in bars:
33     ax.text(
34         bar.get_x() + bar.get_width() / 2,
35         bar.get_height() + 0.3,
36         round(bar.get_height(), 1),
37         horizontalalignment='center',
38         color=bar_color,
39         weight='bold'
40     )
41
42 plt.xticks(rotation = 45)
43 plt.title('Predicted Rate of Return over 5 Years')
44 plt.ylabel('Percent Increase ($)')
45
46 fig.tight_layout()
```

executed in 330ms, finished 14:03:20 2022-05-03



8 Recommendations

8.1 Oakland, CA 94606

- Has a predicted rate of return of 68.1% over the next 5 years (assuming we are in April of 2018 when this dataset was current)
- Has the lowest margin of risk - with the lowest confidence interval value being substantially higher than the models for other zip codes.
- High level of model accuracy:
 - Predictions within 3% of median value on average.

8.2 Lake Worth, FL 33460

- Highest predicted return of 93.0%
- Much higher risk with wide range and large negative potential for the confidence interval price at 5 years
- High level of model accuracy:
 - Predictions within 1.3% of median value on average.

8.3 Boise, ID 83703

- Fairly high predicted return of 55.1%
- Low risk, with lowest confidence interval value being substantially higher than the models for other zip codes.
- Moderate level of model accuracy:
 - Predictions within 8% of median value on average.

9 Further Study

9.1 Adding to the model

- Need more recent data
- Add external influences, such as:
 - Average interest rates on mortgages
 - Overall employment, nationally or locally
 - Percent Increase or decrease in jobs in the county of that zip code

9.2 Ongoing testing

- Need monitoring and testing of price changes in real-time
- Weekly updates for price changes

9.3 Additional zip codes

- Most big cities, such as New York, Los Angeles, and Chicago aren't represented.
 - These big markets should be tracked alongside the top returning markets.

9.4 Top returning zip codes

- Find a variety of time periods to qualify the top returning zip codes for investment options.